instaclustr

Introducing Apache Kafka®

Paul Brebner
Technical Evangelist

February 2019

What is Kafka?

instaclustr

Distributed streams Processing System

Messages sent by
Distributed Producers
to
Distributed Consumers
Consumers
via
Distributed Kafka Cluster
Cluster

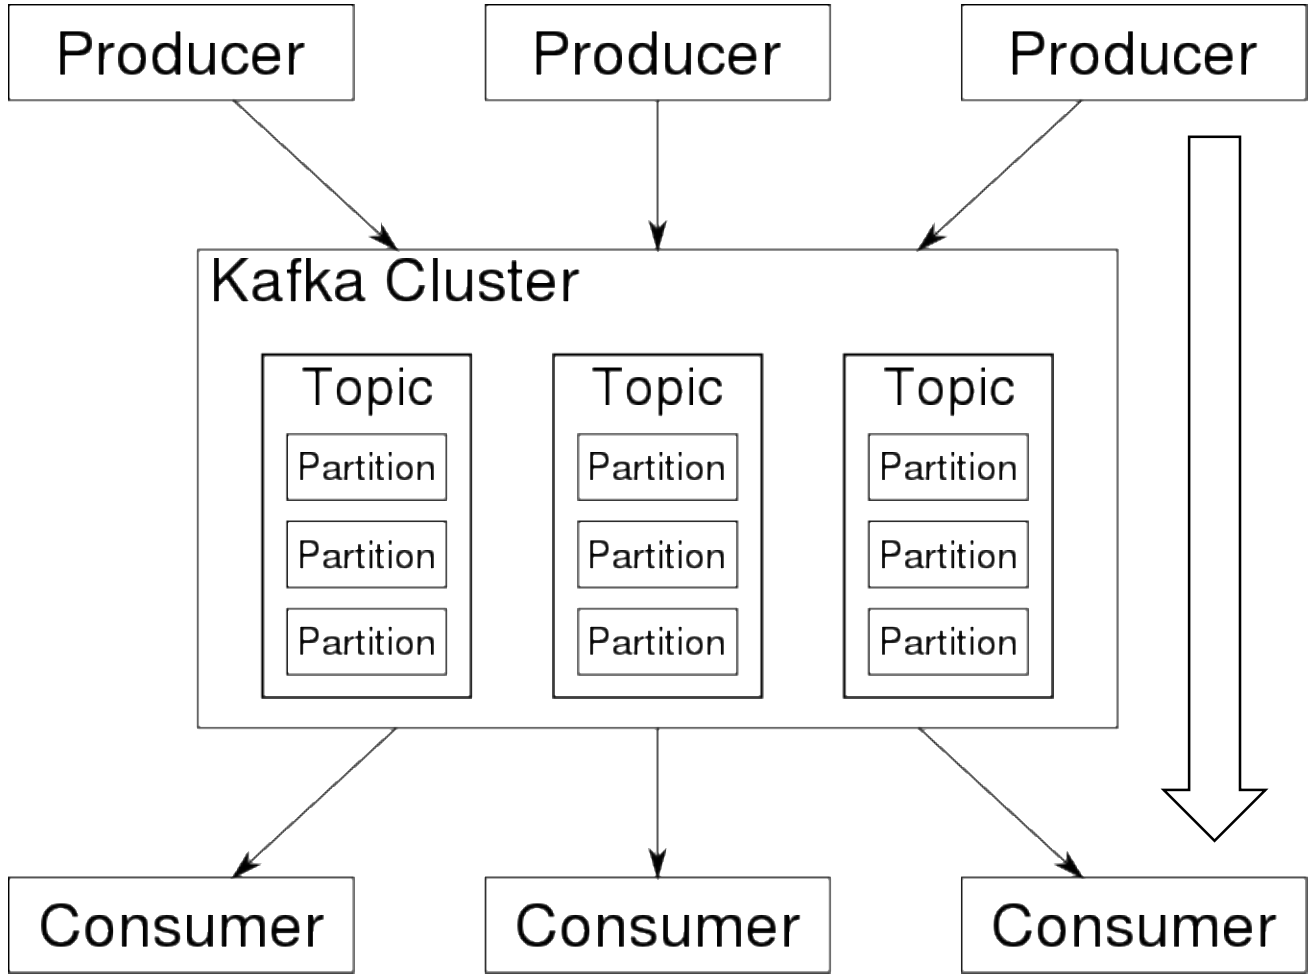

Kafka benefits

instaclustr

- Fast
- Scalable
- Reliable
- Durable
- Open Source
- Managed Service

Kafka benefits

- Fast – high throughput and low latency
- Scalable – horizontally scalable with nodes and partitions
- Reliable – distributed and fault tolerant
- Durable - zero data loss, messages persisted to disk with immutable log
- Open Source – An Apache project
- Available as a Managed Service - on multiple cloud platforms

instaclustr

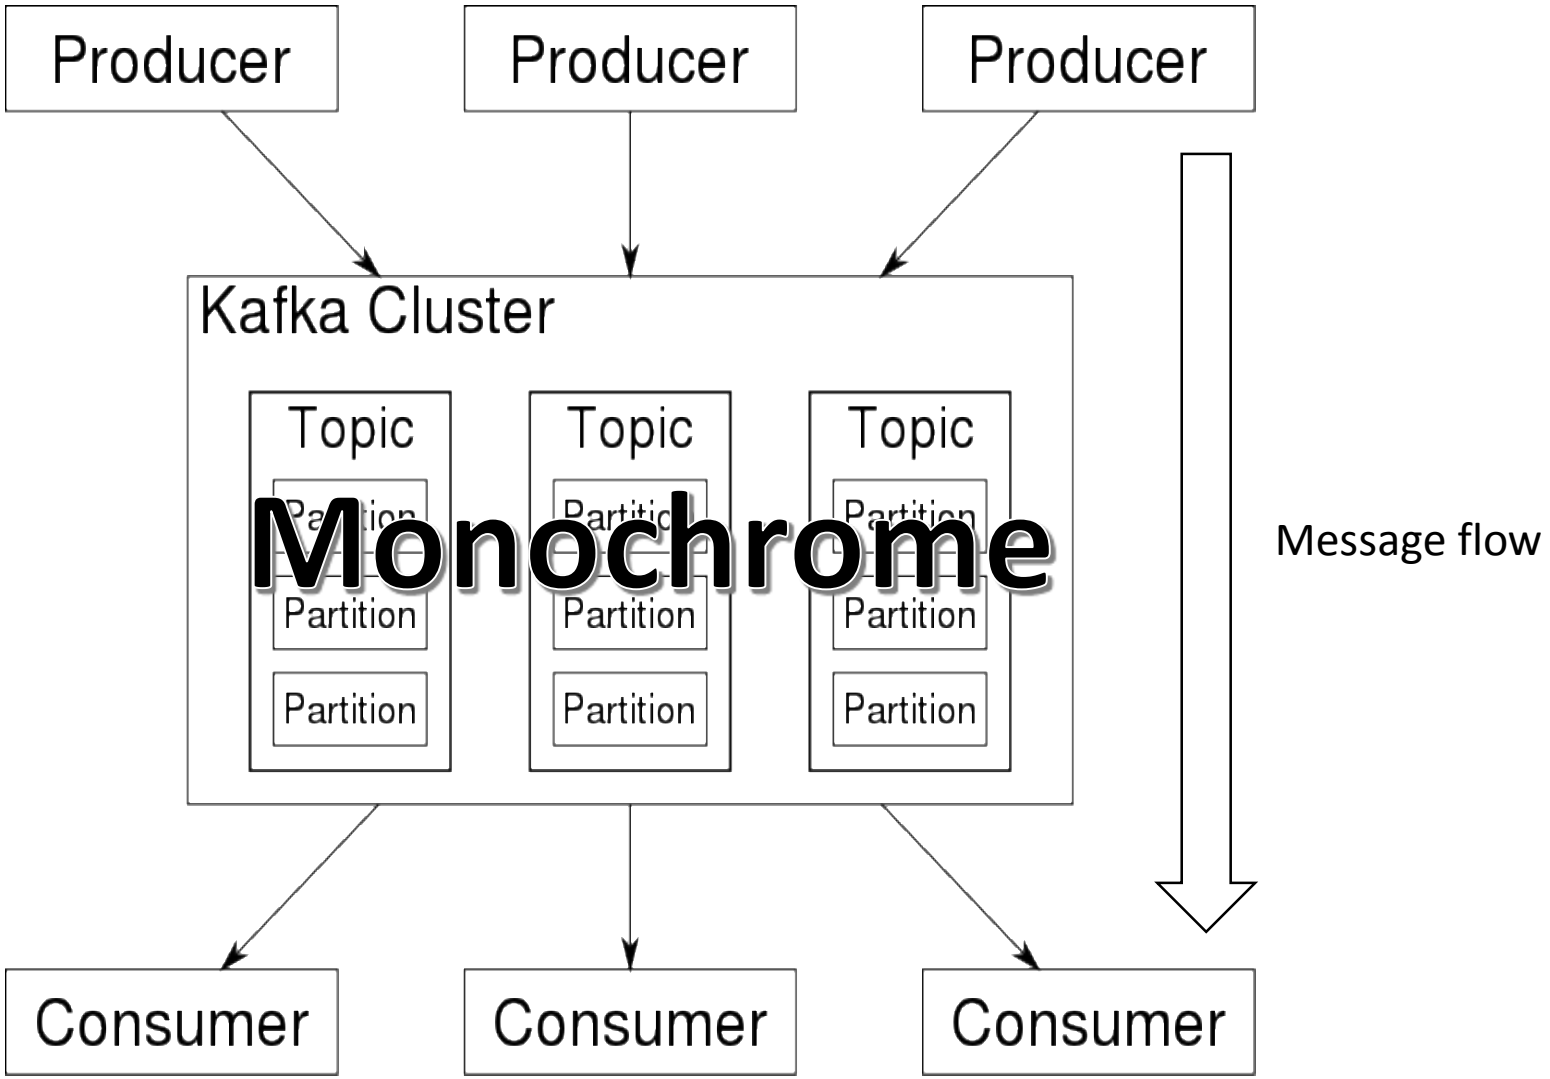

Monochrome

This talk

Colourful

Extended Story

instaclustr

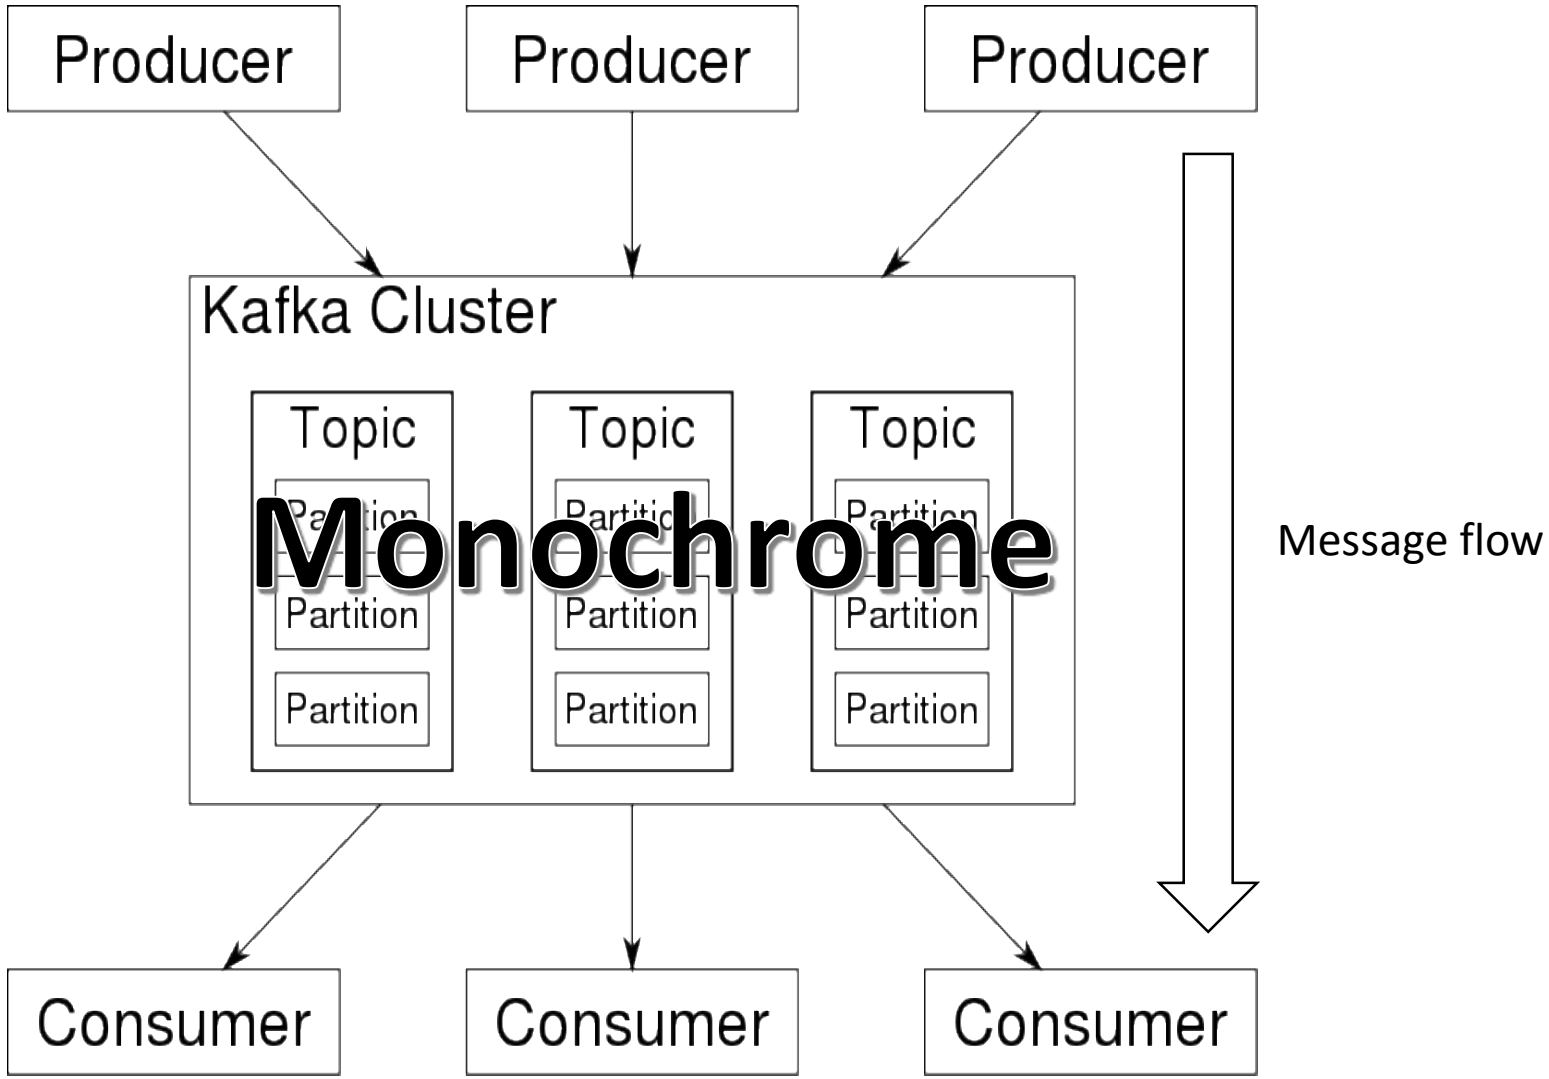

Monochrome

Kafka Postal Service

Let's build the

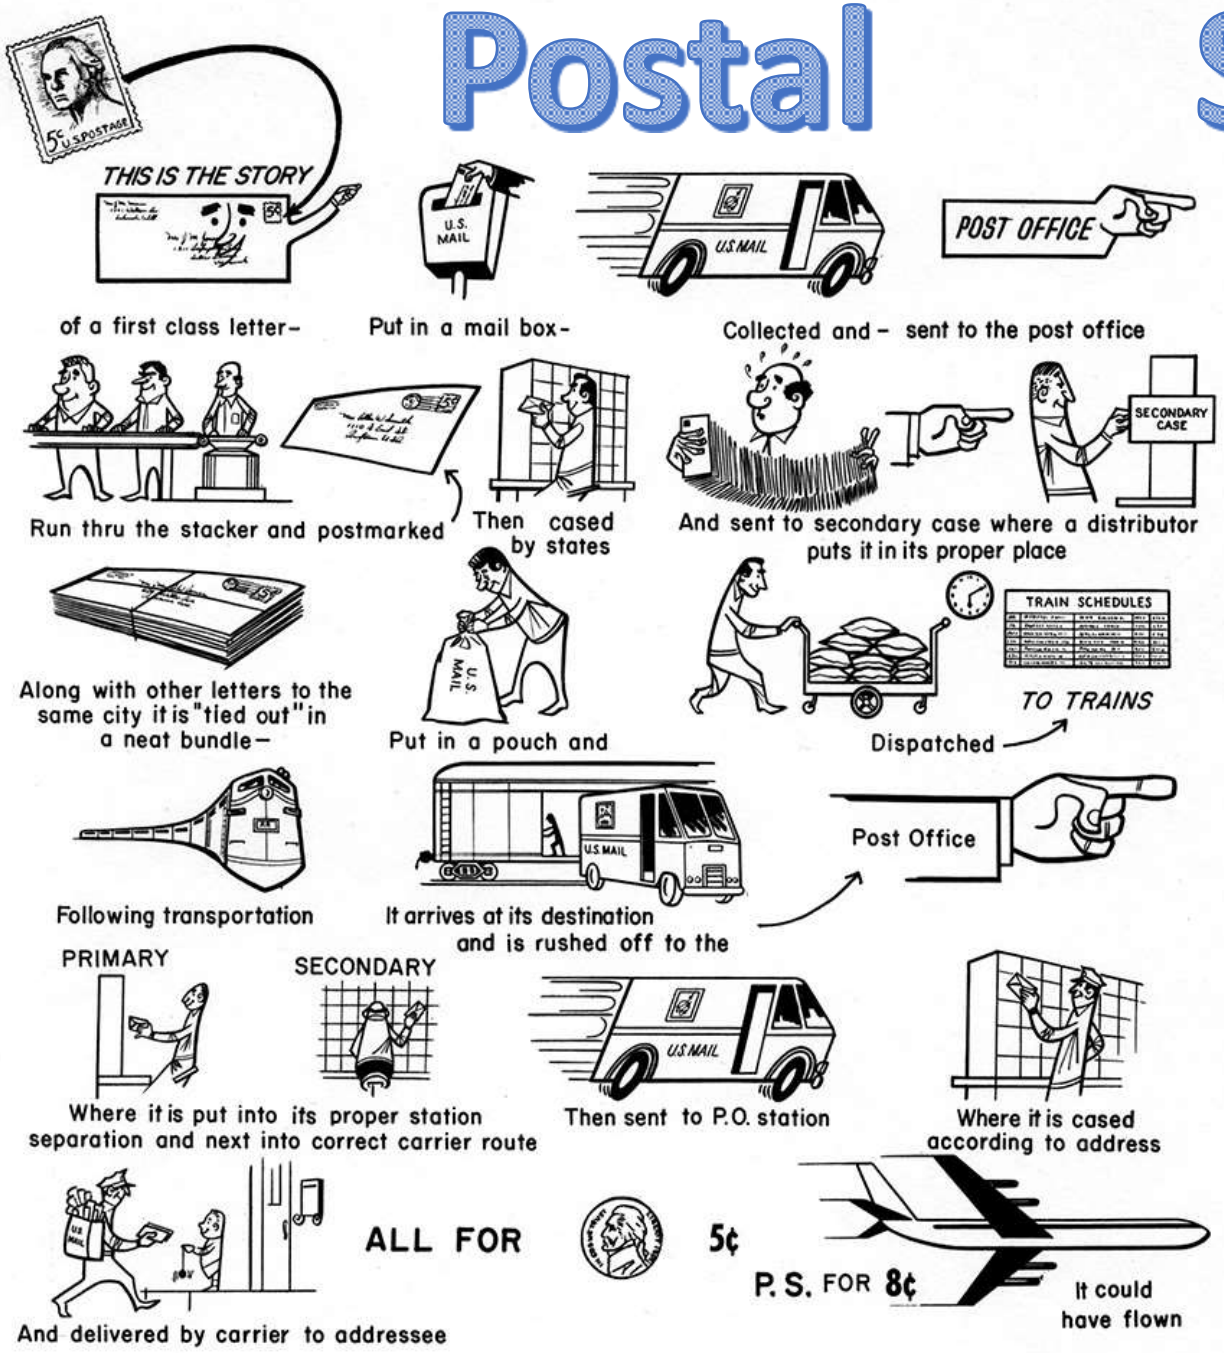

To send messages from A to B

A

B

“A” is the Producer – sends a message, to

“B” the Consumer (recipient) of the message

Due to decline in “snail mail” direct deliveries

Due to decline in “snail mail” direct deliveries

Instead … “Poste Restante”

“Poste Restante”?

- Not a post office in a restaurant
- General delivery (in the US)
- The mail is delivered to a post office, they hold it for you until you call

Consumers “poll” for messages by visiting the Poste Restante counter at the post office

Kafka topics act like a Post Office

Benefits include

- Disconnected delivery – consumer doesn't need to be available to receive messages
- Less effort for the messaging service – only has to deliver to a few locations not every consumer
- Can scale better and handle more complex delivery semantics!

Scalability? Many consumers for a topic?

A single counter introduces delays

AUSTRALIAN WAR MEMORIAL P00784.042

More counters increases concurrency

Kafka topics have >= 1 Partitions (“counters”)

- Partitions increase consumer concurrency
- Increase throughput
- Reduce latency

What's a Kafka message?

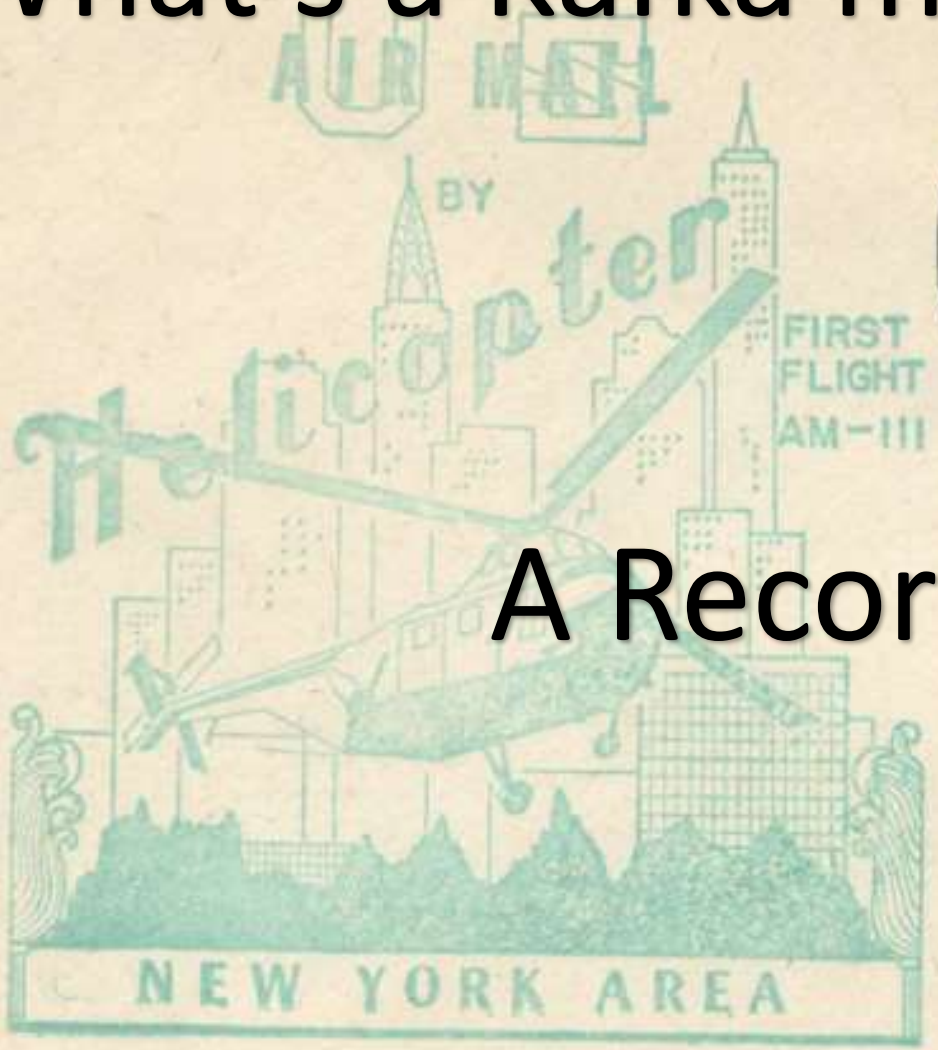

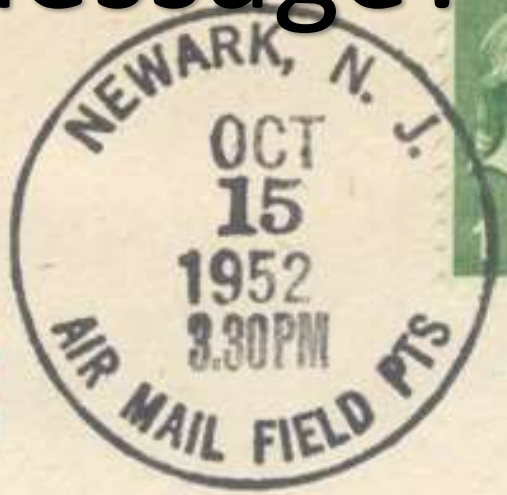

A Record – like a letter

Santa

North Pole

Topic is the destination

Santa

Topic → North Pole

The “Postmark”

Timestamp, offset, partition

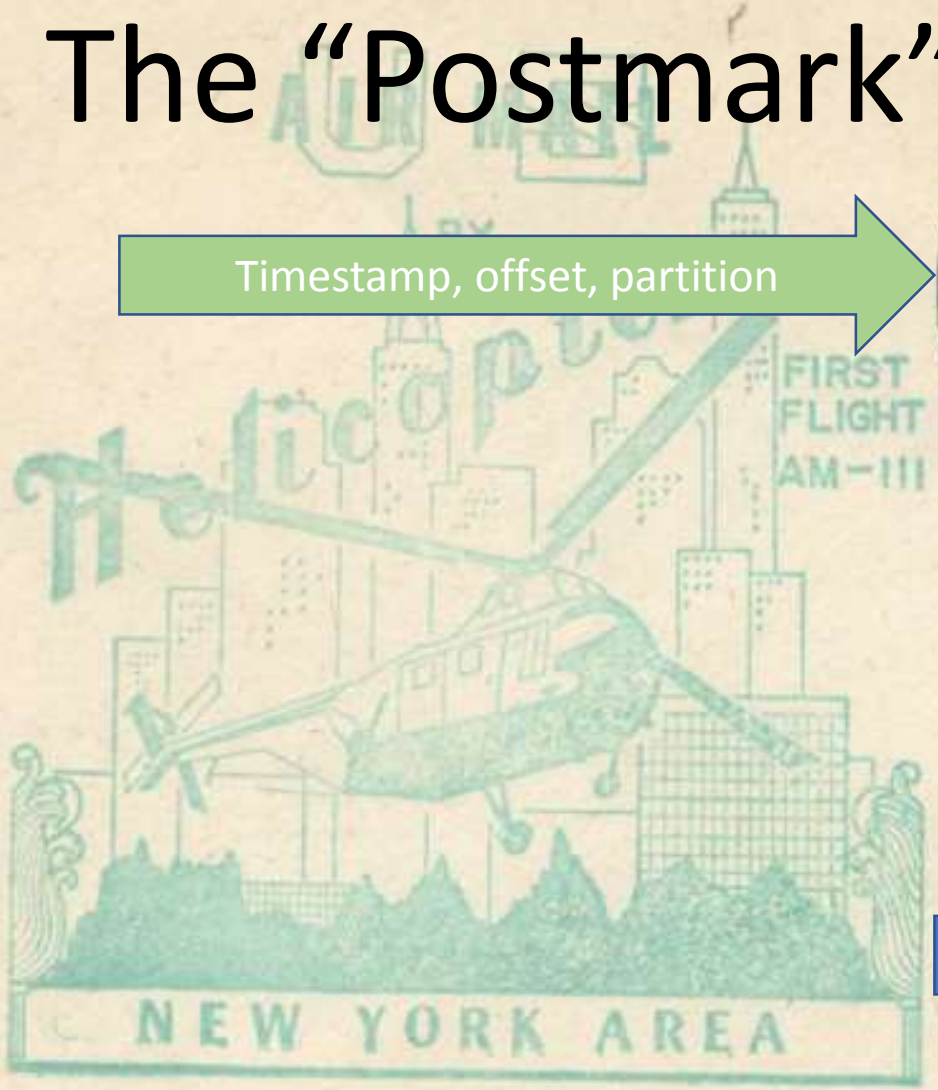

Topic

Santa

North Pole

The “Postmark”

Timestamp, offset, partition

Time semantics are flexible

time of event creation,

ingestion,

or processing.

Topic

Santa

North Pole

Key (optional)

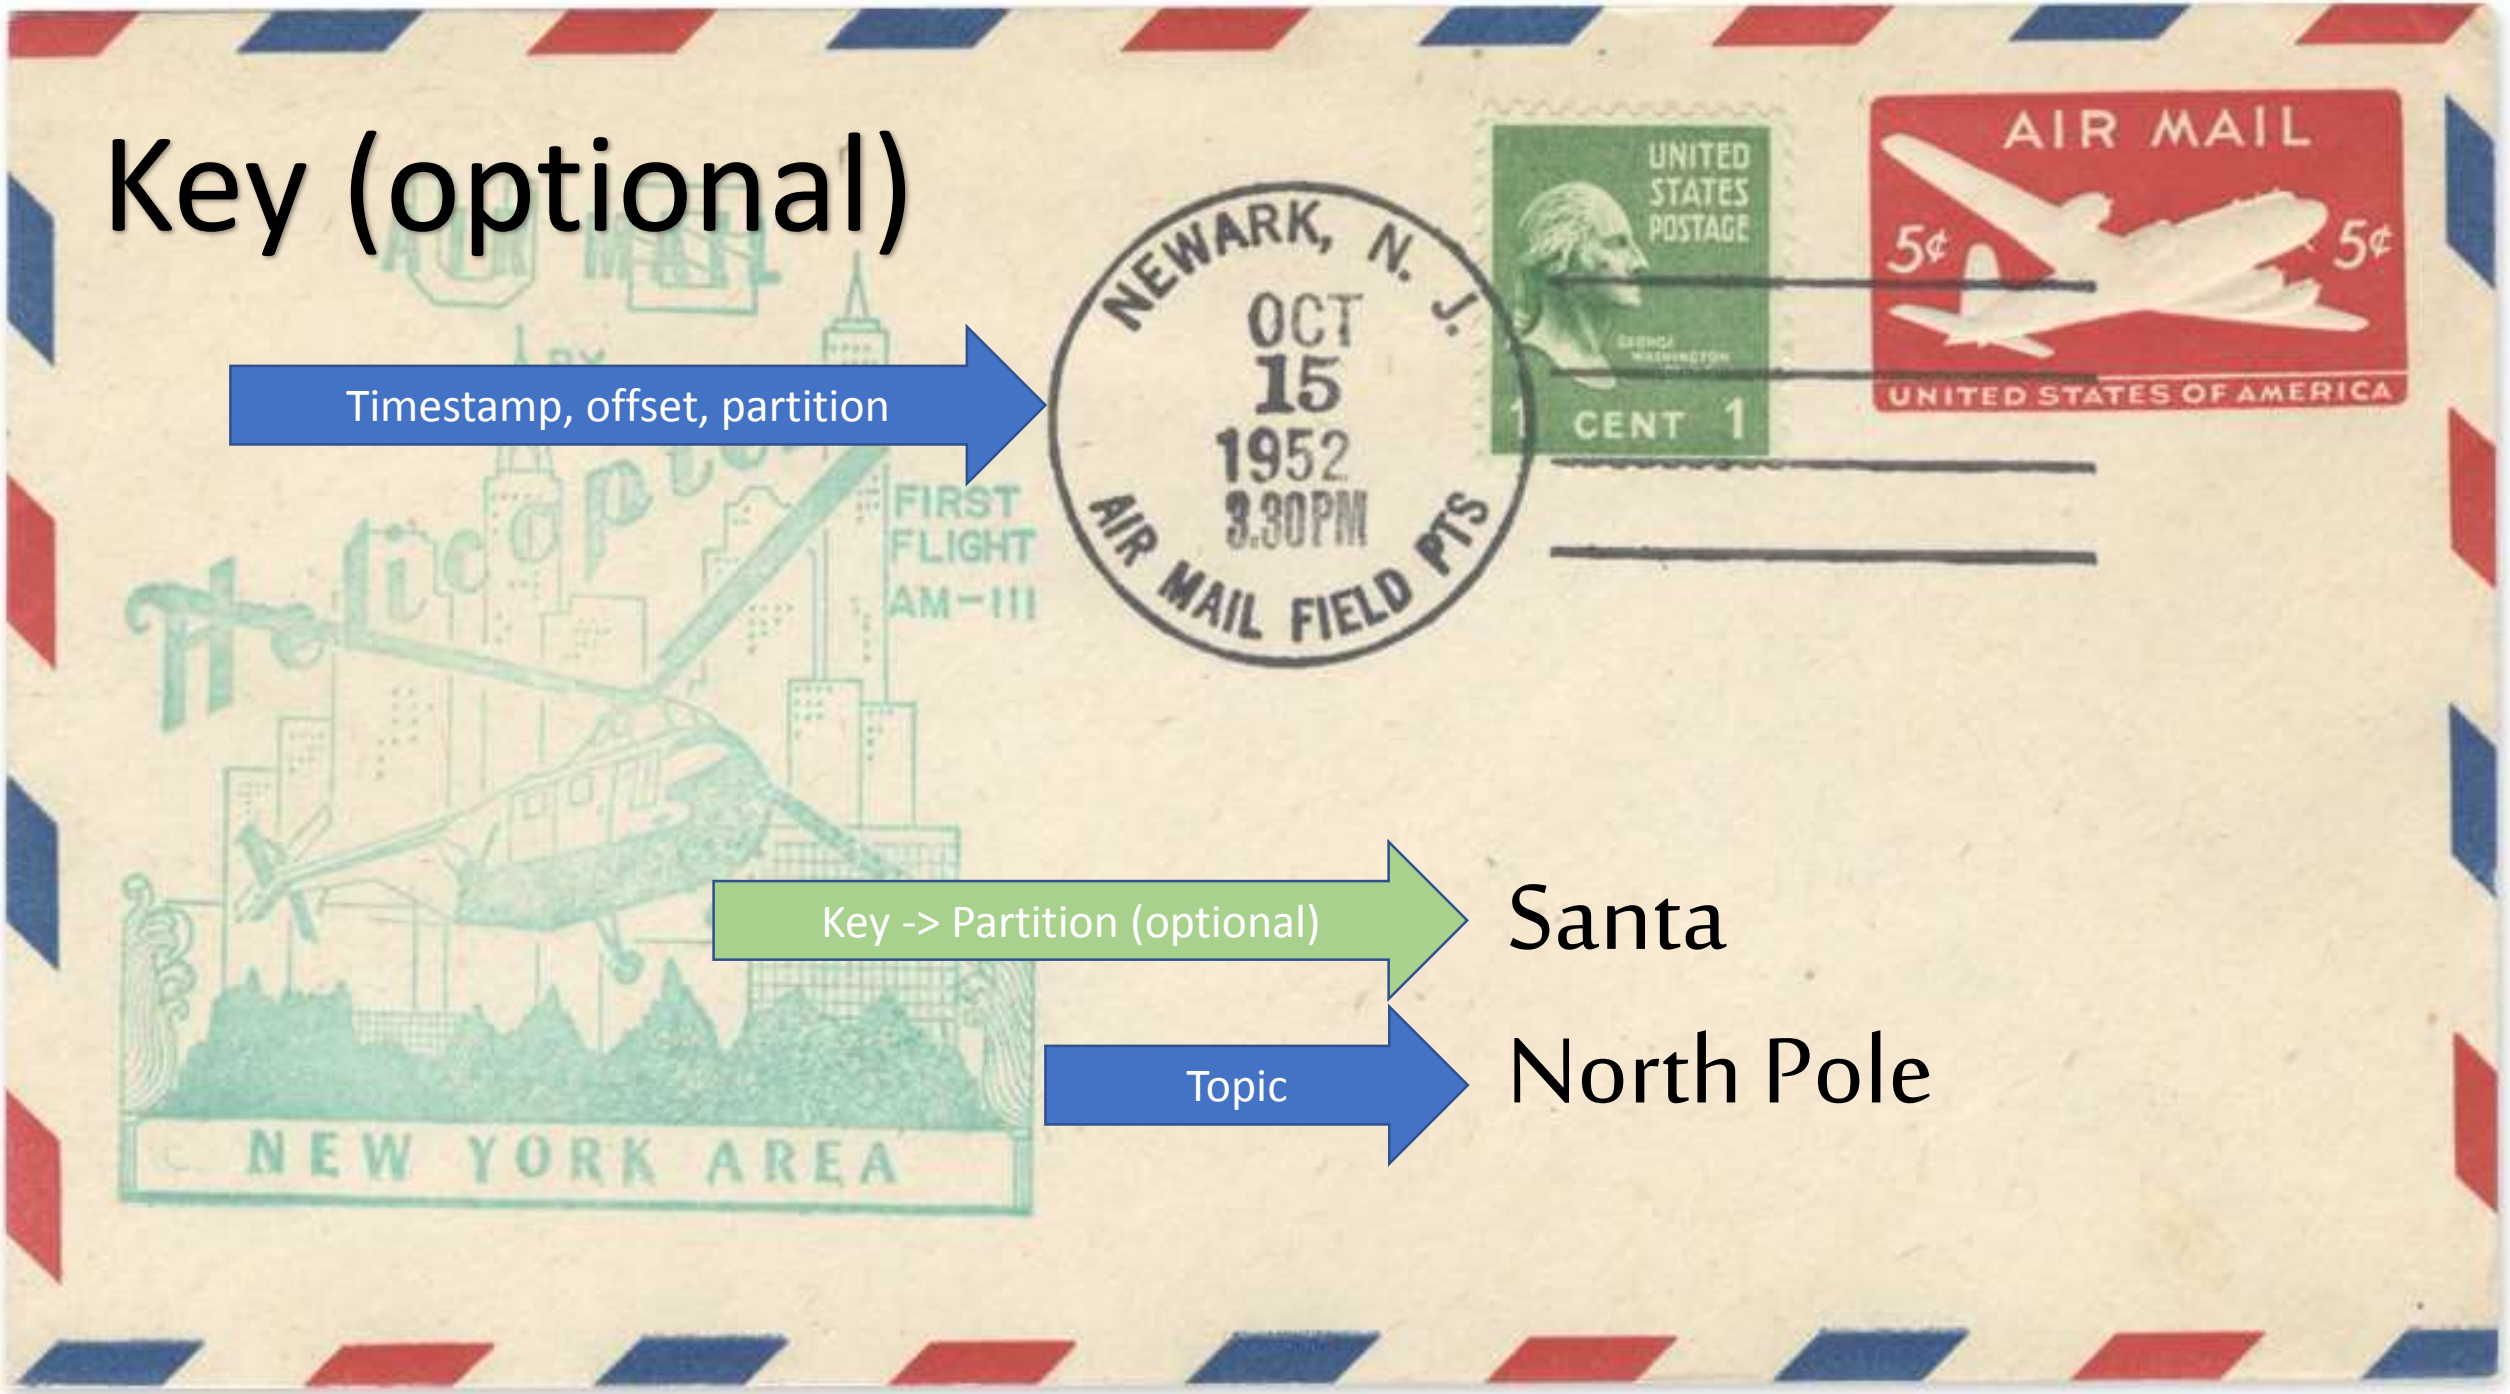

Timestamp, offset, partition

Key -> Partition (optional)

Santa

Topic

North Pole

Key (optional)

Timestamp, offset, partition

Refines the destination

Send to Santa not just any Elf

Key -> Partition (optional)

Santa

Topic

North Pole

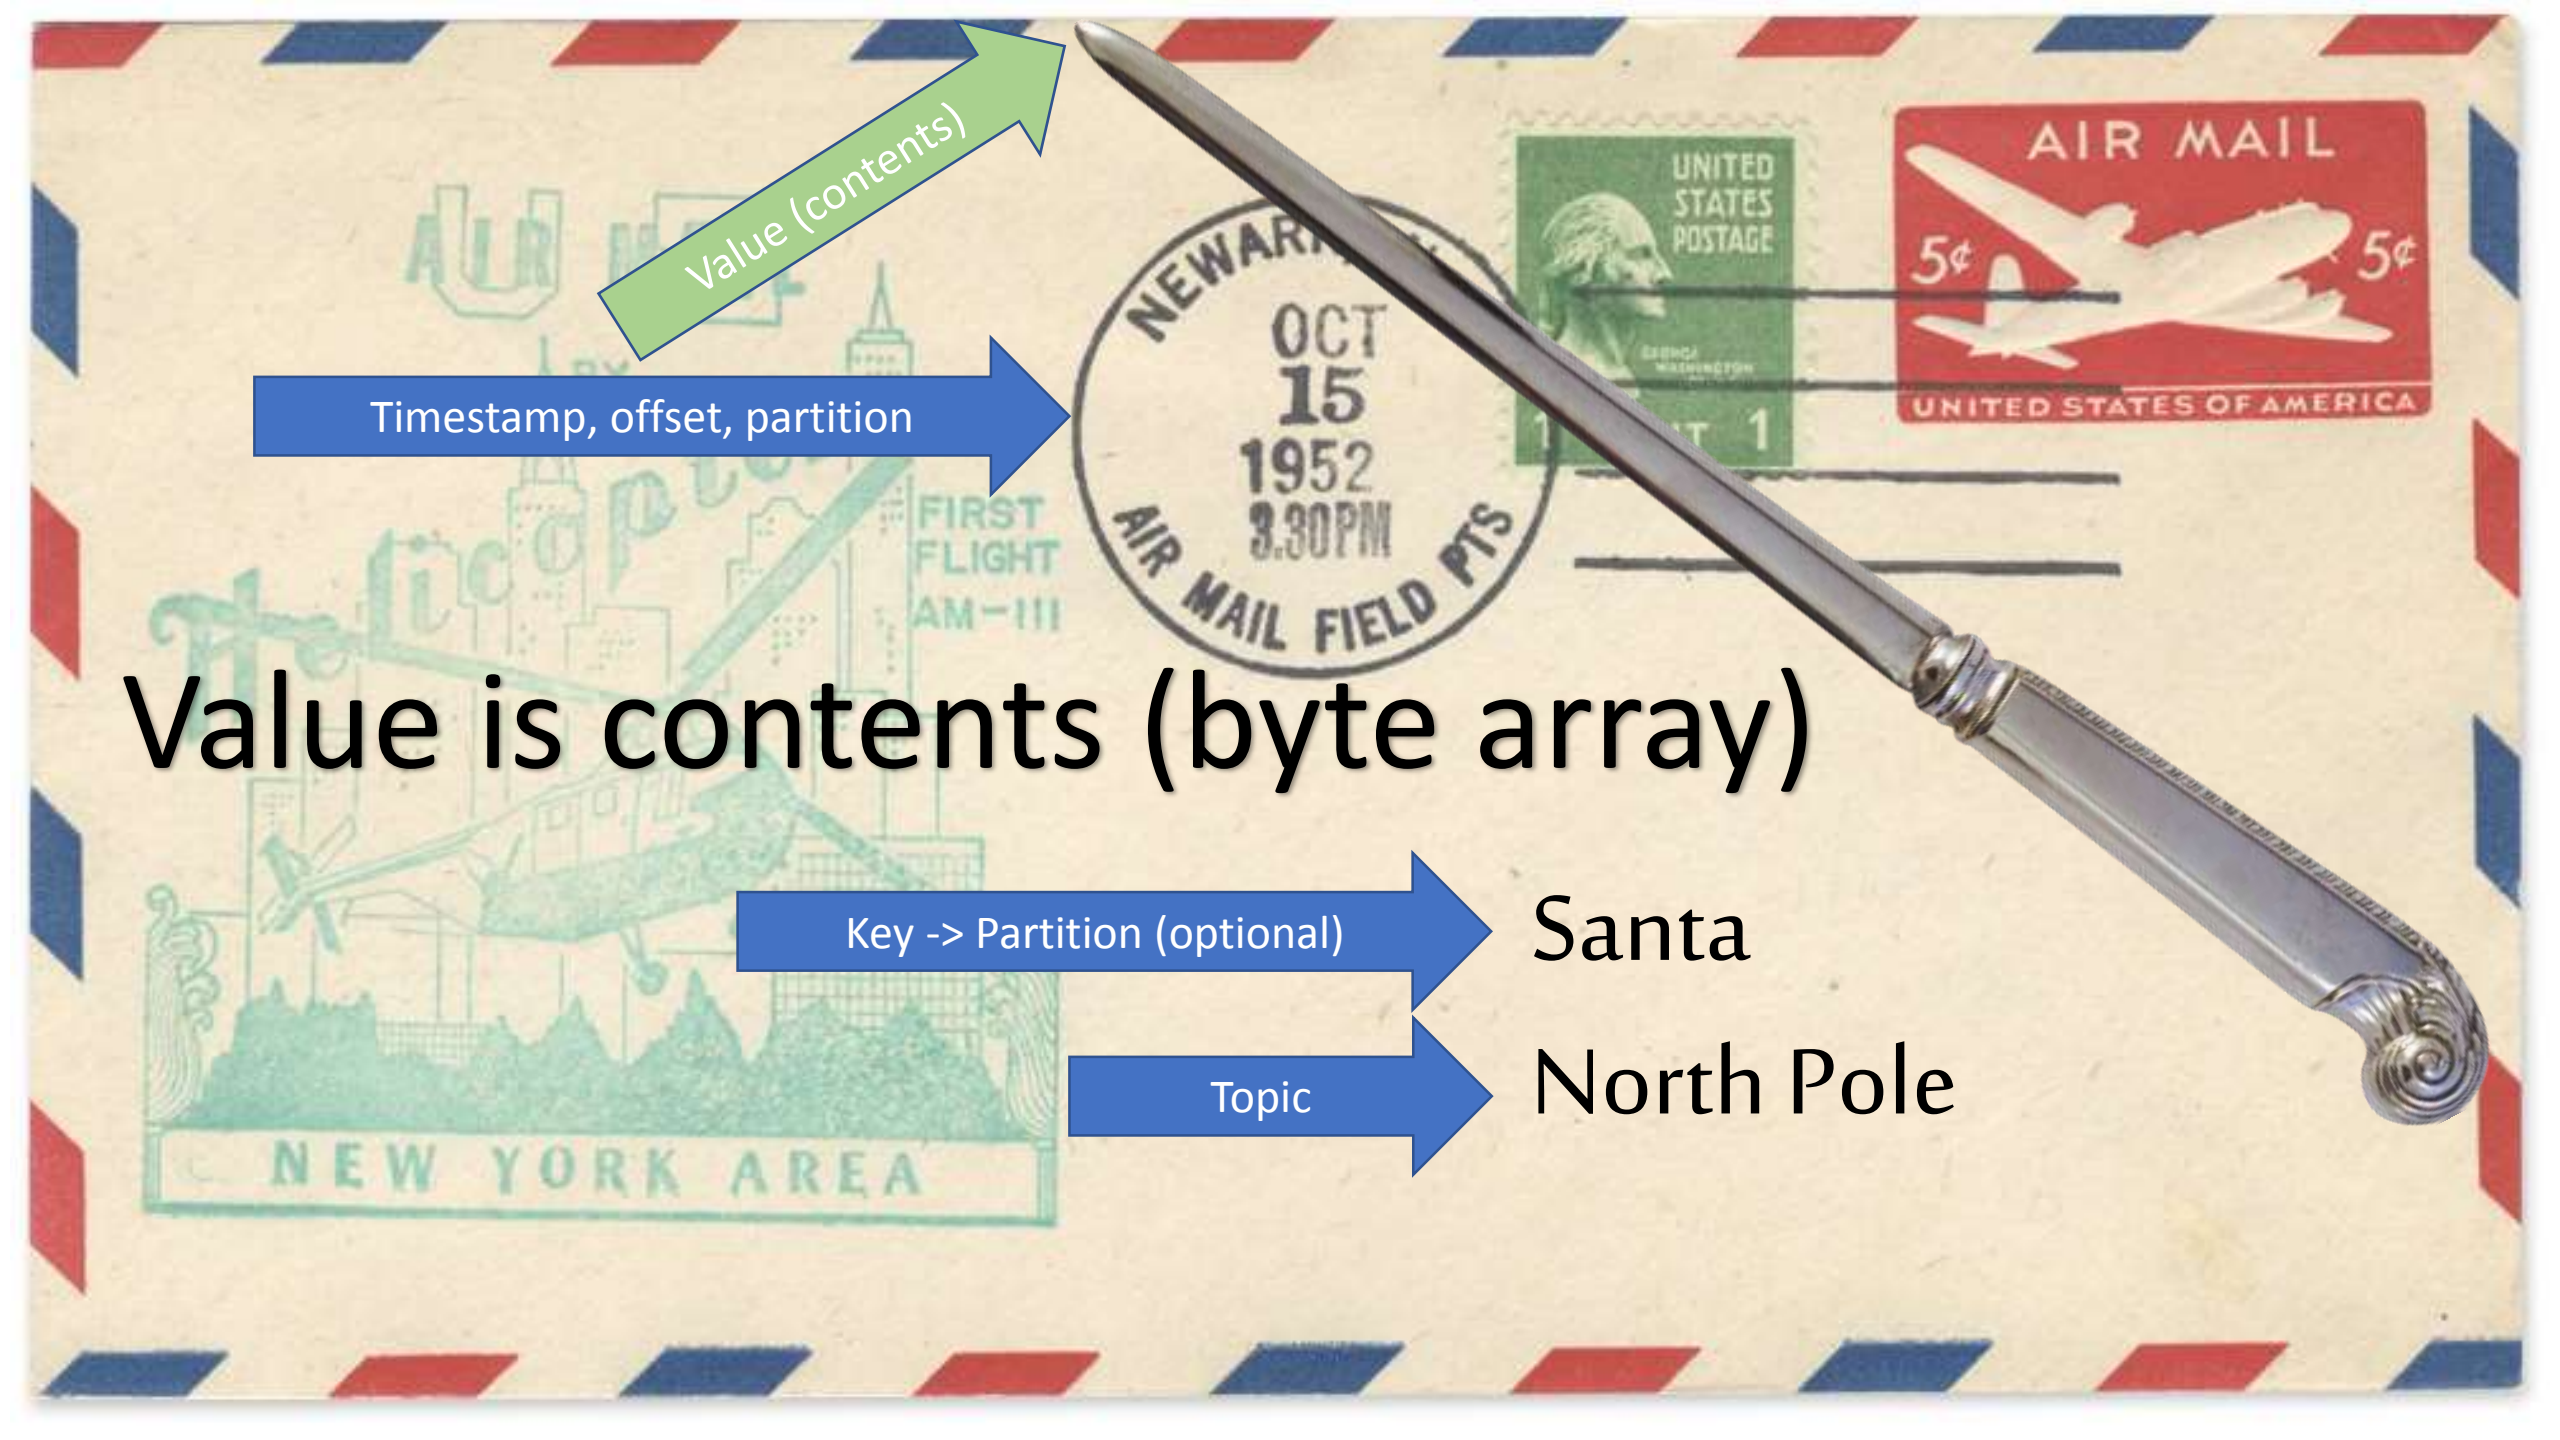
Value (contents)
Timestamp, offset, partition
Value is contents (byte array)
Key -> Partition (optional)
Santa
Topic
North Pole
AIR MAIL
5¢
5¢
UNITED STATES OF AMERICA
UNITED STATES POSTAGE
OCT
15
1952
3.30PM
AIR MAIL FIELD PTS
FIRST
FLIGHT
AM-111
NEW YORK AREA

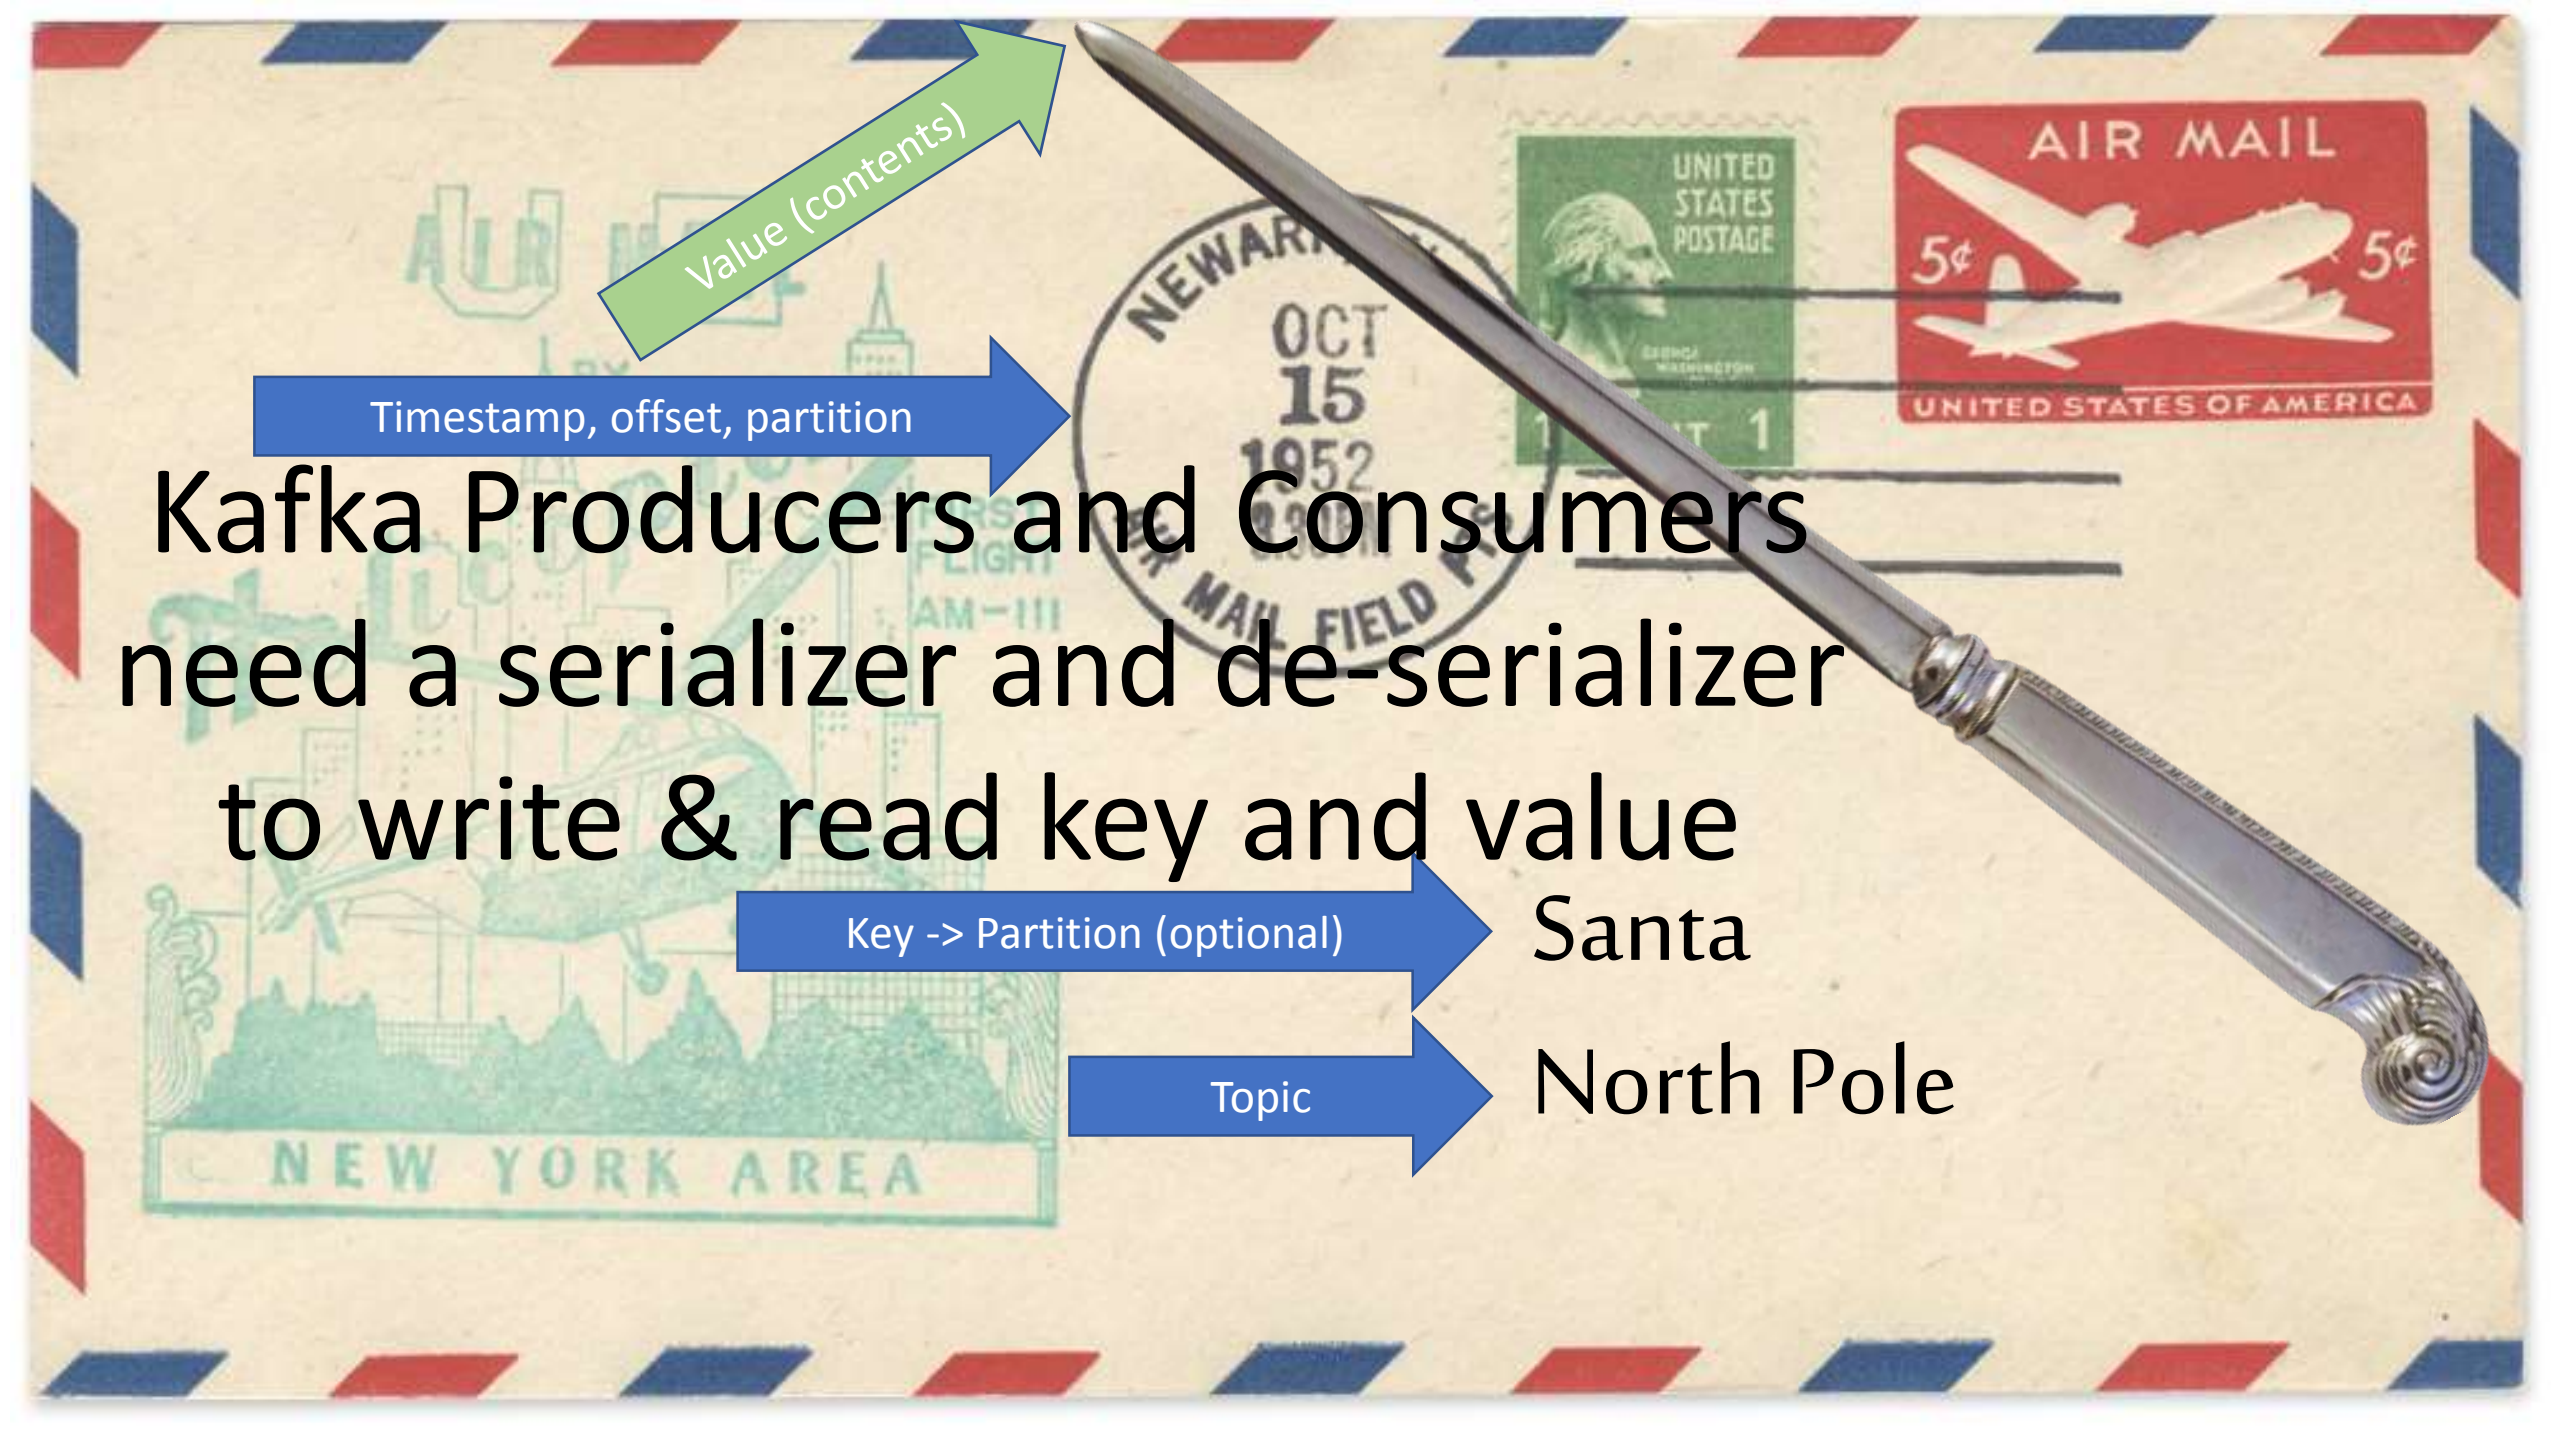
Value (contents)
Timestamp, offset, partition
Kafka Producers and Consumers need a serializer and de-serializer to write & read key and value
Key -> Partition (optional)
Santa
Topic
North Pole
AIR MAIL
UNITED STATES OF AMERICA
UNITED STATES POSTAGE
NEW YORK AREA

- Kafka doesn’t look at the value
- Consumer can read value
- And try to make sense of the message
- What will Santa be delivering?!

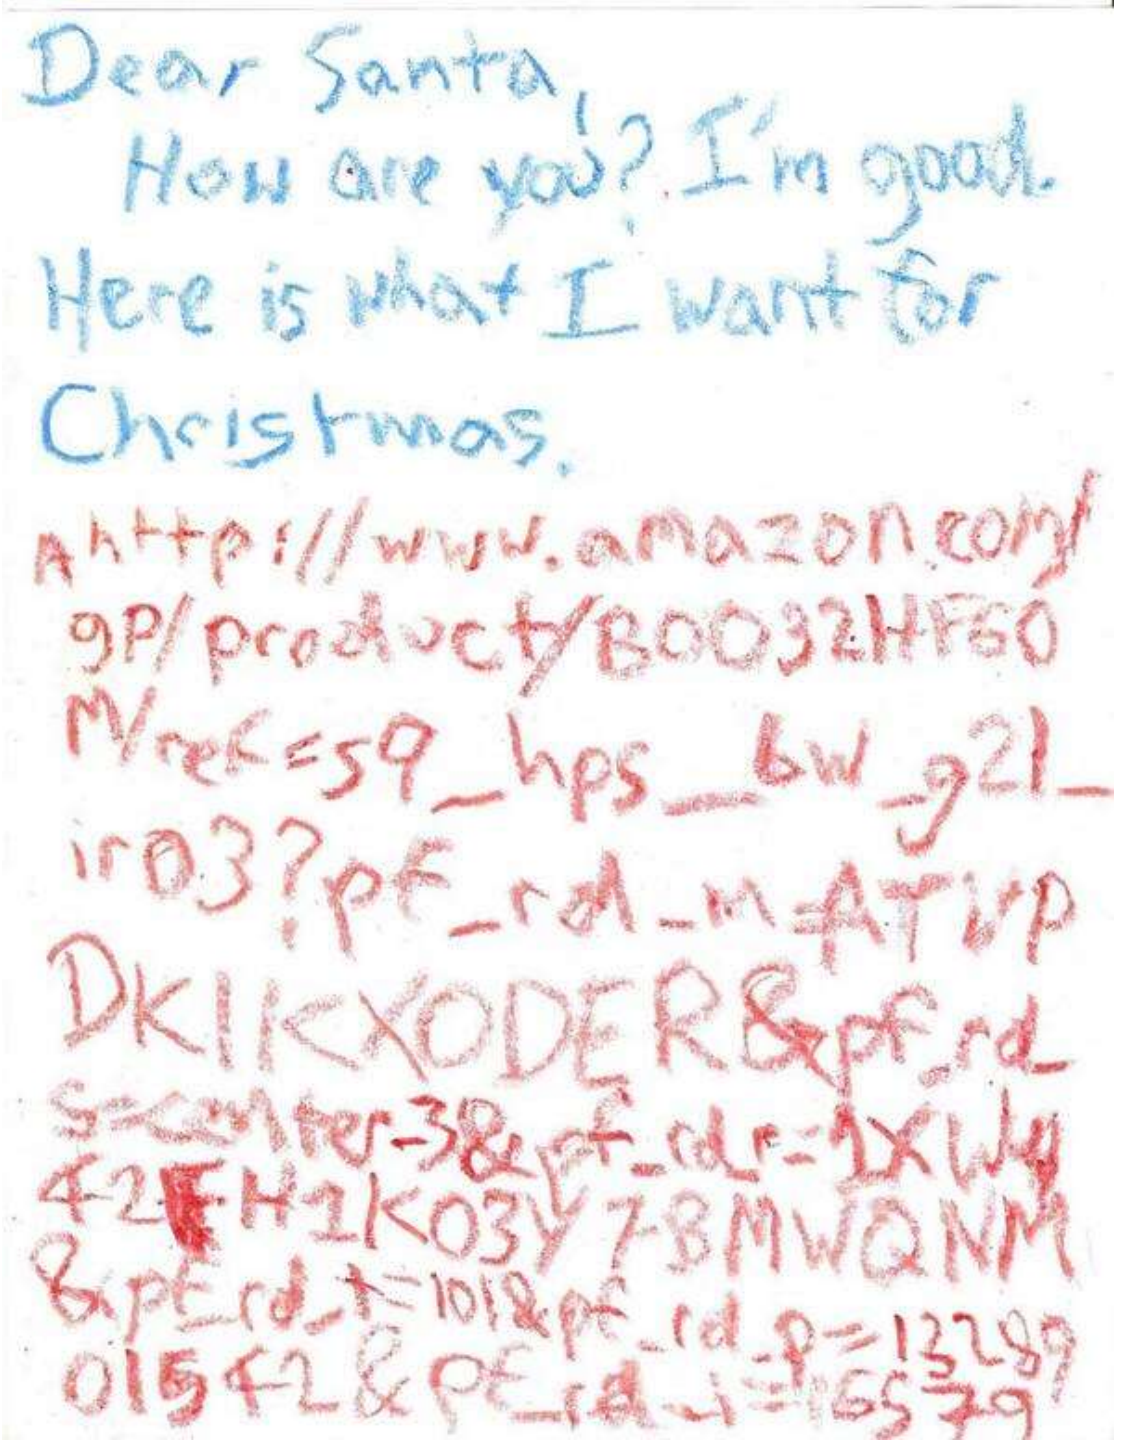

Next…
Delivery Semantics

Do we care
if the message arrives?

Yes! Guaranteed delivery is desirable

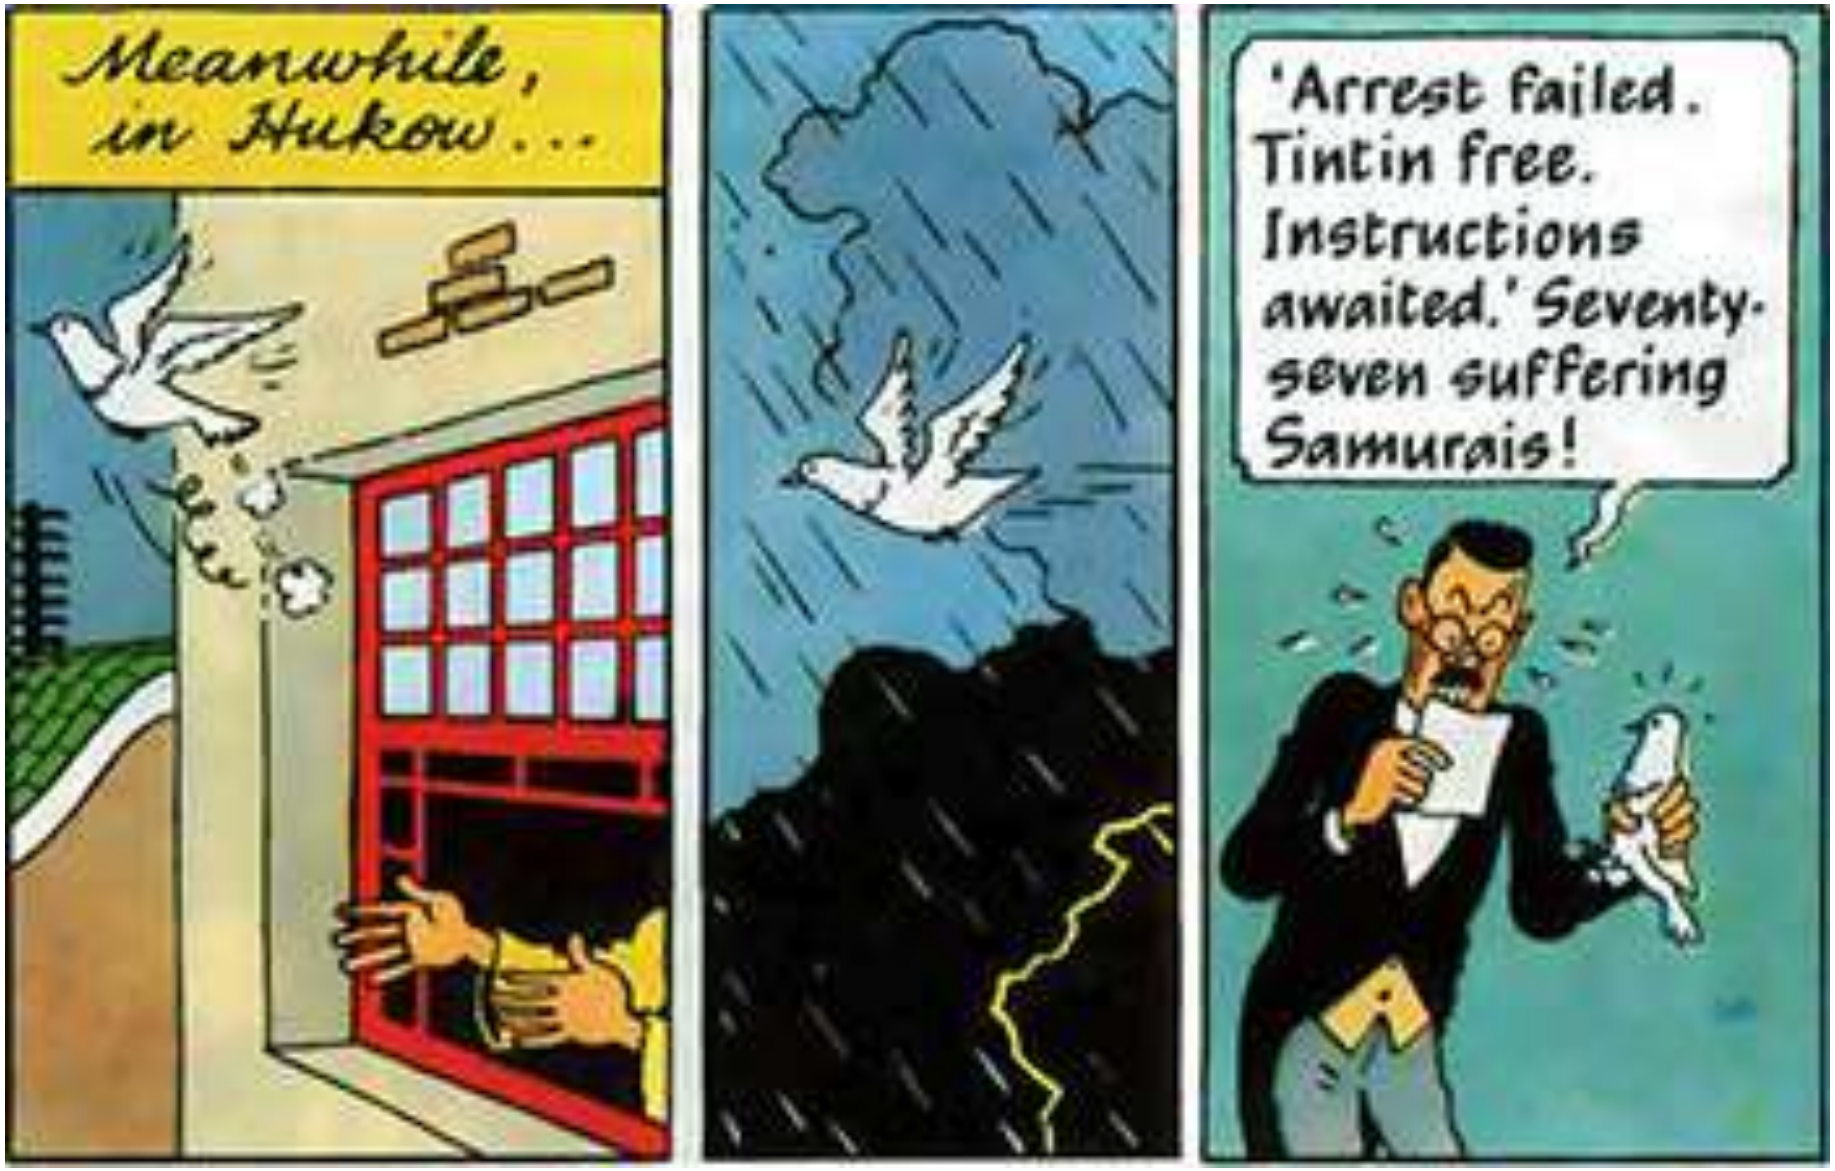

But homing pigeons got lost or eaten

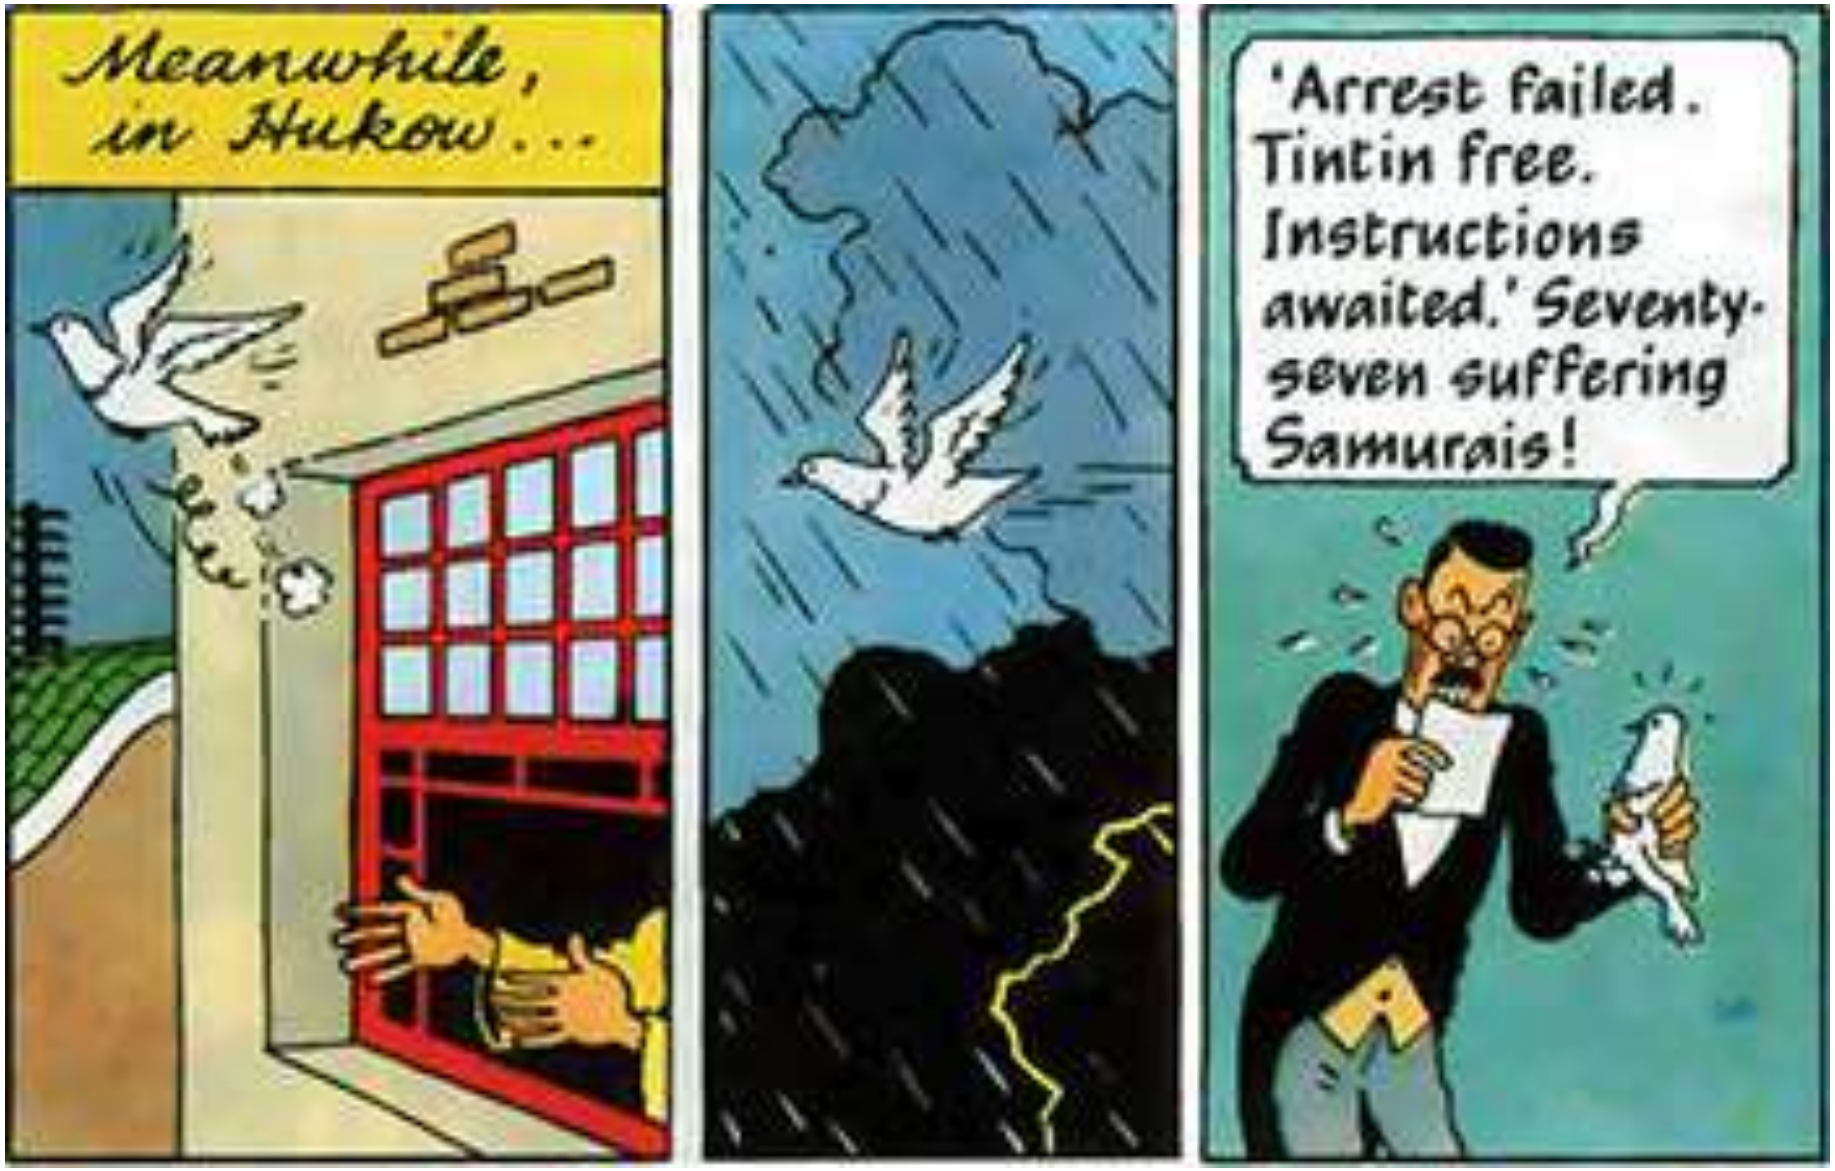

Send multiple pigeons

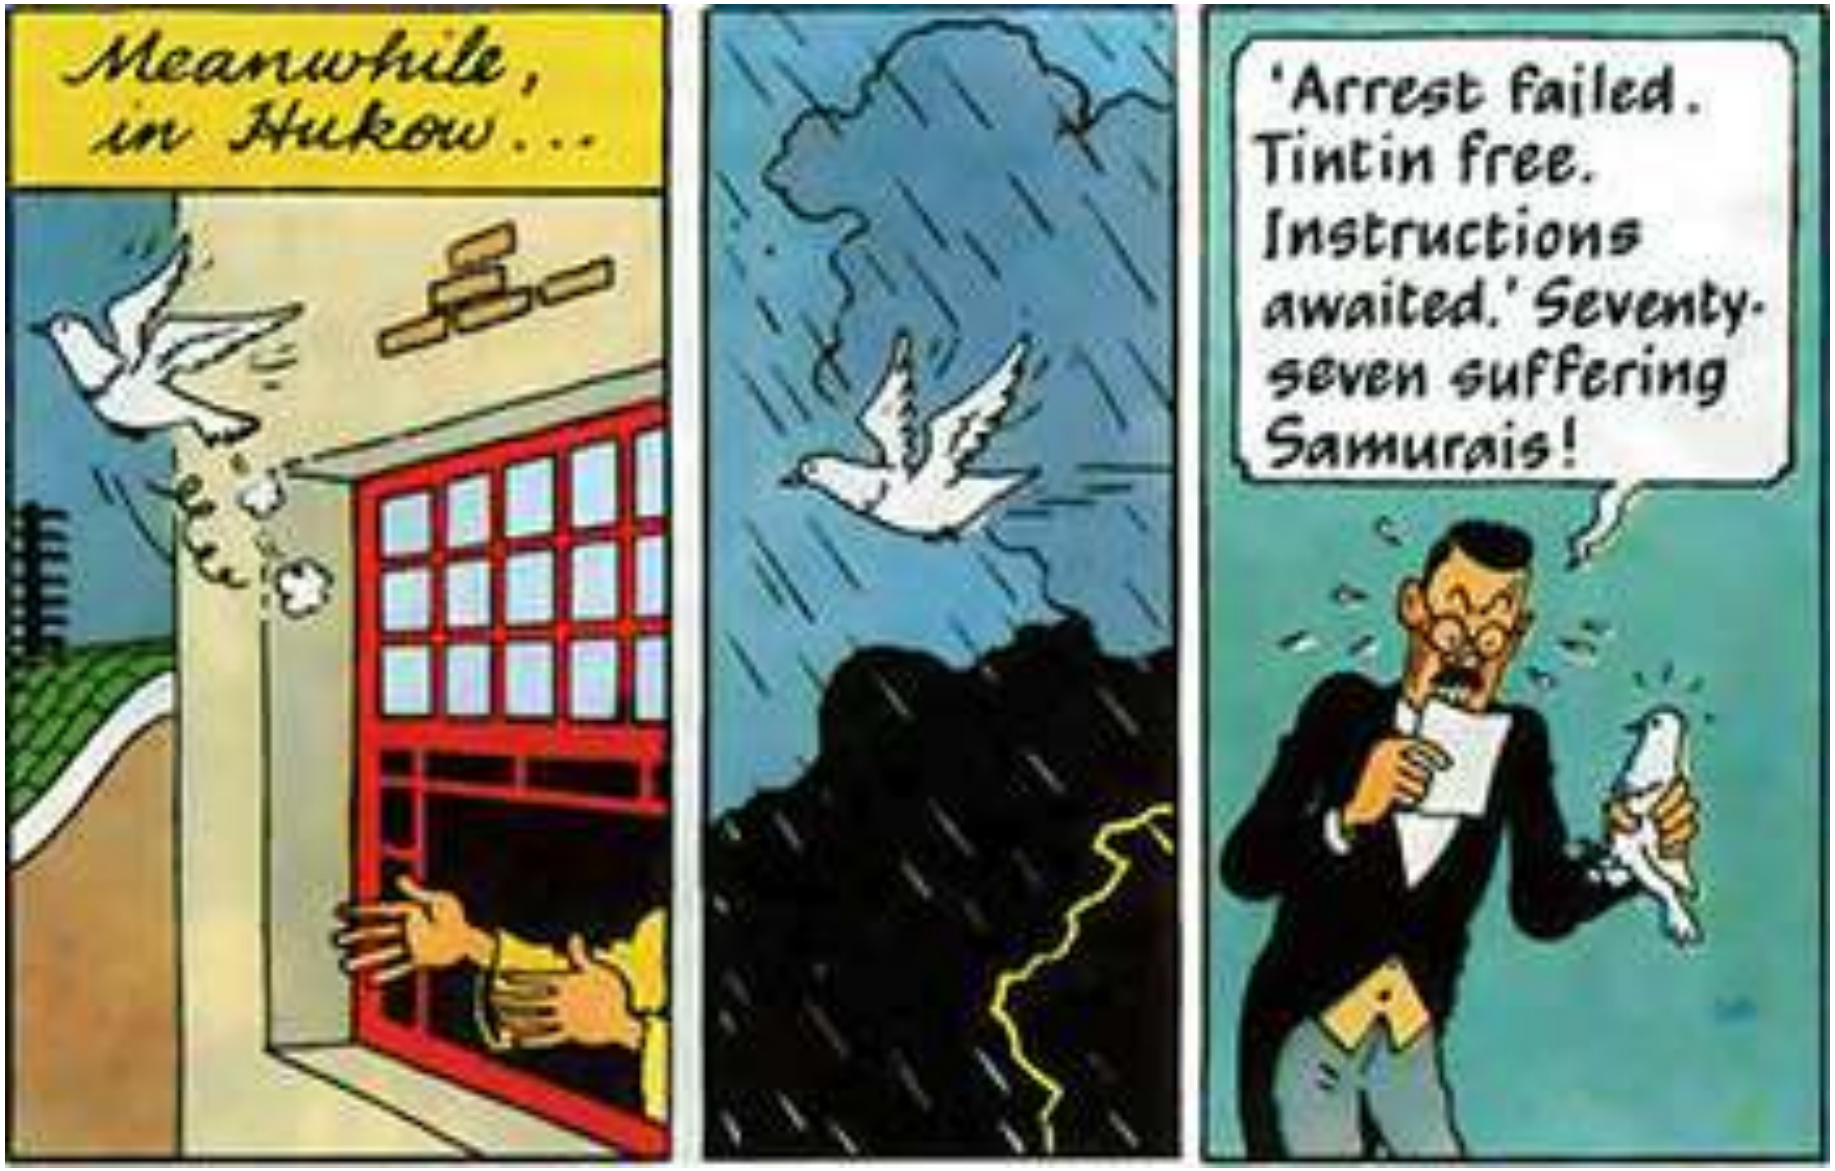

One pigeon may make it

Lost

Eaten

How does Kafka guarantee delivery?

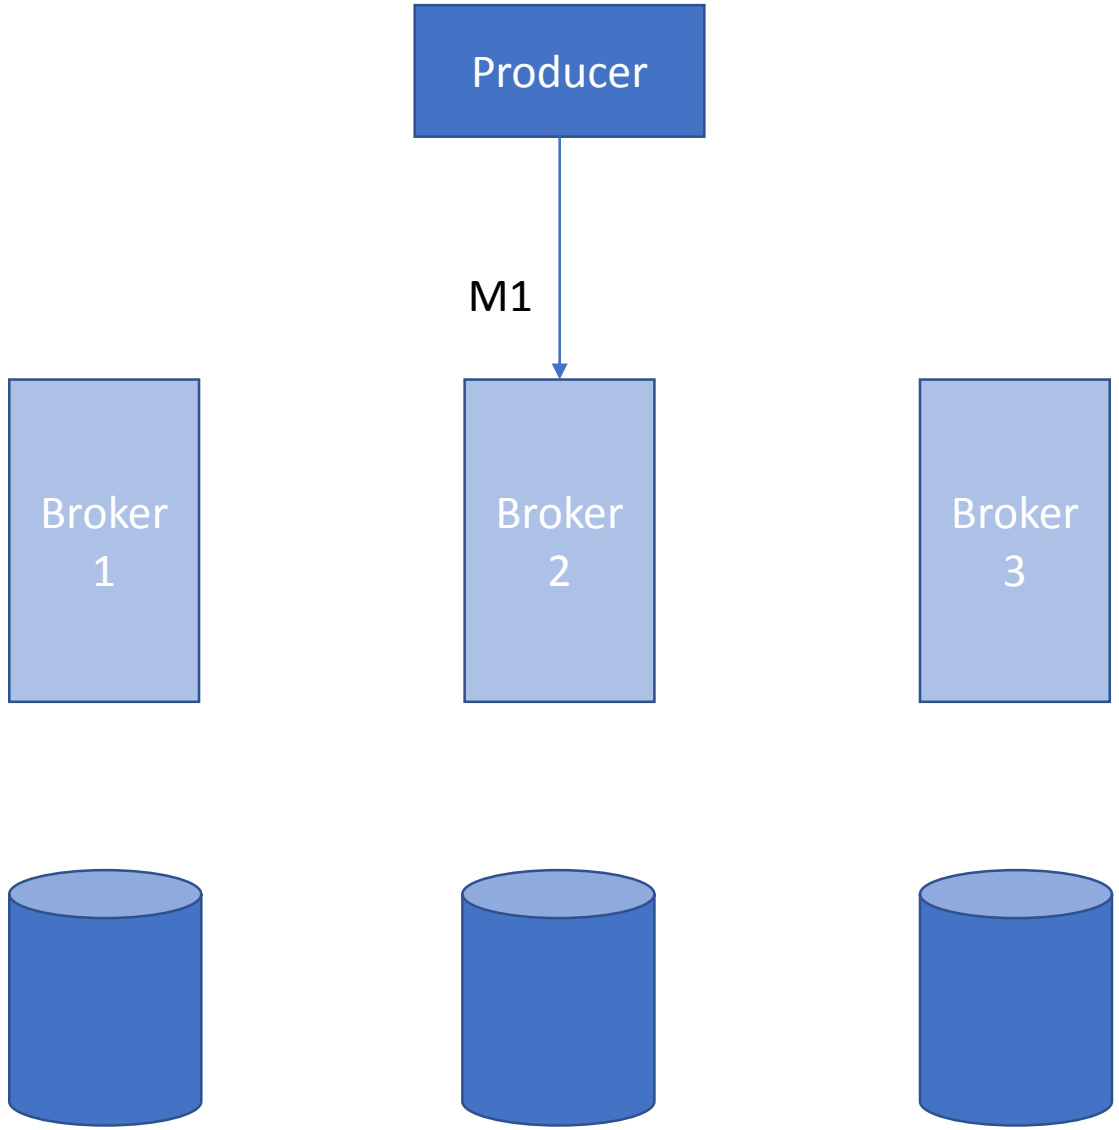

A Message (M1) is written to a broker (2)

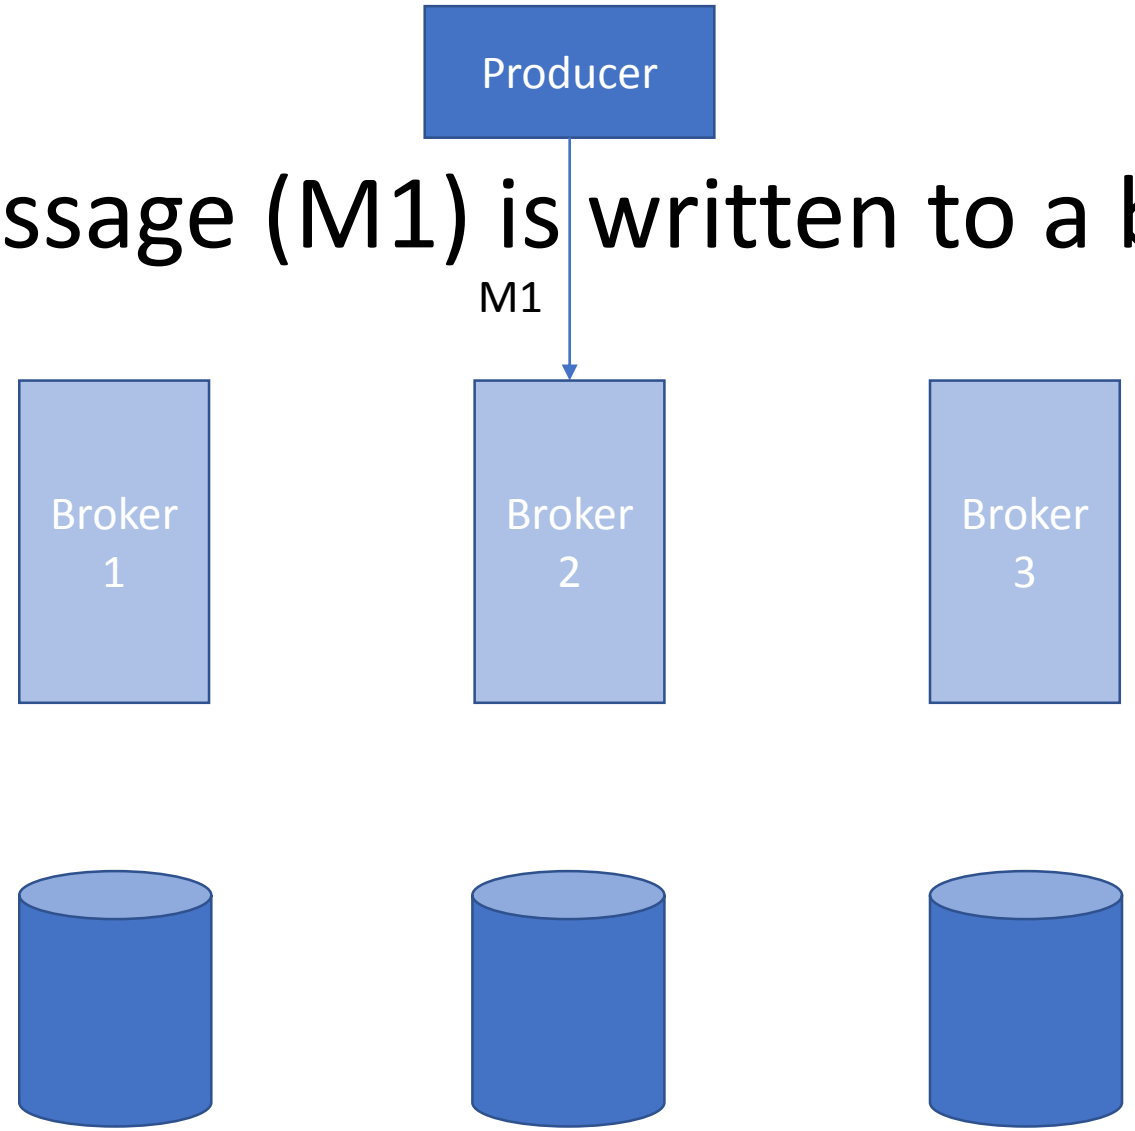

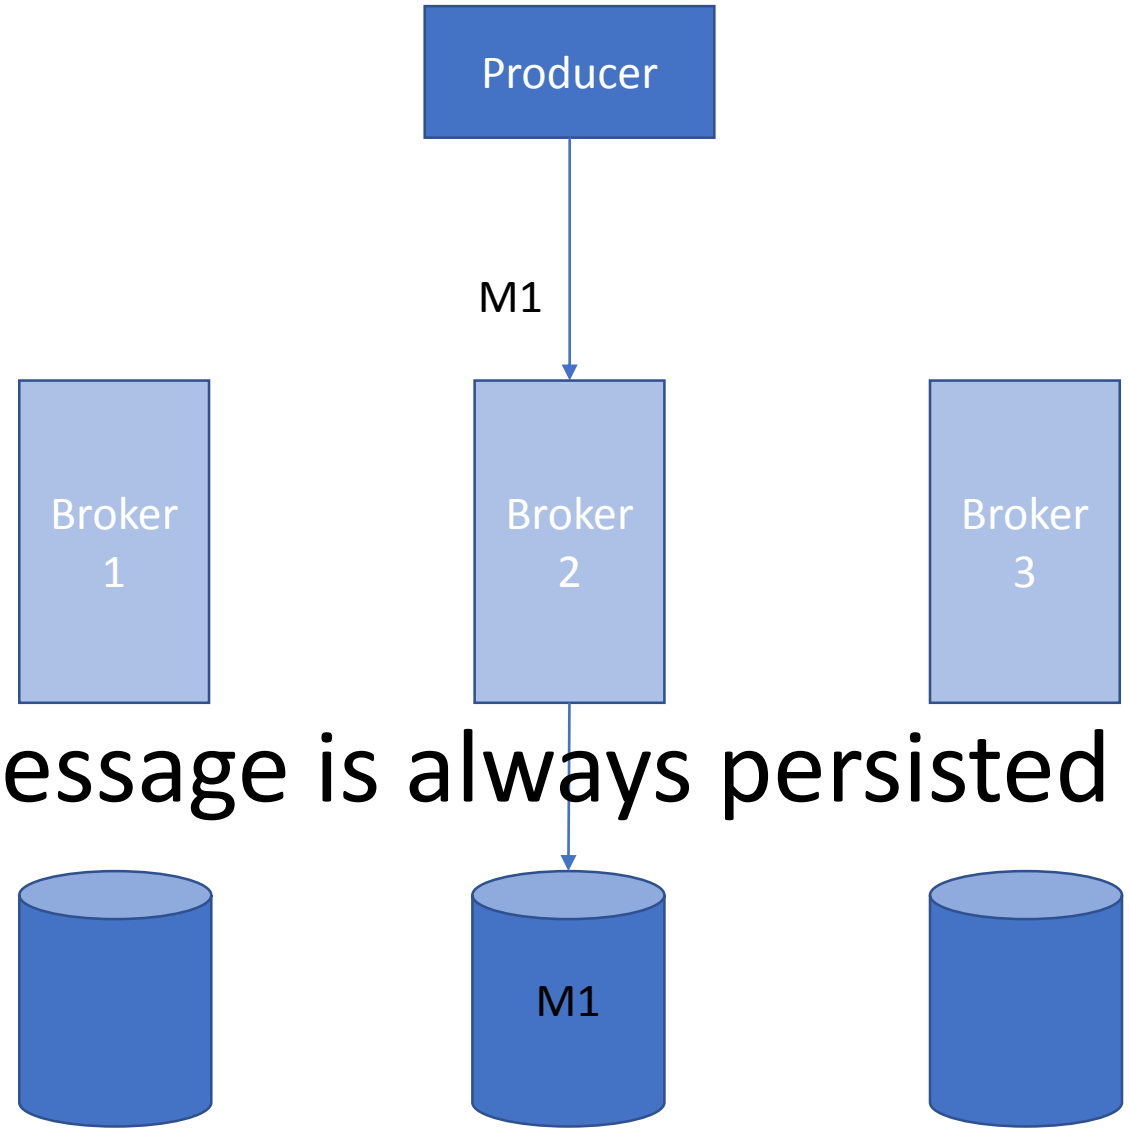

The message is always persisted to disk.

This makes it resilient to loss of power.

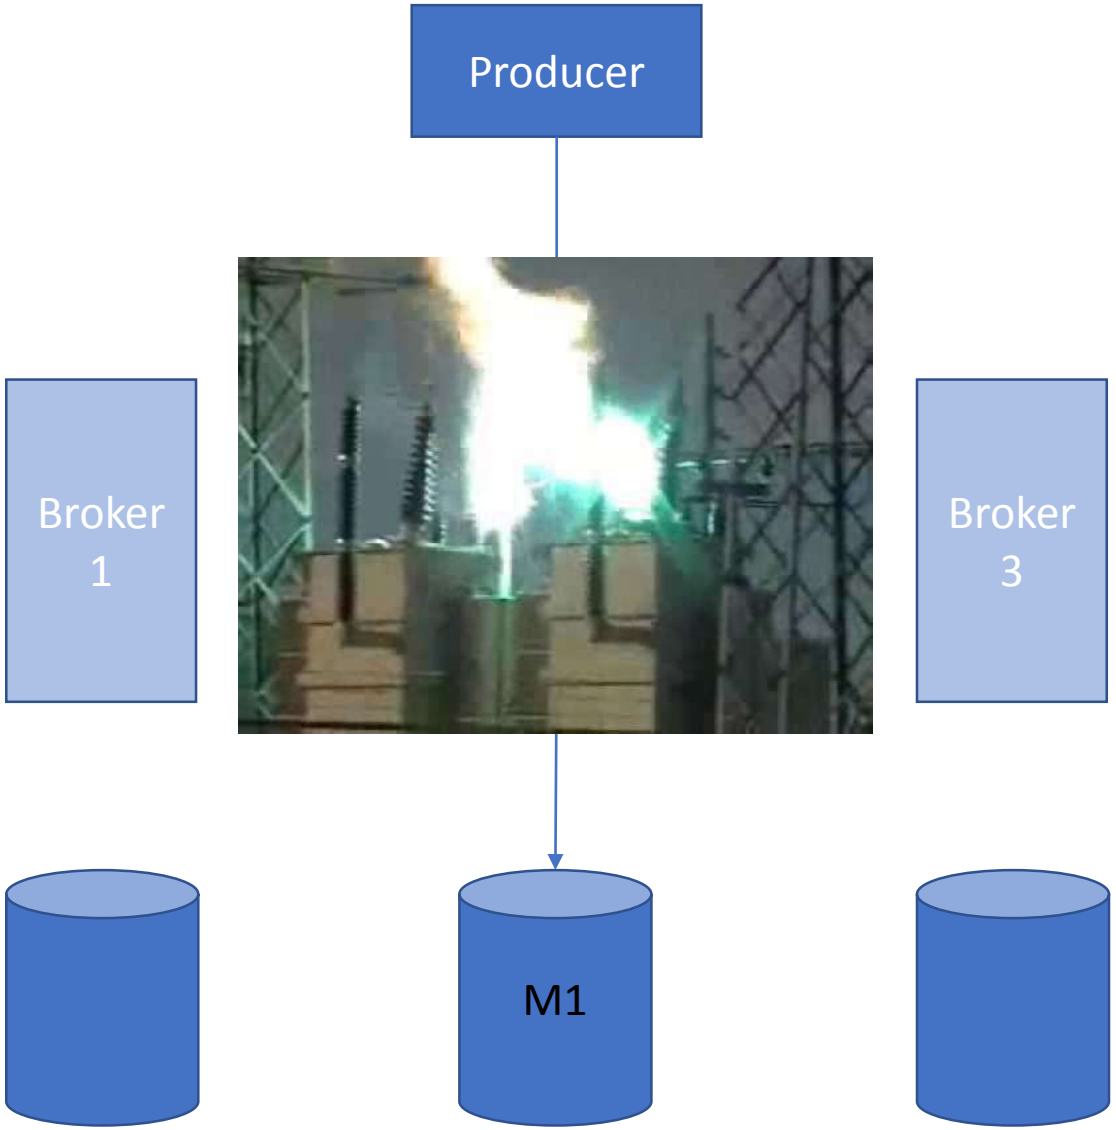

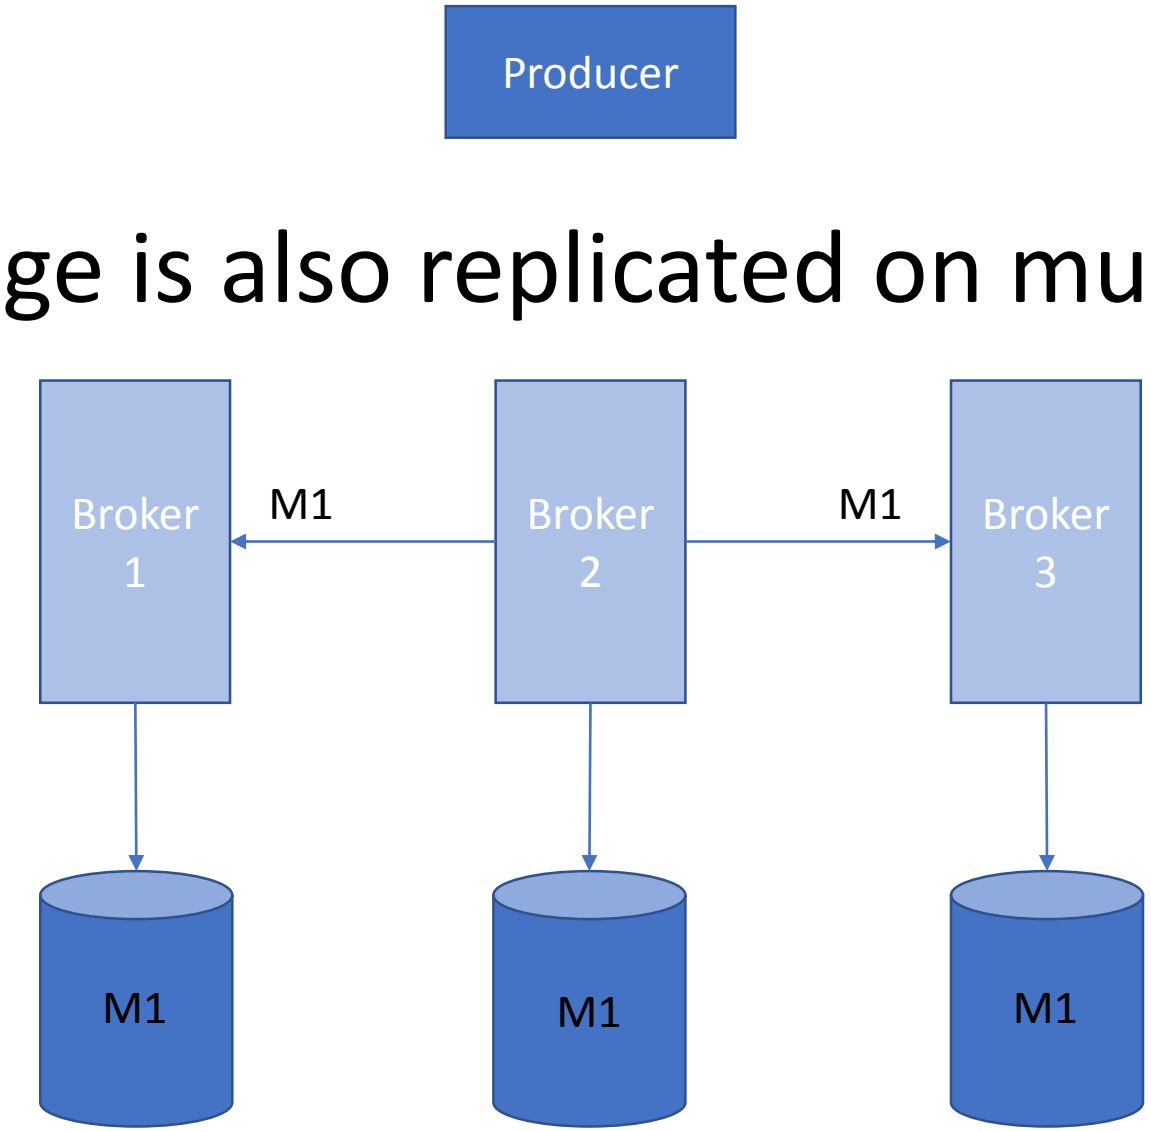

The message is also replicated on multiple “brokers”

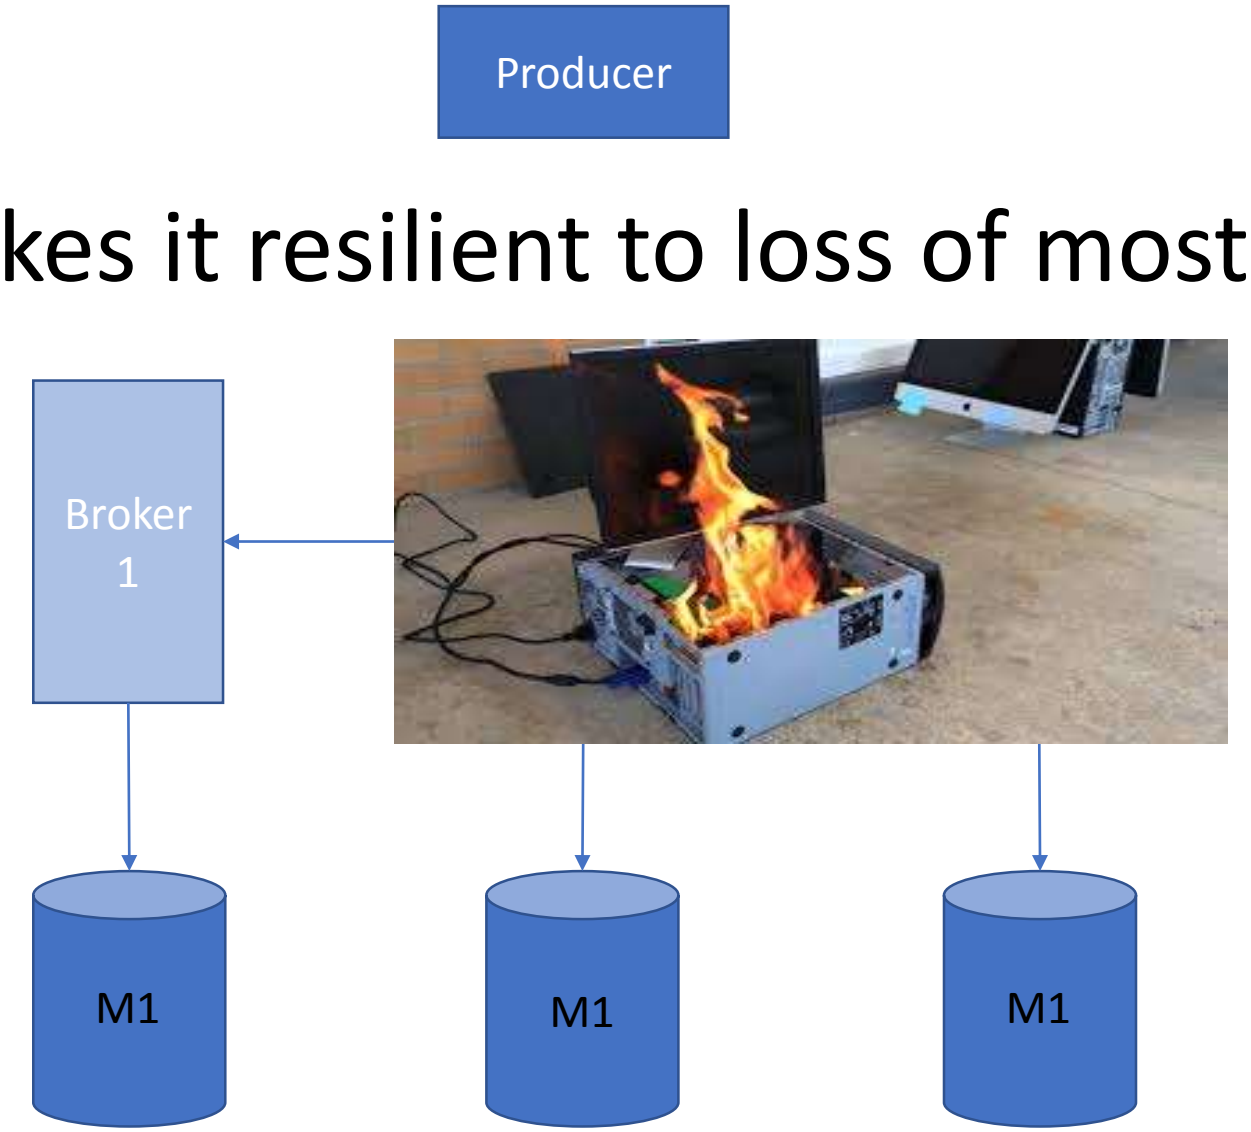

Which makes it resilient to loss of most servers

Producer gets acknowledgement once the message is persisted and replicated (configurable)

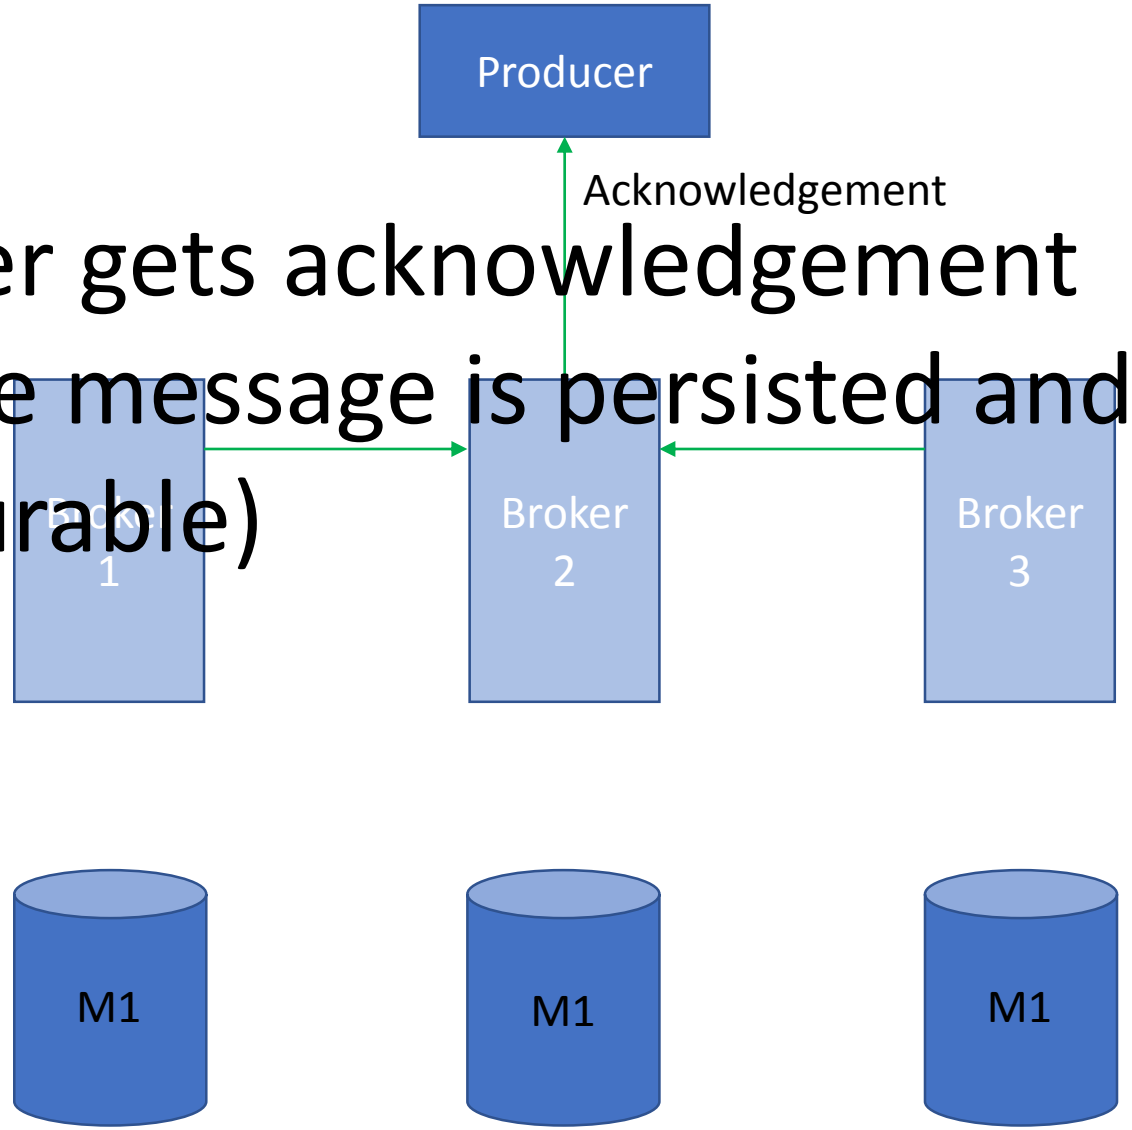

Multiple Brokers and Partitions → increased read availability and concurrency

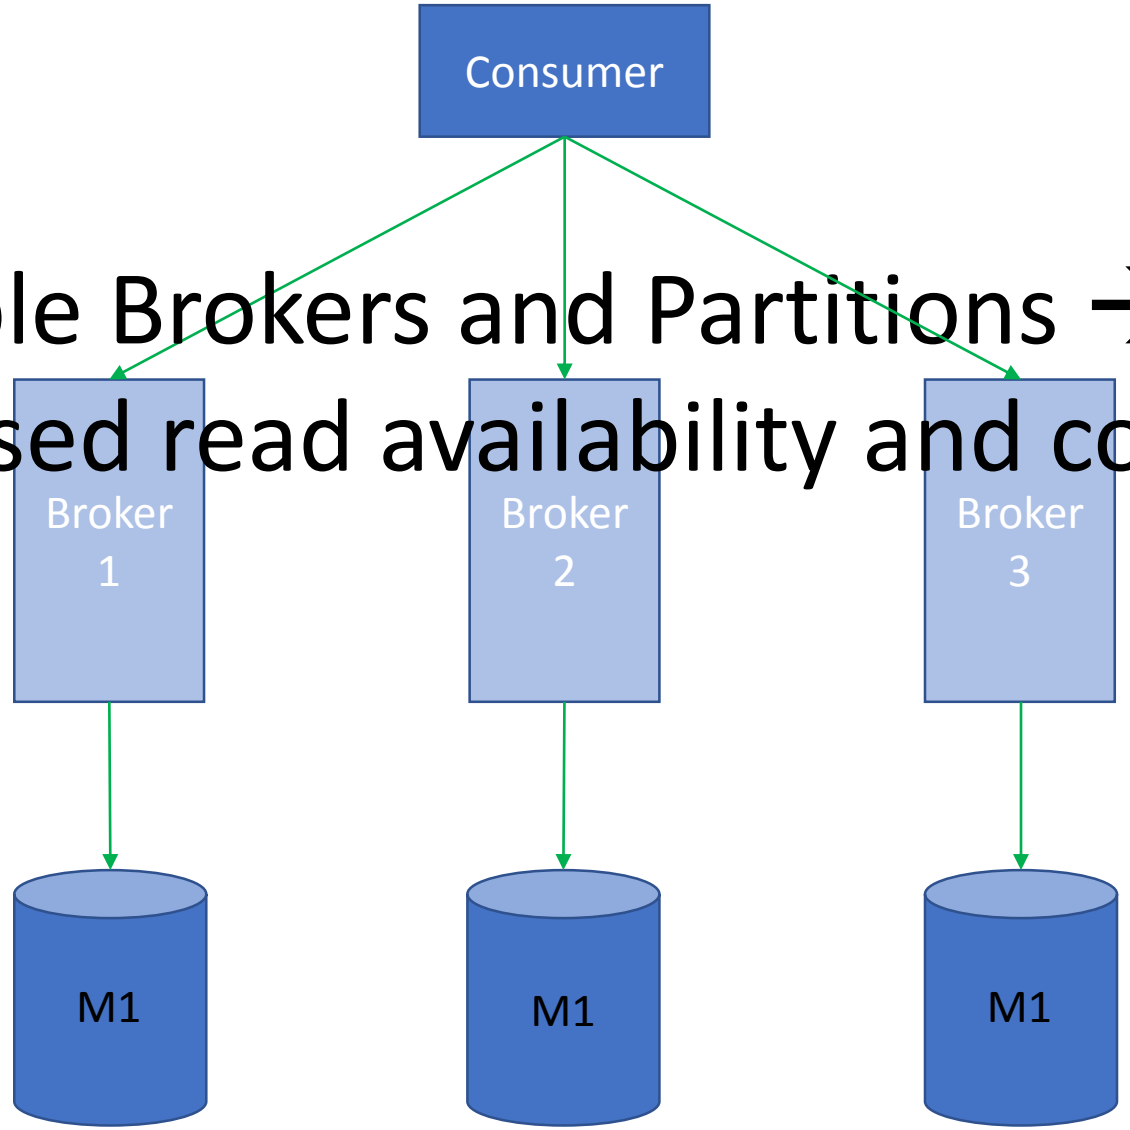

The 2nd aspect of delivery semantics:
Who gets the messages?
How many times are messages delivered?

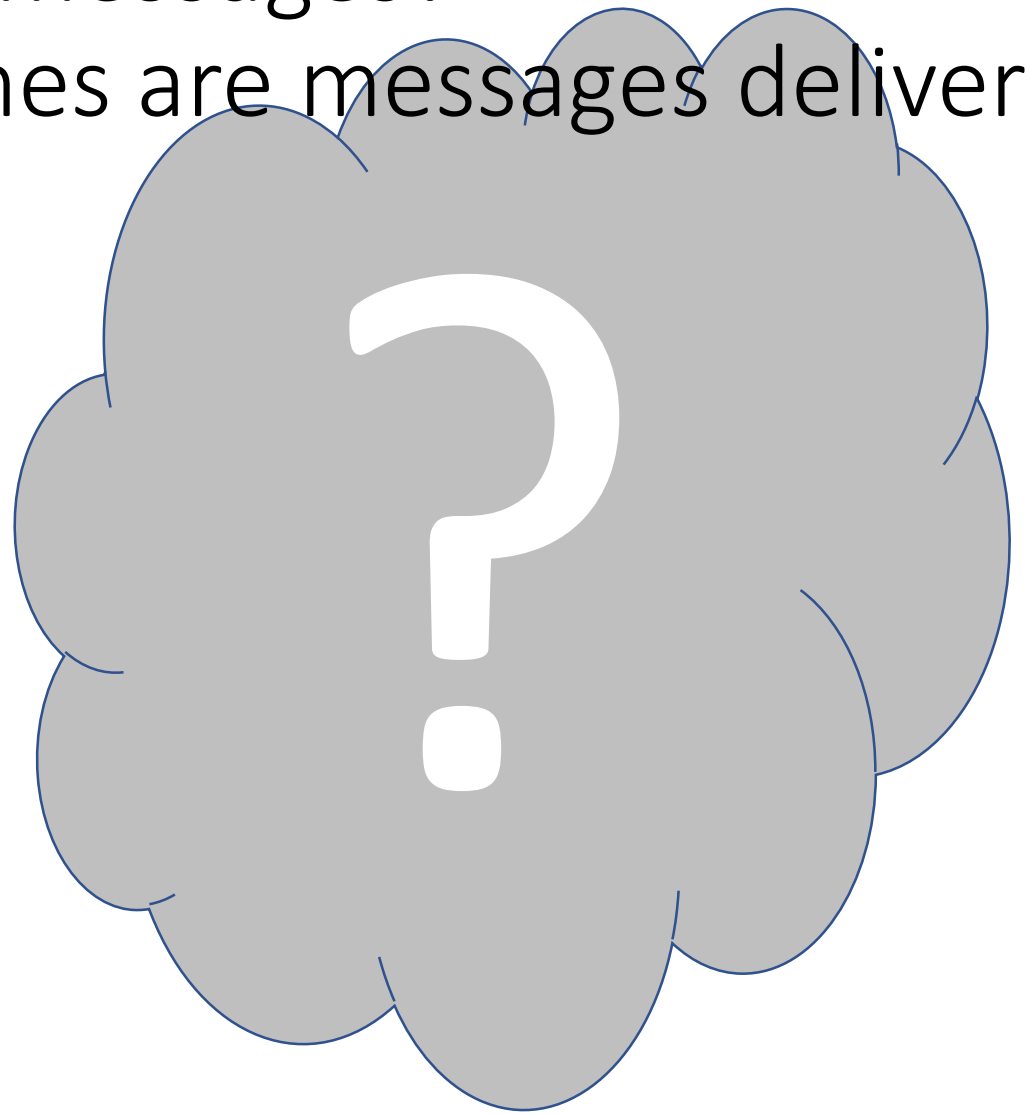

Producer

Consumer

Consumer

Consumer

Consumer

Delivery Semantics - Kafka is “pub-sub”

- Loosely coupled
- Producers and consumers don’t know about each other

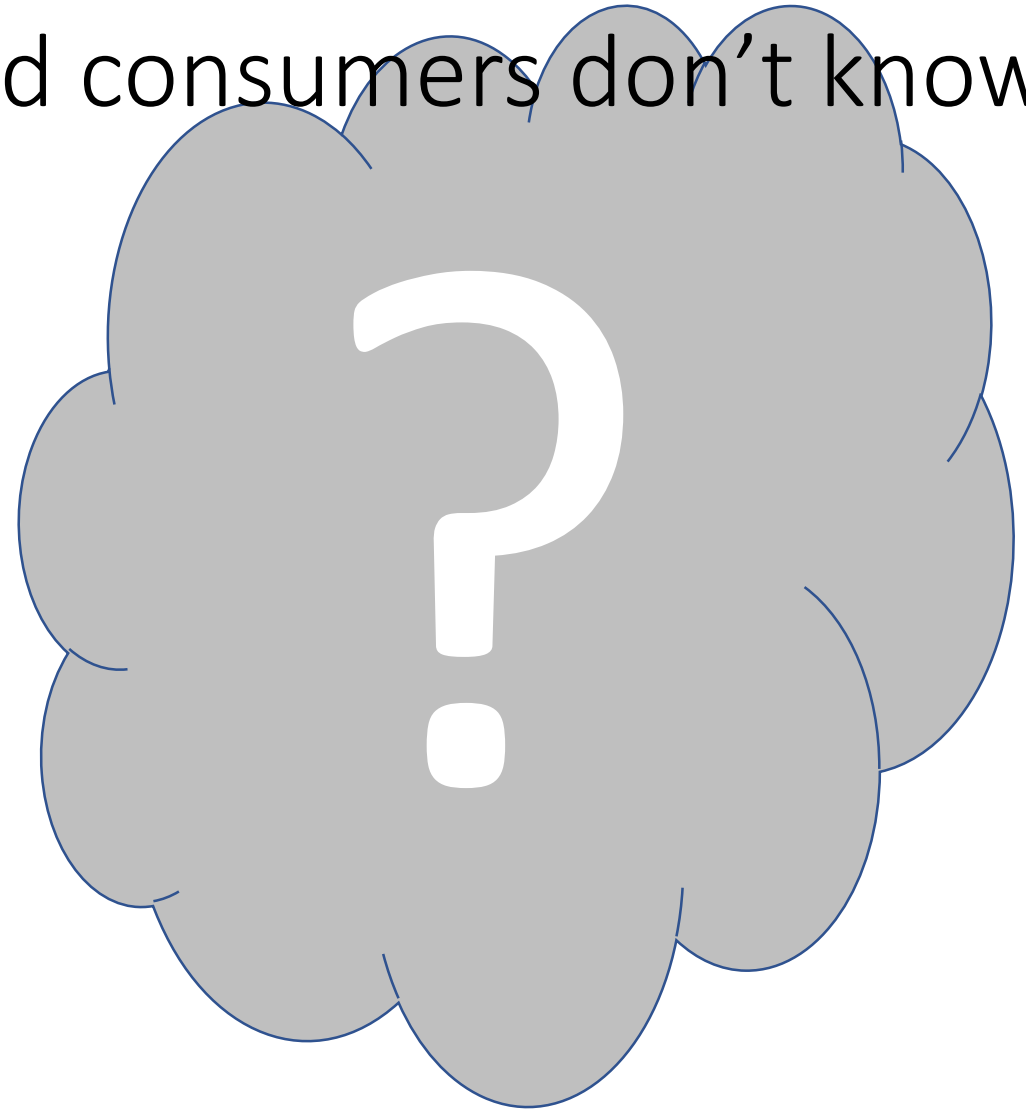

Producer

Consumer

Consumer

Consumer

Consumer

Which consumers get which messages (filtering), is topic based

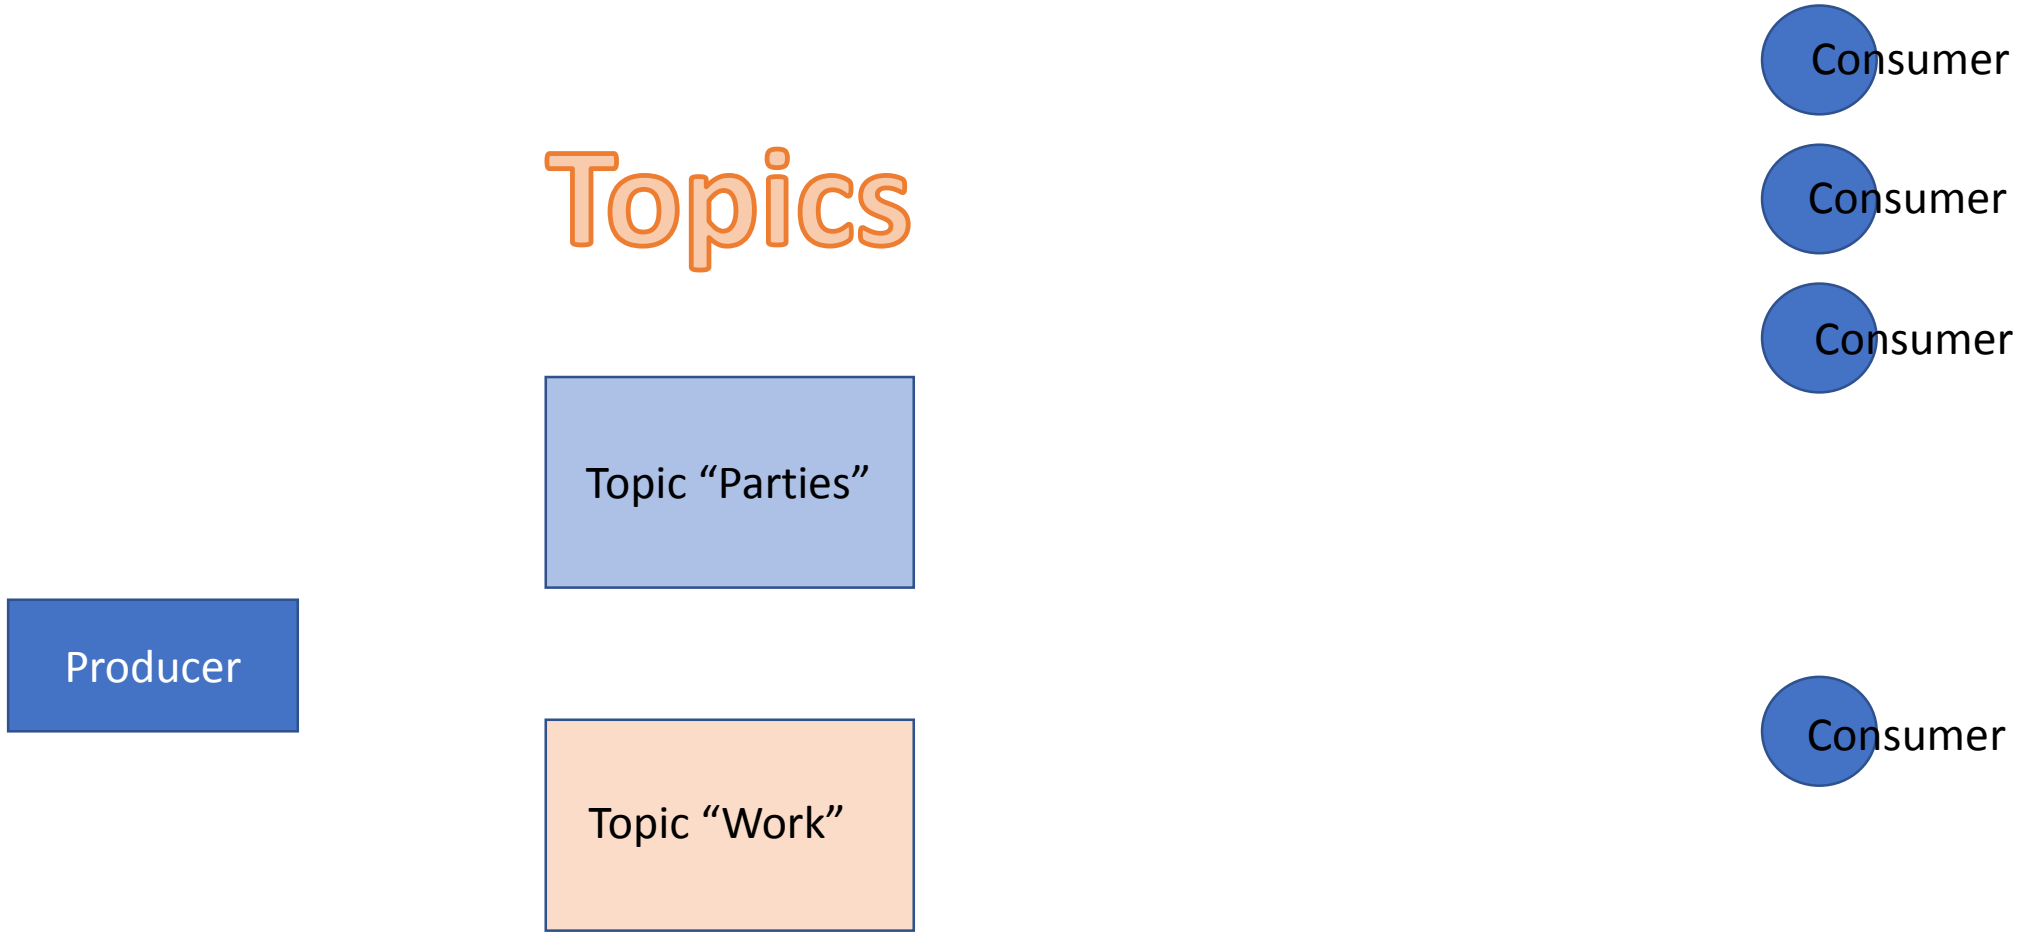

Consumers Subscribe to topic “Parties”

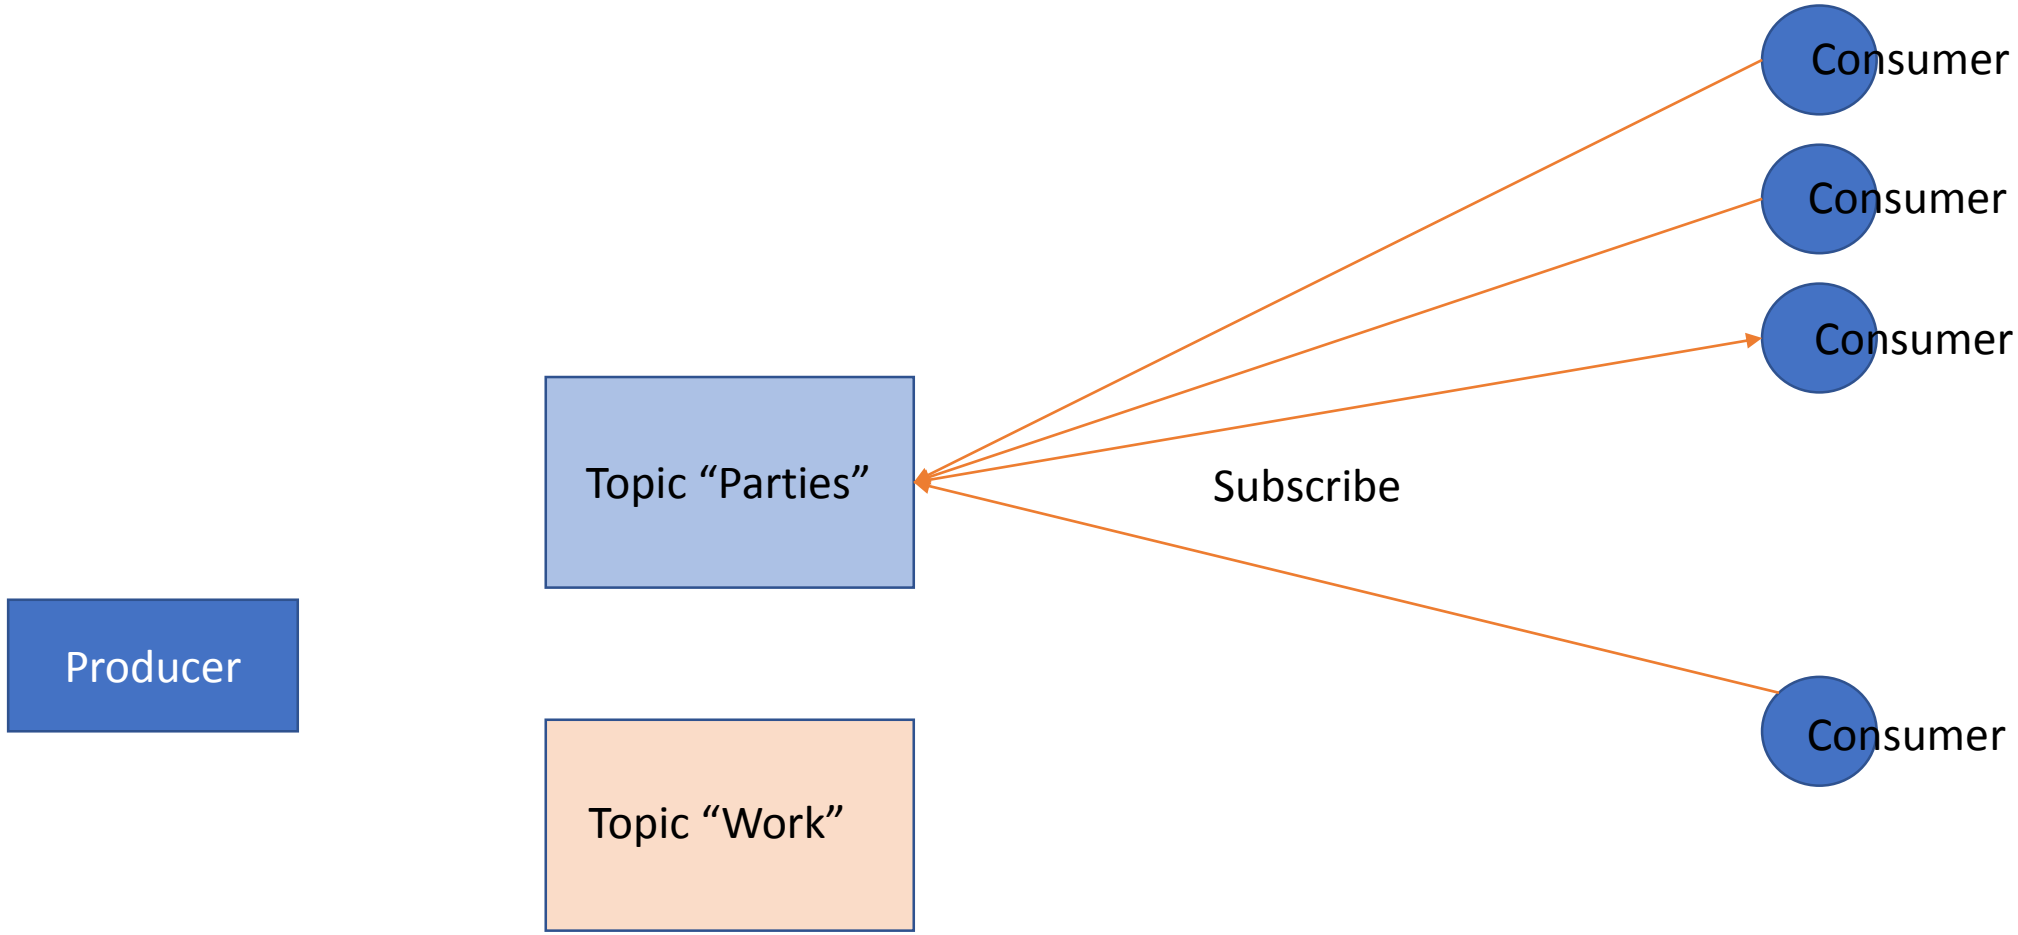

Publishers send messages to topics

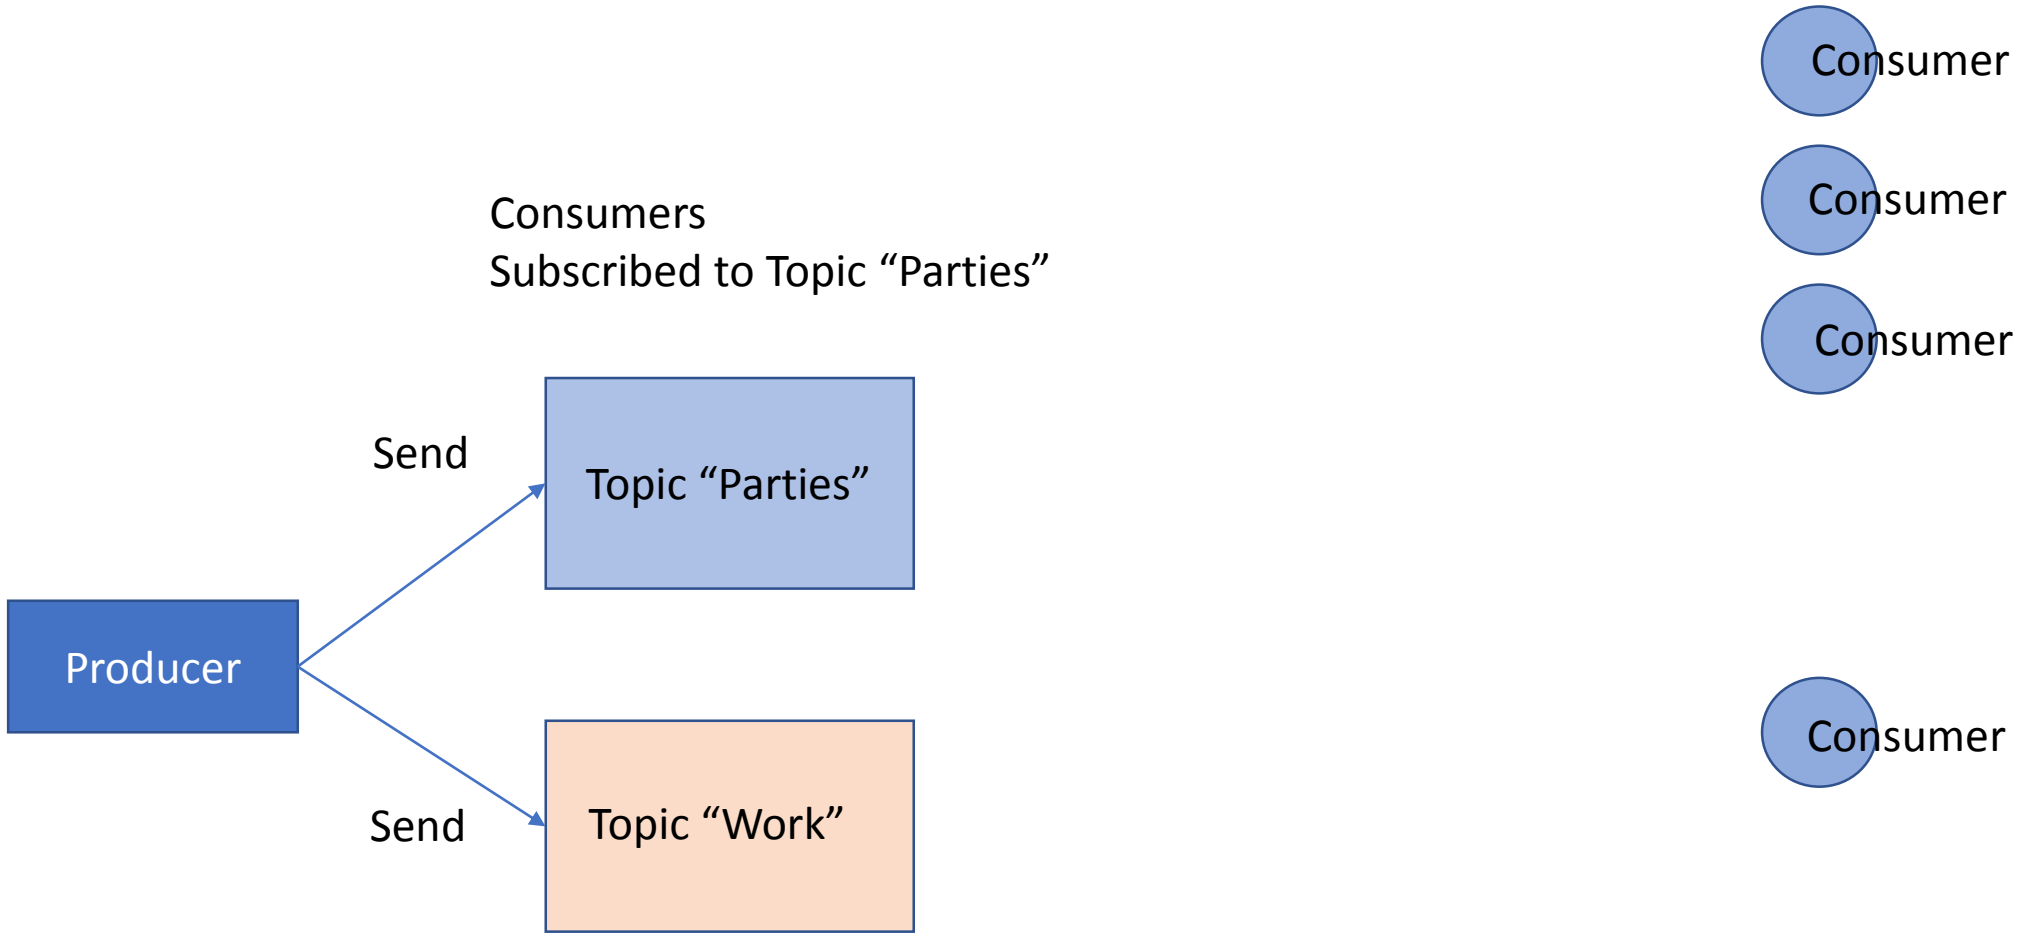

Consumers only receive messages from subscribed topics

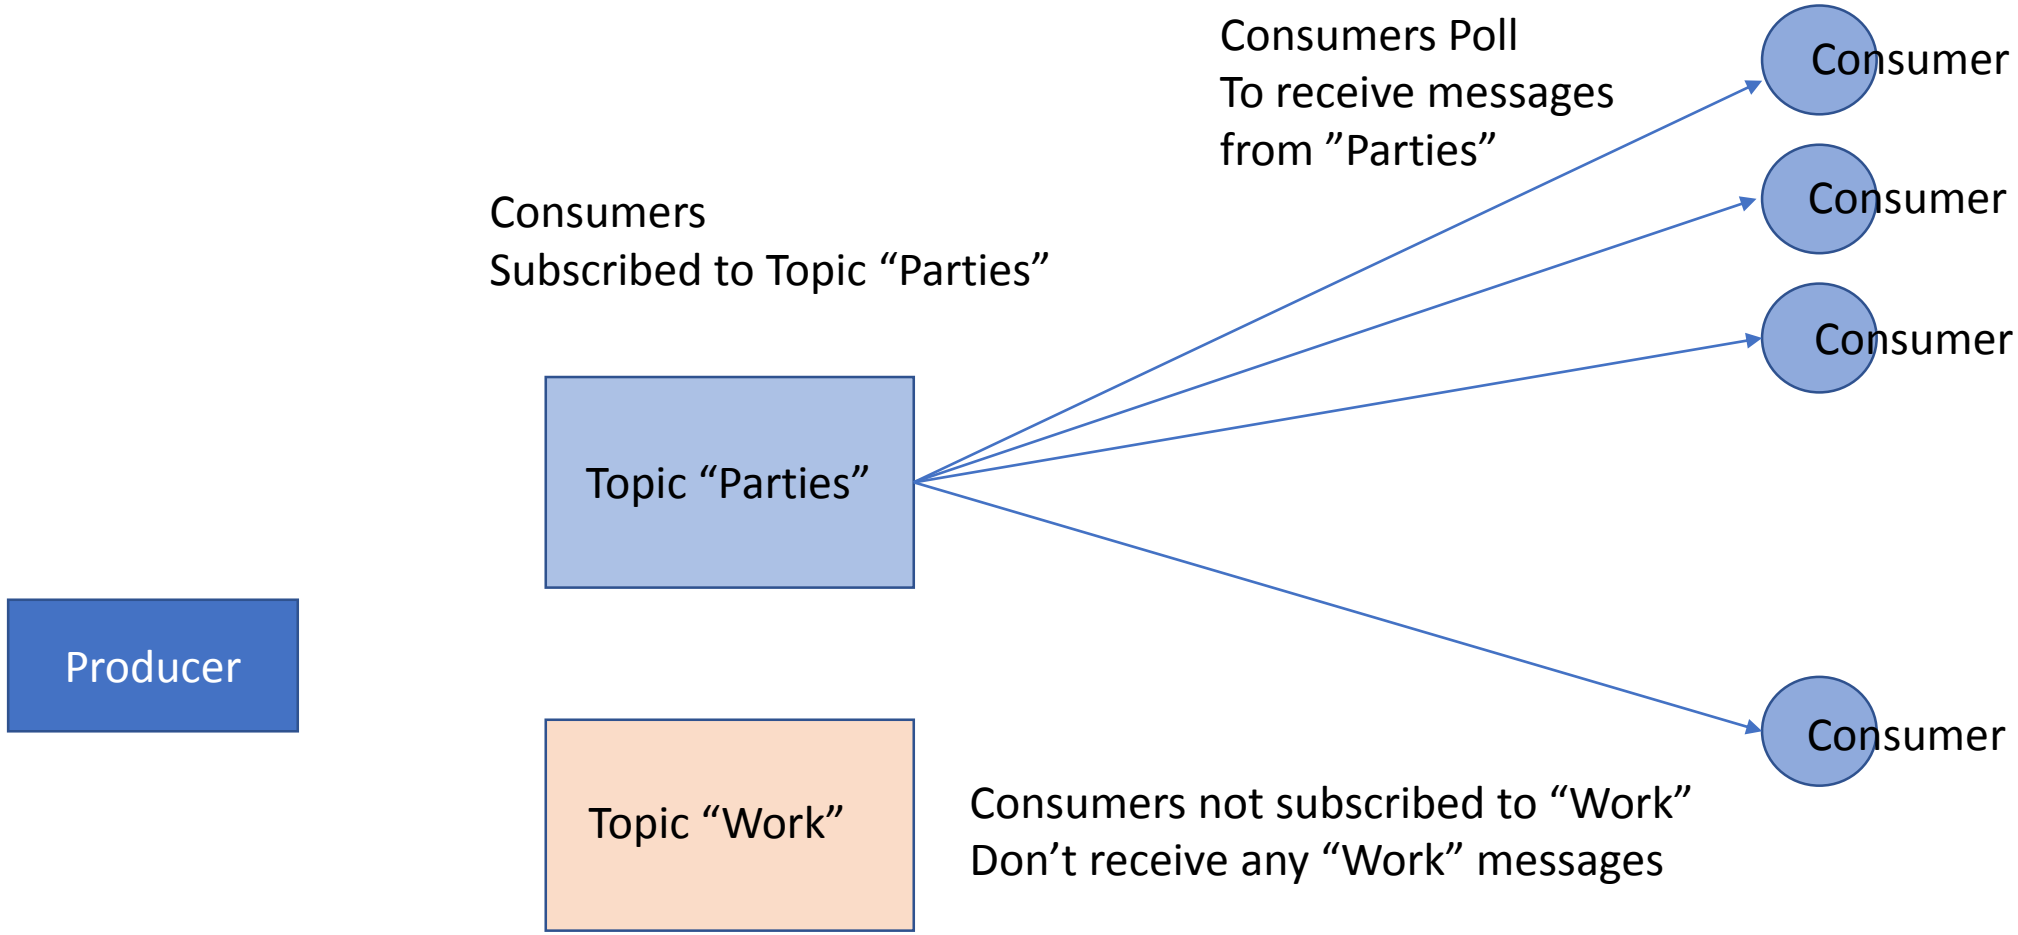

Partitions and Consumer Groups

Enable sharing of work across consumers

Duplicate message delivery

Each message is delivered to each subscribed consumer group

Consumers subscribed to topic are allocated partitions
They will only get messages from their allocated partitions.

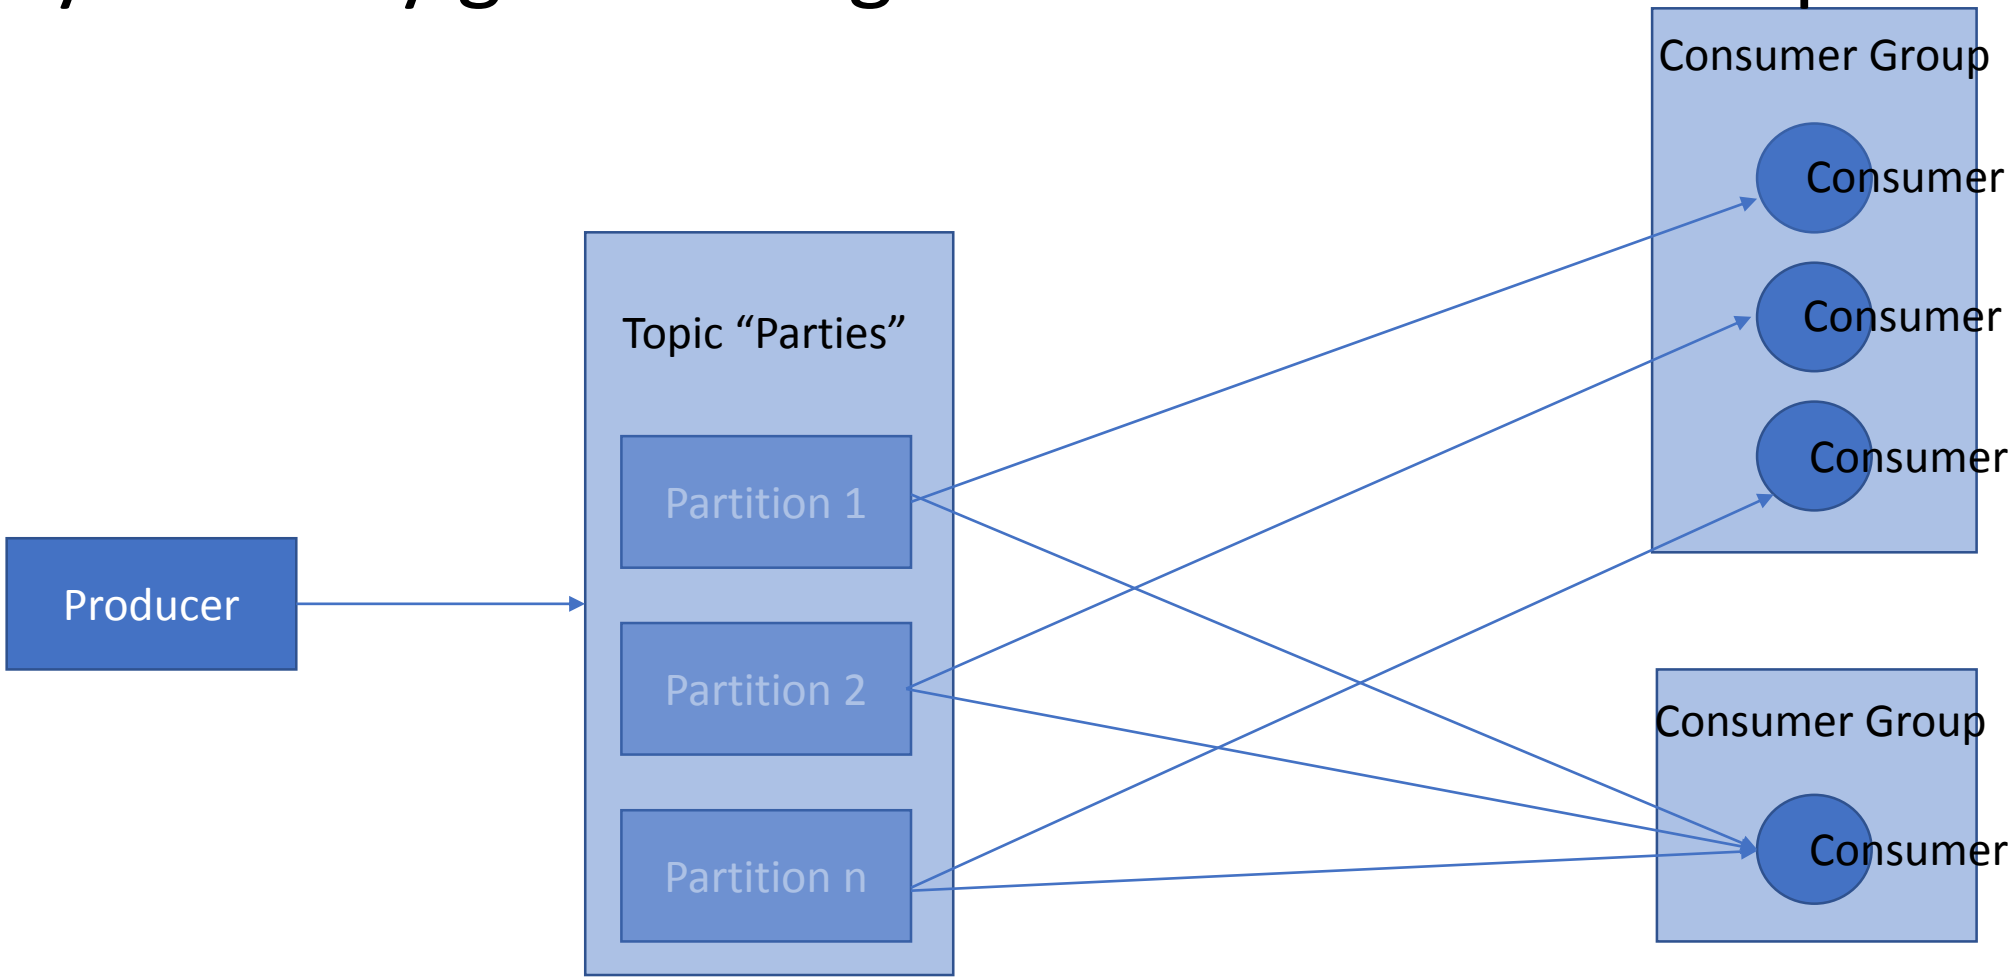

Consumers in the same group share the work around Each consumer gets only a subset of messages

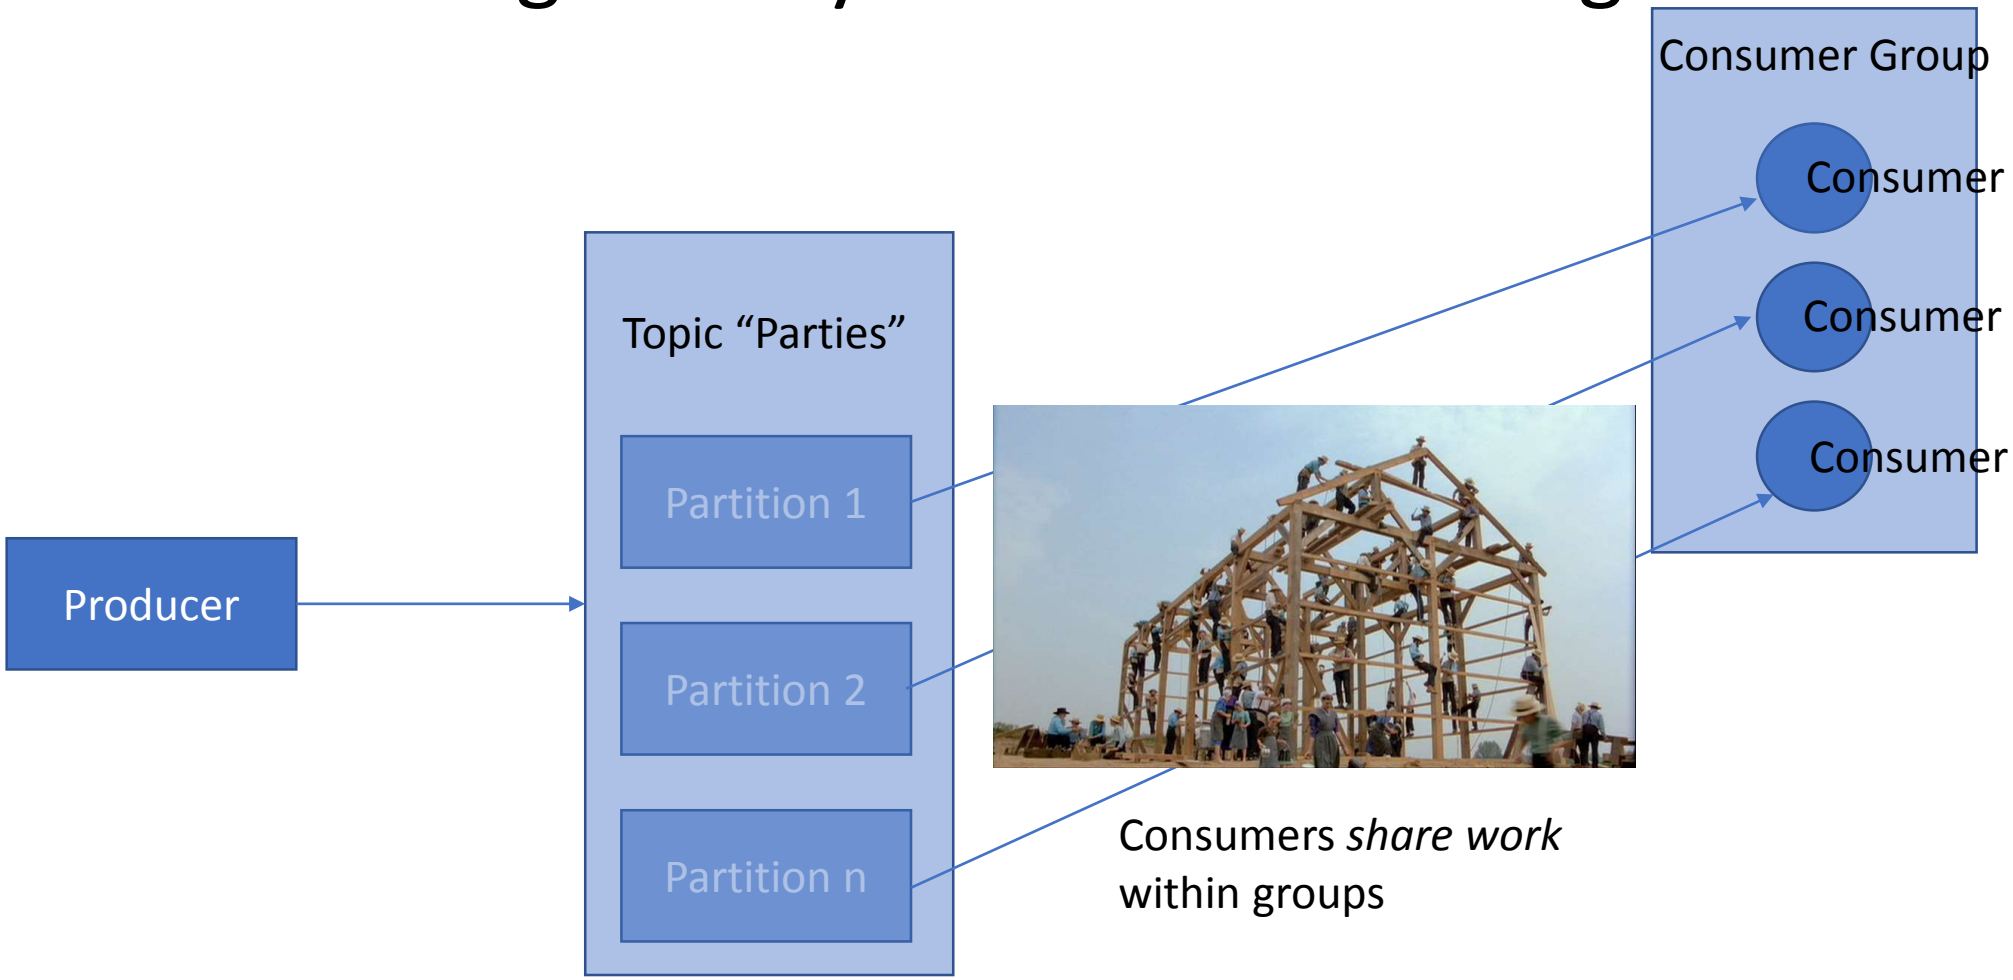

Multiple groups enable message broadcasting

Messages are duplicated across groups, each consumer group receives a copy of each message.

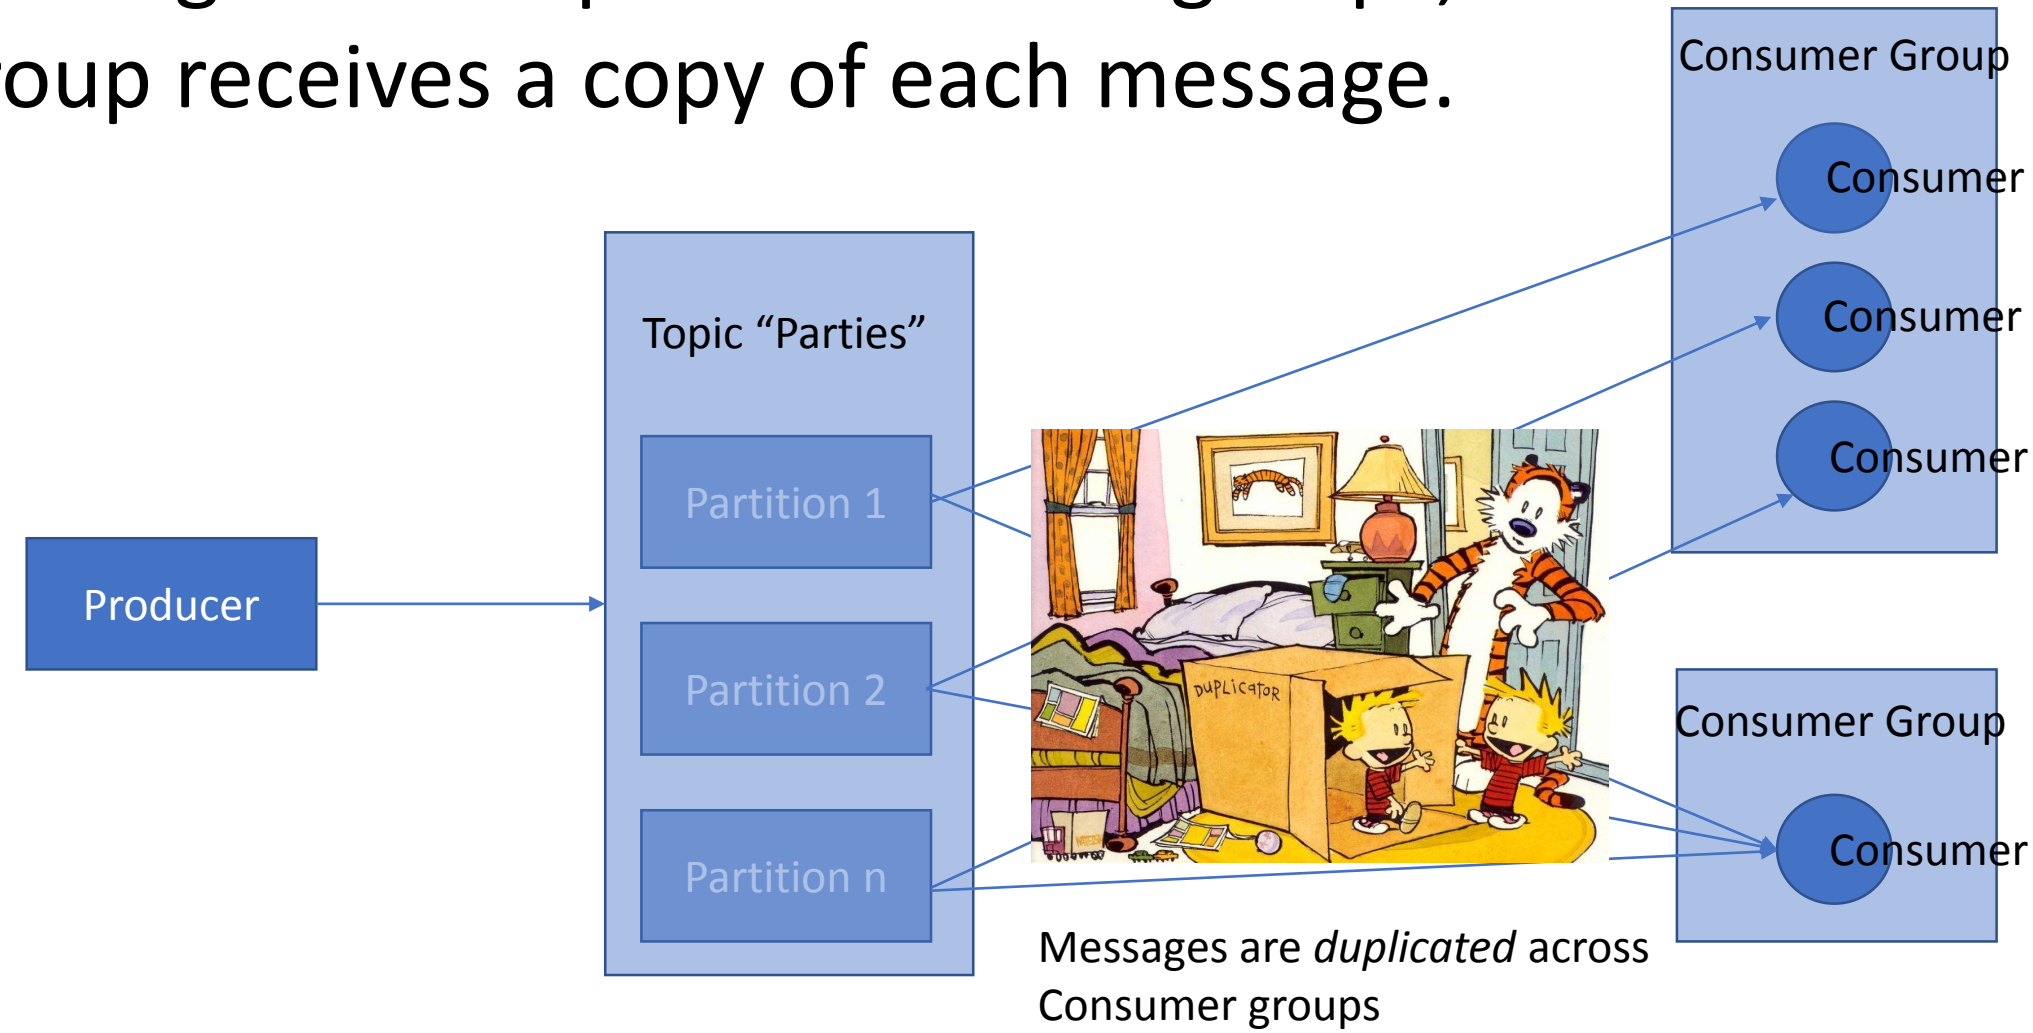

Key

Partition based delivery

Which messages are delivered to which consumers?

If a message has a key, then Kafka uses Partition based delivery.

Messages with the same key are always sent to the same partition and therefore the same consumer.

And the order is guaranteed.

If the key is null, then Kafka uses round robin delivery

Each message is delivered to the next partition

No Key

Round robin delivery

Time for an Example, with 2 consumer groups.

Consumer Group = Nerds
Multiple consumers

Consumer Group = Nerds
Multiple consumers

Consumer Group = Hairy
Single consumer

Case 1: No Key

Consumers
Subscribed to “Parties”

Group “Nerds”

Consumer 1 (Bill)

Consumer 2 (Paul)

Consumer n

No Key
Round Robin

Topic “Parties”

M1
Partition 1

M2
Partition 2

Partition n

Producer

M1

M2

etc

M1

M2

Group “Hairy”

M1

M2

Consumer 1 (Chewy)

Message (M1, M2, etc) sent to the next partition
All consumers allocated to that partition will receive a message when they poll next.

Here’s what happens (not showing producer or topics, have to imagine them)

1. Both Groups subscribe to Topic “parties” (11 partitions, so 1 consumer per partition).

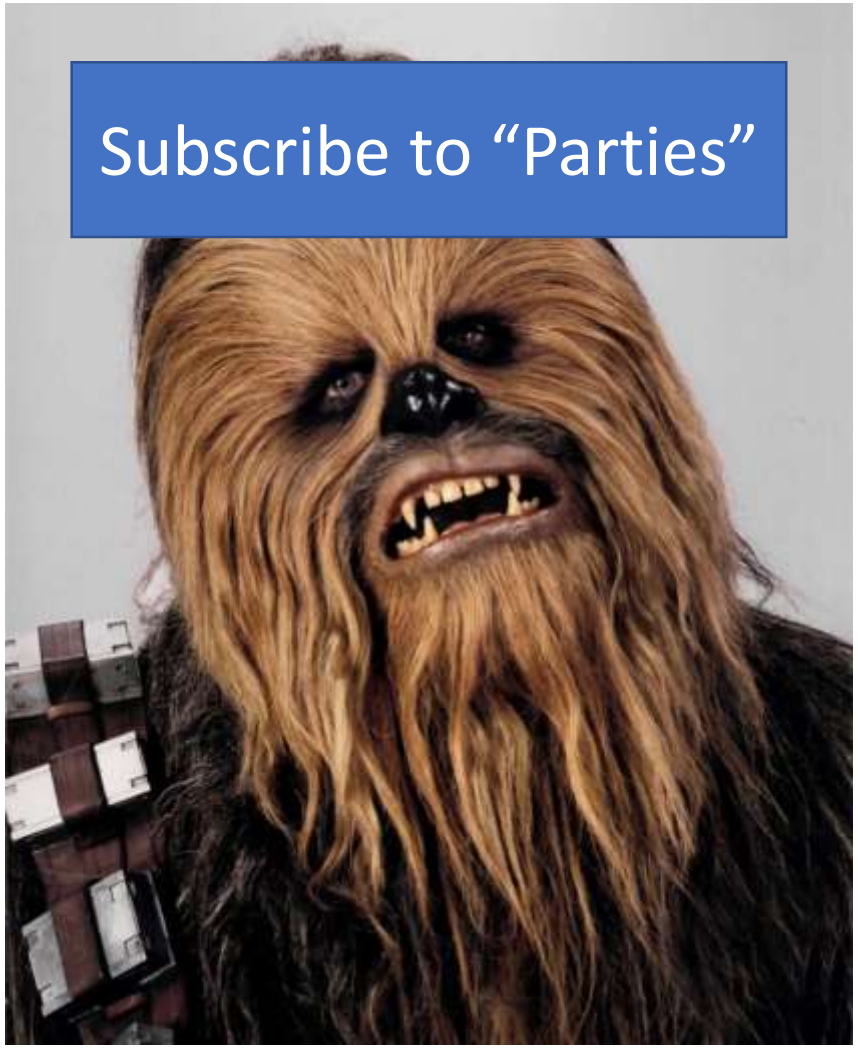

1. Both Groups subscribe to Topic “parties” (11 partitions, so 1 consumer per partition).
2. Producer sends record “Cool pool party – Invitation”
   <key=null, value=“Cool pool party - Invitation”> to “parties” topic (no key)

1. Both Groups subscribe to Topic “parties” (11 partitions, so 1 consumer per partition).
2. Producer sends record “Cool pool party - Invitation”> to “parties” topic
3. Bill and Chewbacca receive a copy of the invitation and plan to attend

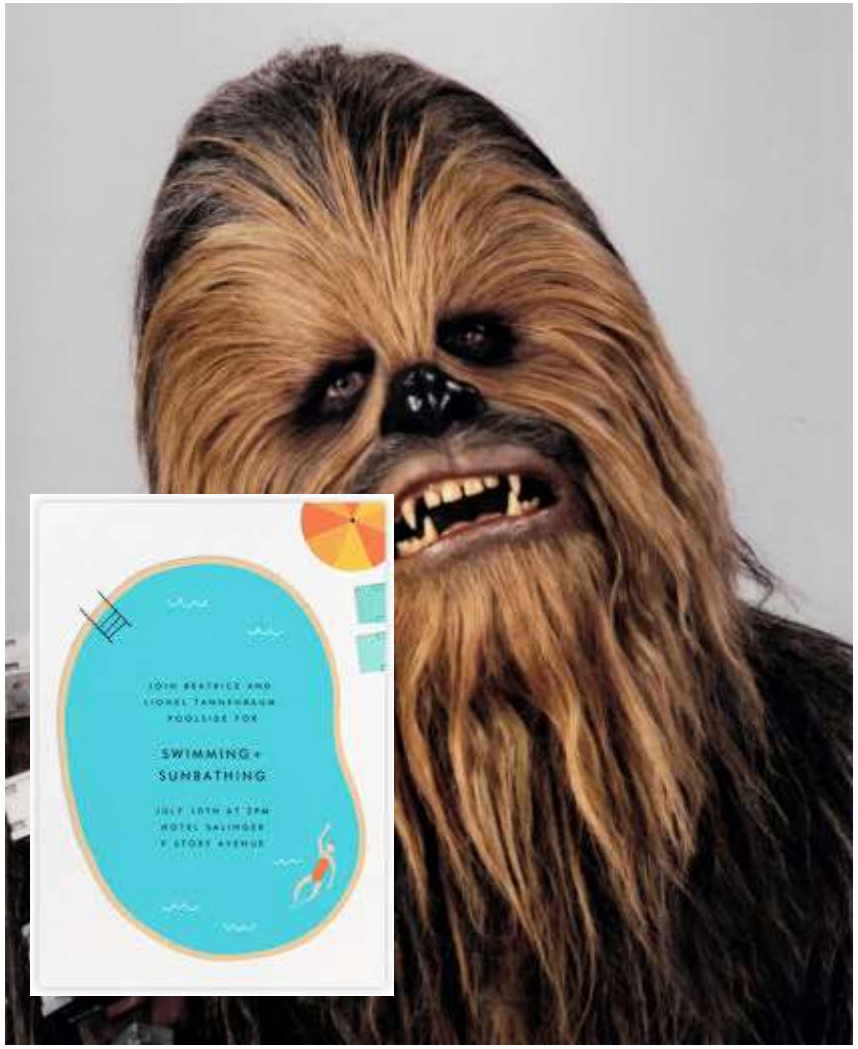

4. Producer sends another record “Cool pool party – Cancelled” <key=null, value=“Cool pool party - Cancelled”> to “parties” topic

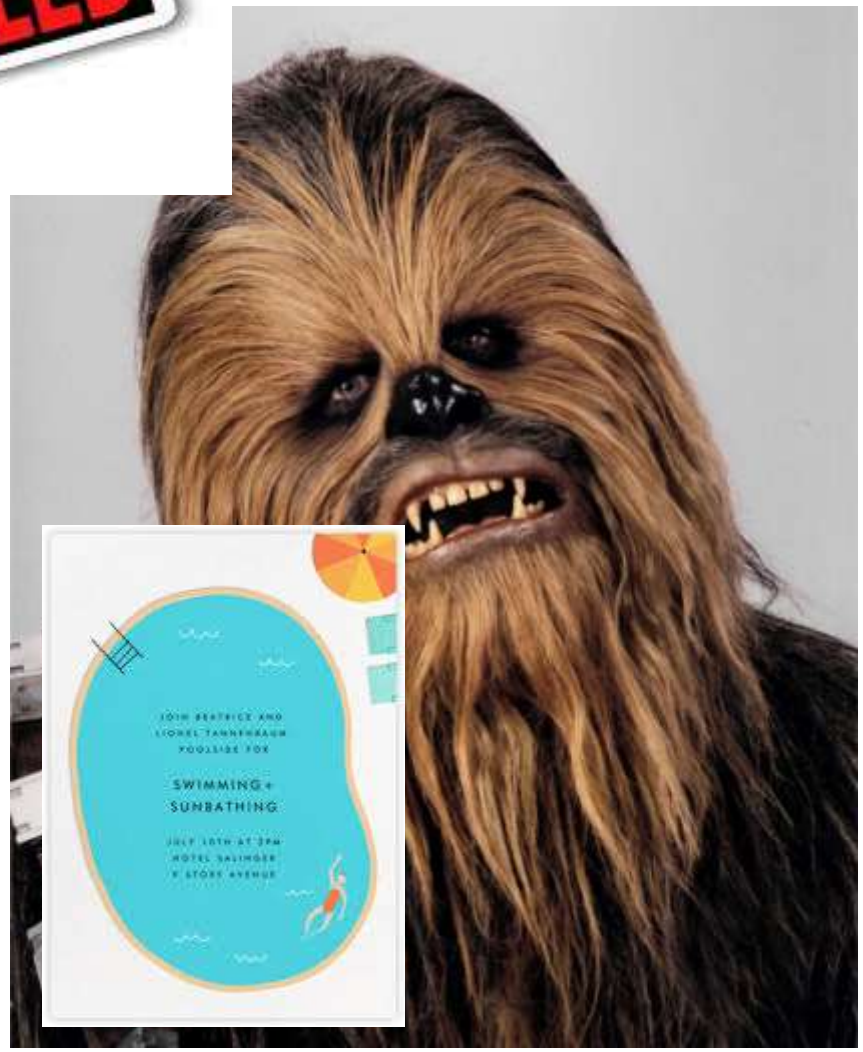

4. Producer sends another record <key=null, value="Cool pool party - Cancelled"> to "parties" topic
5. Paul and Chewbacca receive the cancellation.

Paul gets the message this time as it's round robin, ignores it as he didn't get the invitation. Bill wastes his time trying to go to cancelled party. The rest of the gang aren't surprised at not receiving any party invites and stay at home to do some hacking. Chewy is only consumer in his group so gets all messages, plans something fun instead…

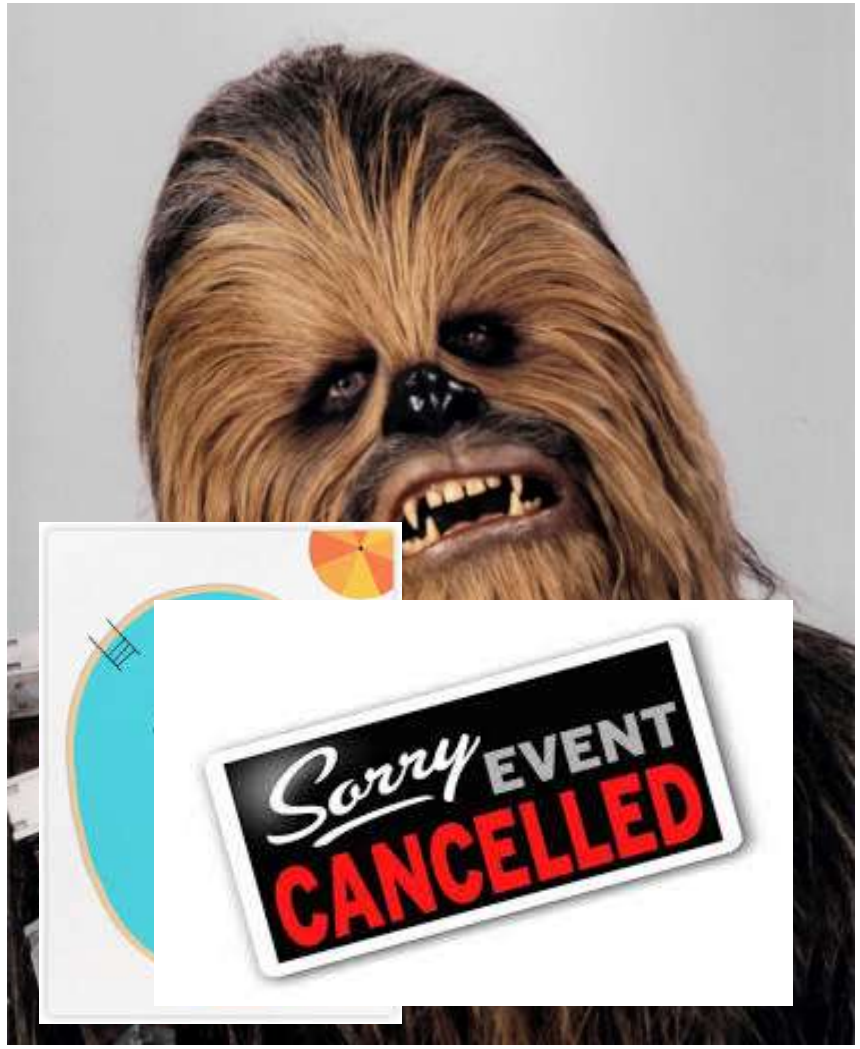


A visit to the hairdressers!

Case 2: If there is a Key

Consumers
Subscribed to “Parties”

Group “Nerds”

Consumer 1 (Bill)

Consumer 2 (Paul)

Consumer n

Key
Hashed to partition

Topic “Parties”

Partition 1

Partition 2

Partition n

Producer

M1, M2

M3

etc

M1, M2

M3

M1, M2

M3

Group “Hairy”

Consumer 1 (Chewy)

A key is hashed to a partition, and a Message with that key is always sent to that partition.
Assume there are 3 messages, and messages 1 and 2 are hashed to same partition.

Here's what happens with a key: key is "title" of the message (e.g. "Cool pool party")
Same set up as before:

1. Both Groups subscribe to Topic "parties" (11 partitions).

1. Both Groups subscribe to Topic “parties” (11 partitions).
2. Producer sends record <key=“Cool pool party”, value=“Invitation”> to “parties” topic

1. Both Groups subscribe to Topic “parties” (11 partitions).
2. Producer sends record <key=“Cool pool party”, value=“Invitation”> to “parties” topic
3. As before Bill and Chewbacca receive a copy of the invitation and plan to attend

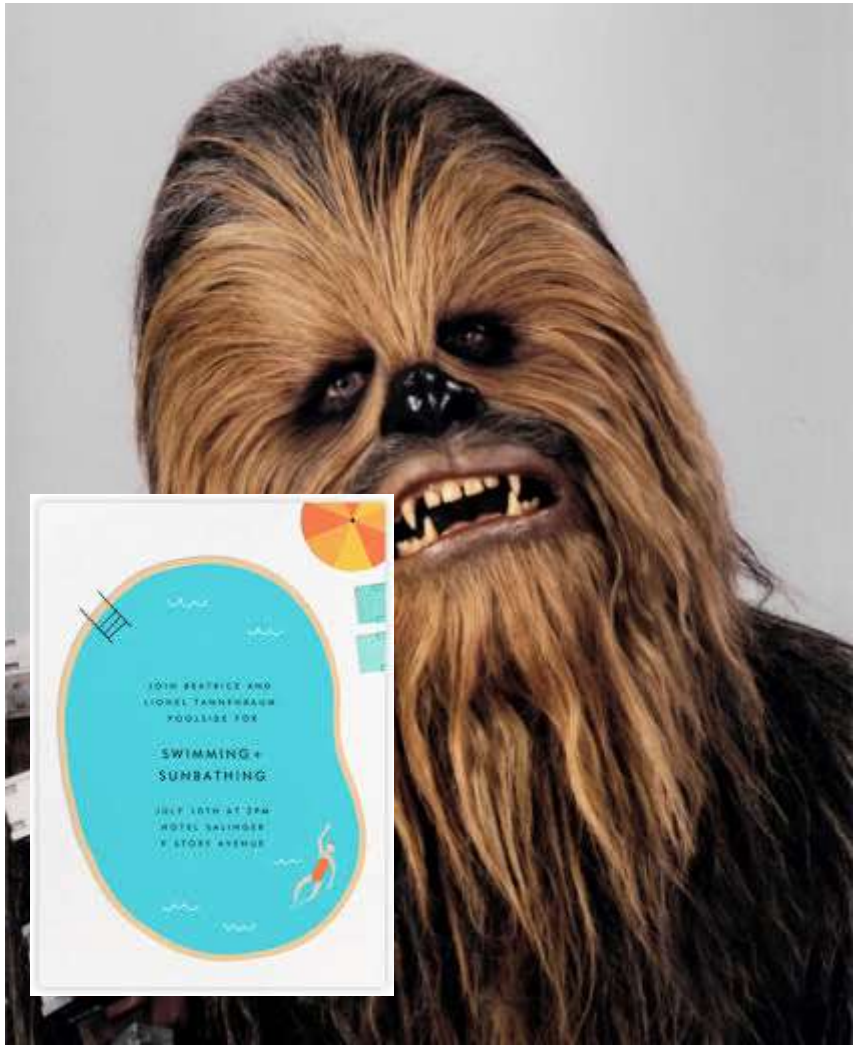

4. Producer sends another record <key="Cool pool party", value="Cancelled"> to "parties" topic

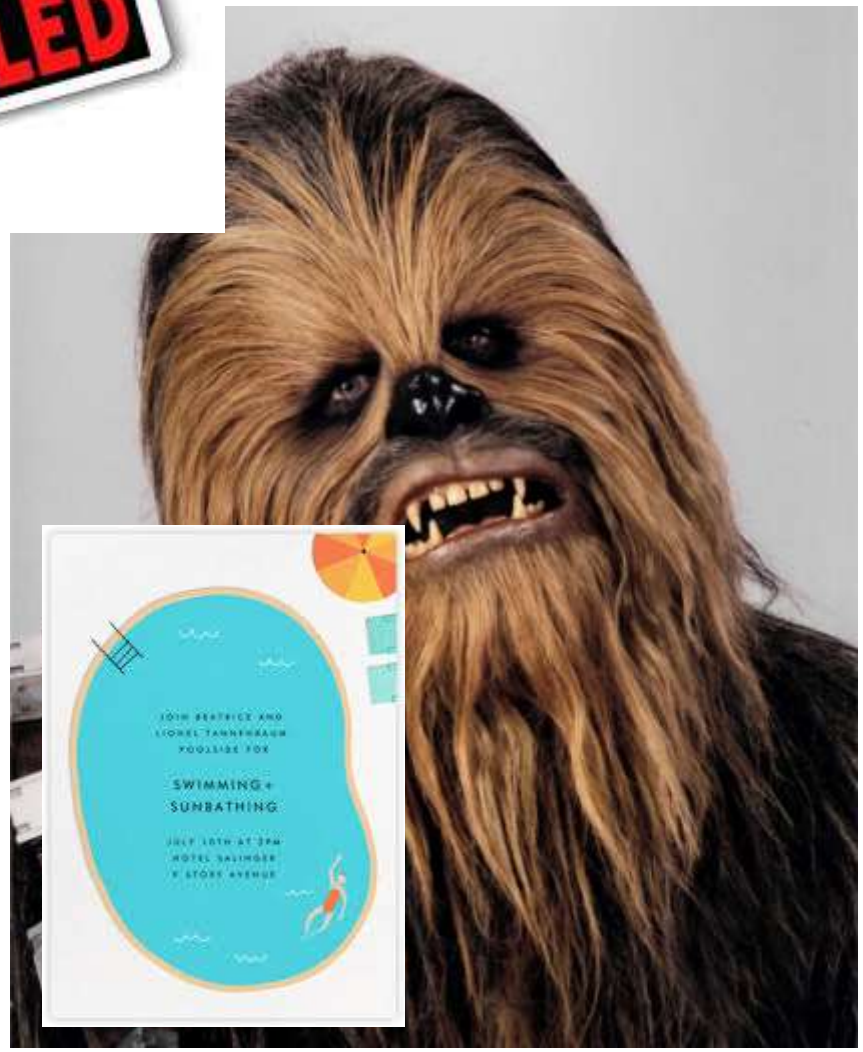

4. Producer sends another record <key="Cool pool party", value="Cancelled"> to "parties" topic
5. Bill and Chewbacca receive the cancellation (same consumers this time, as identical key)

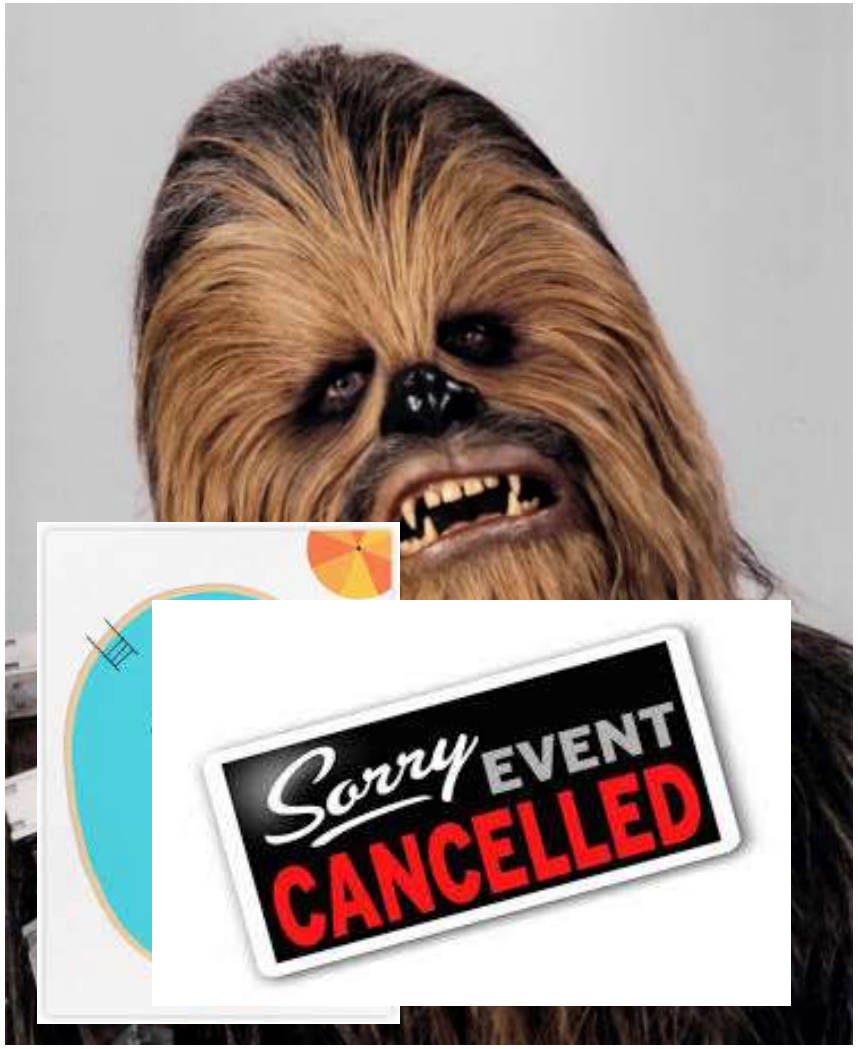

6. Producer sends another record <key="Horrible Halloween party", value="Invitation"> to "parties" topic

6. Producer sends another record <key="Horrible Halloween party", value="Invitation"> to "parties" topic
7. Paul and Chewy receive the invitation

Paul receives the Halloween invitation as the key is different and the record is sent to the partition that Paul is allocated to

Chewy is the only consumer in his group so he gets every record no matter what partition it's sent to


Chewy gets dressed up
And goes to the party

Who did you imagine was producing the invites?

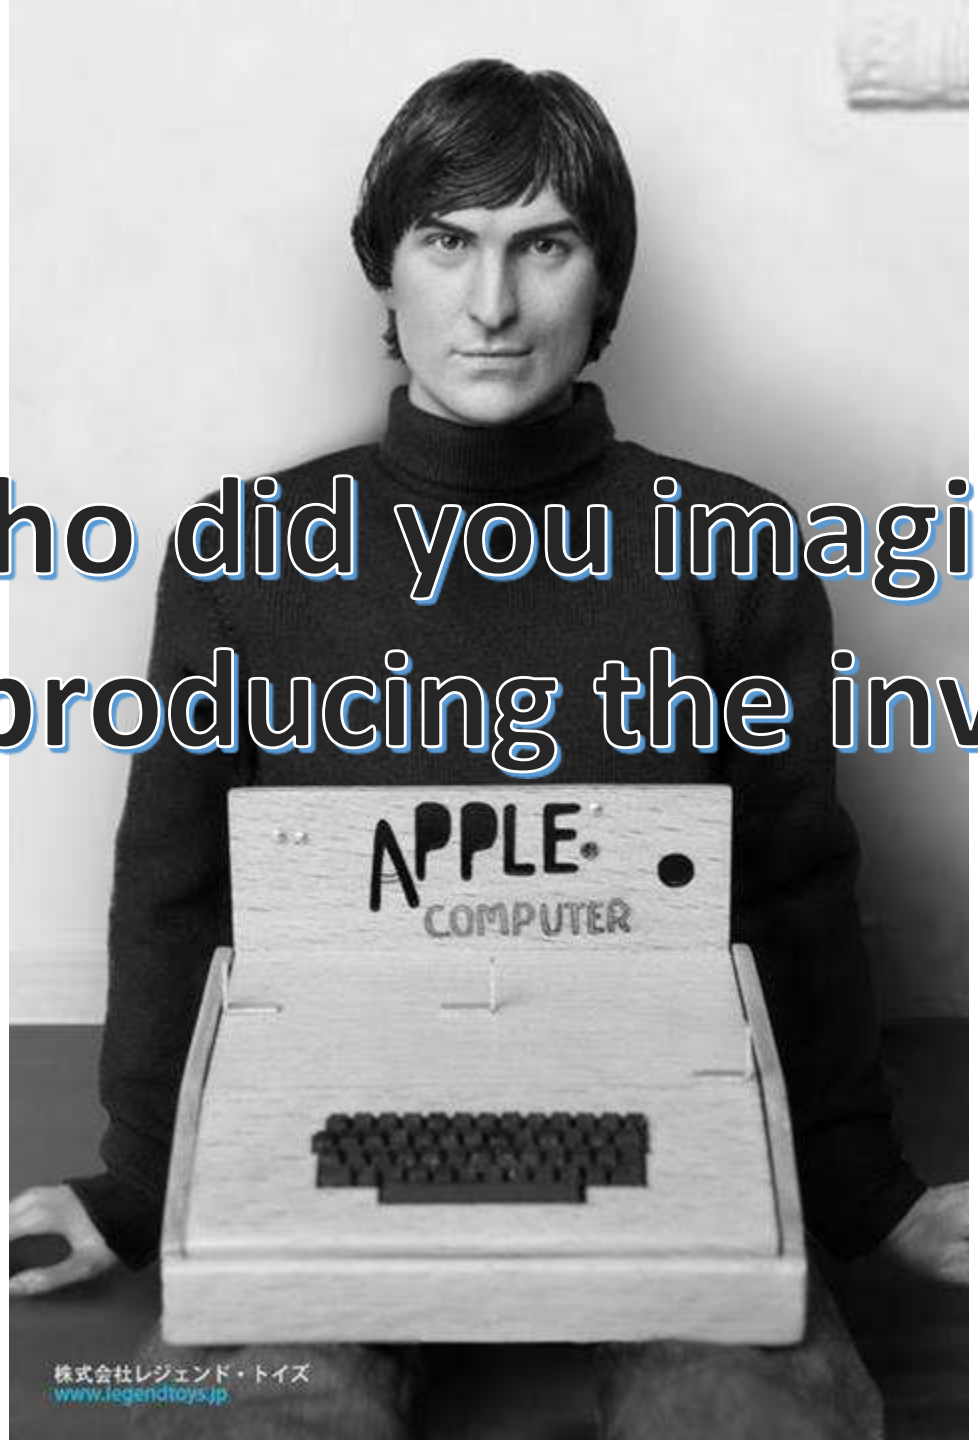

Example Kafka Use Cases

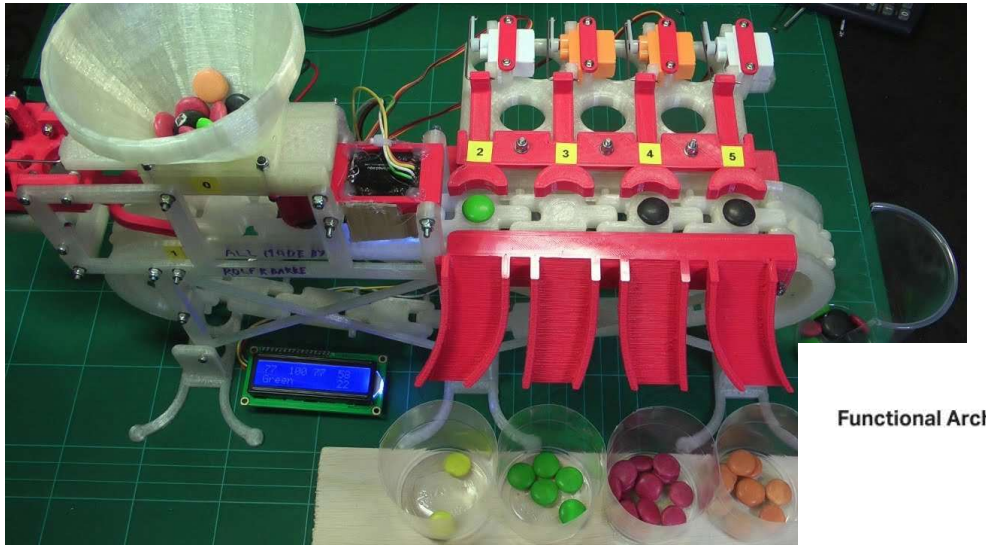

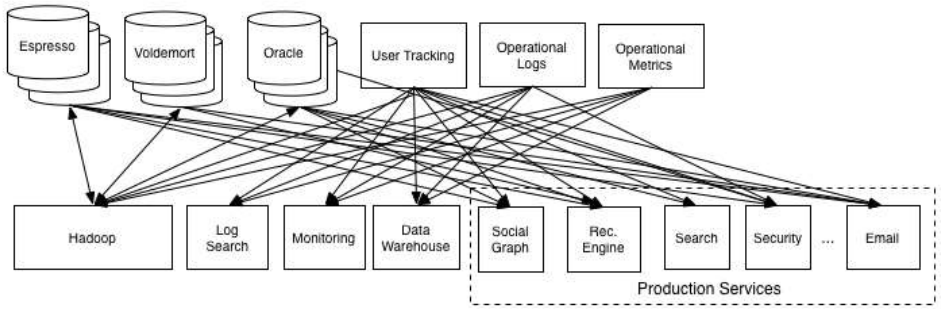

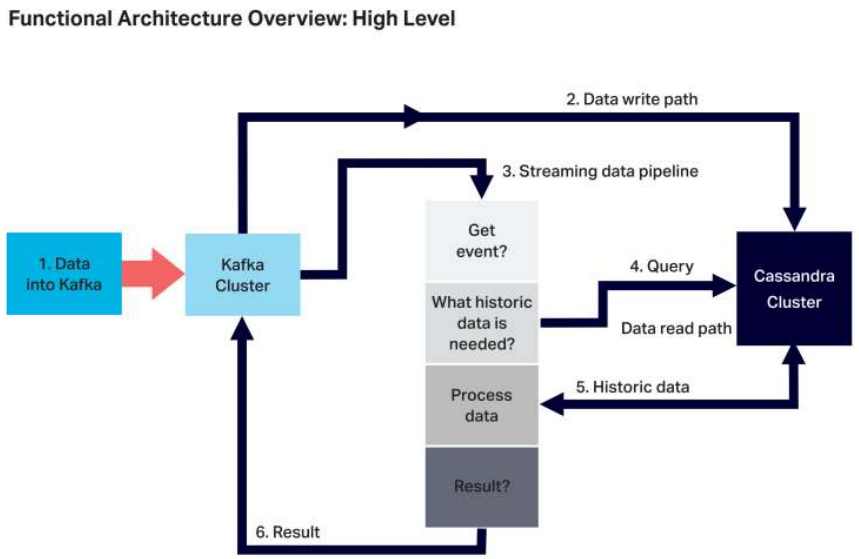

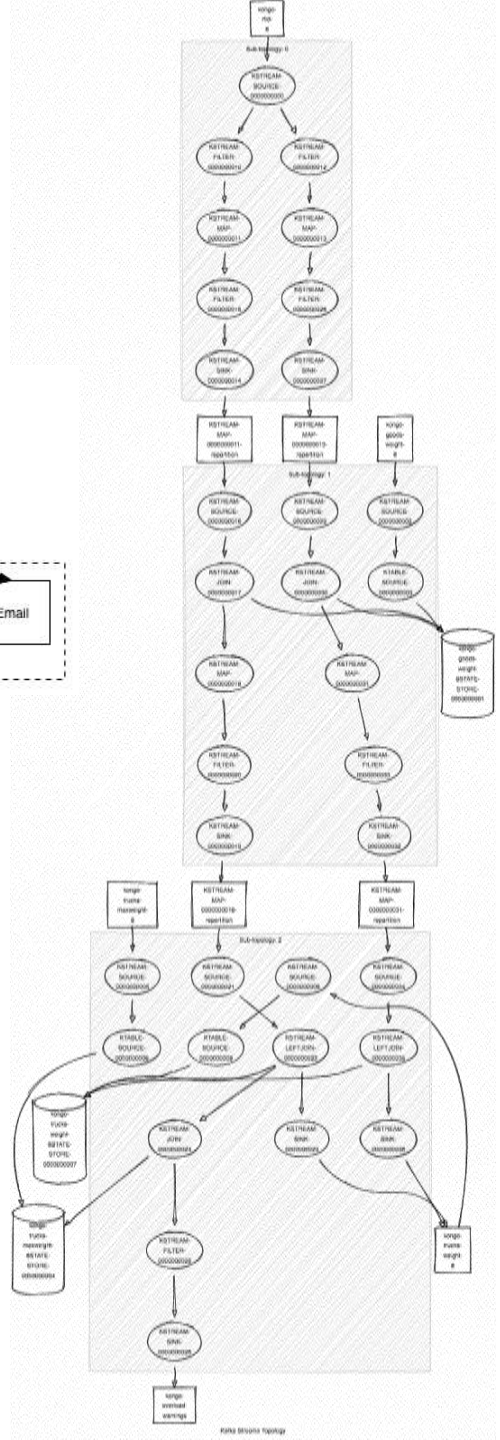

Real-time data pipeline

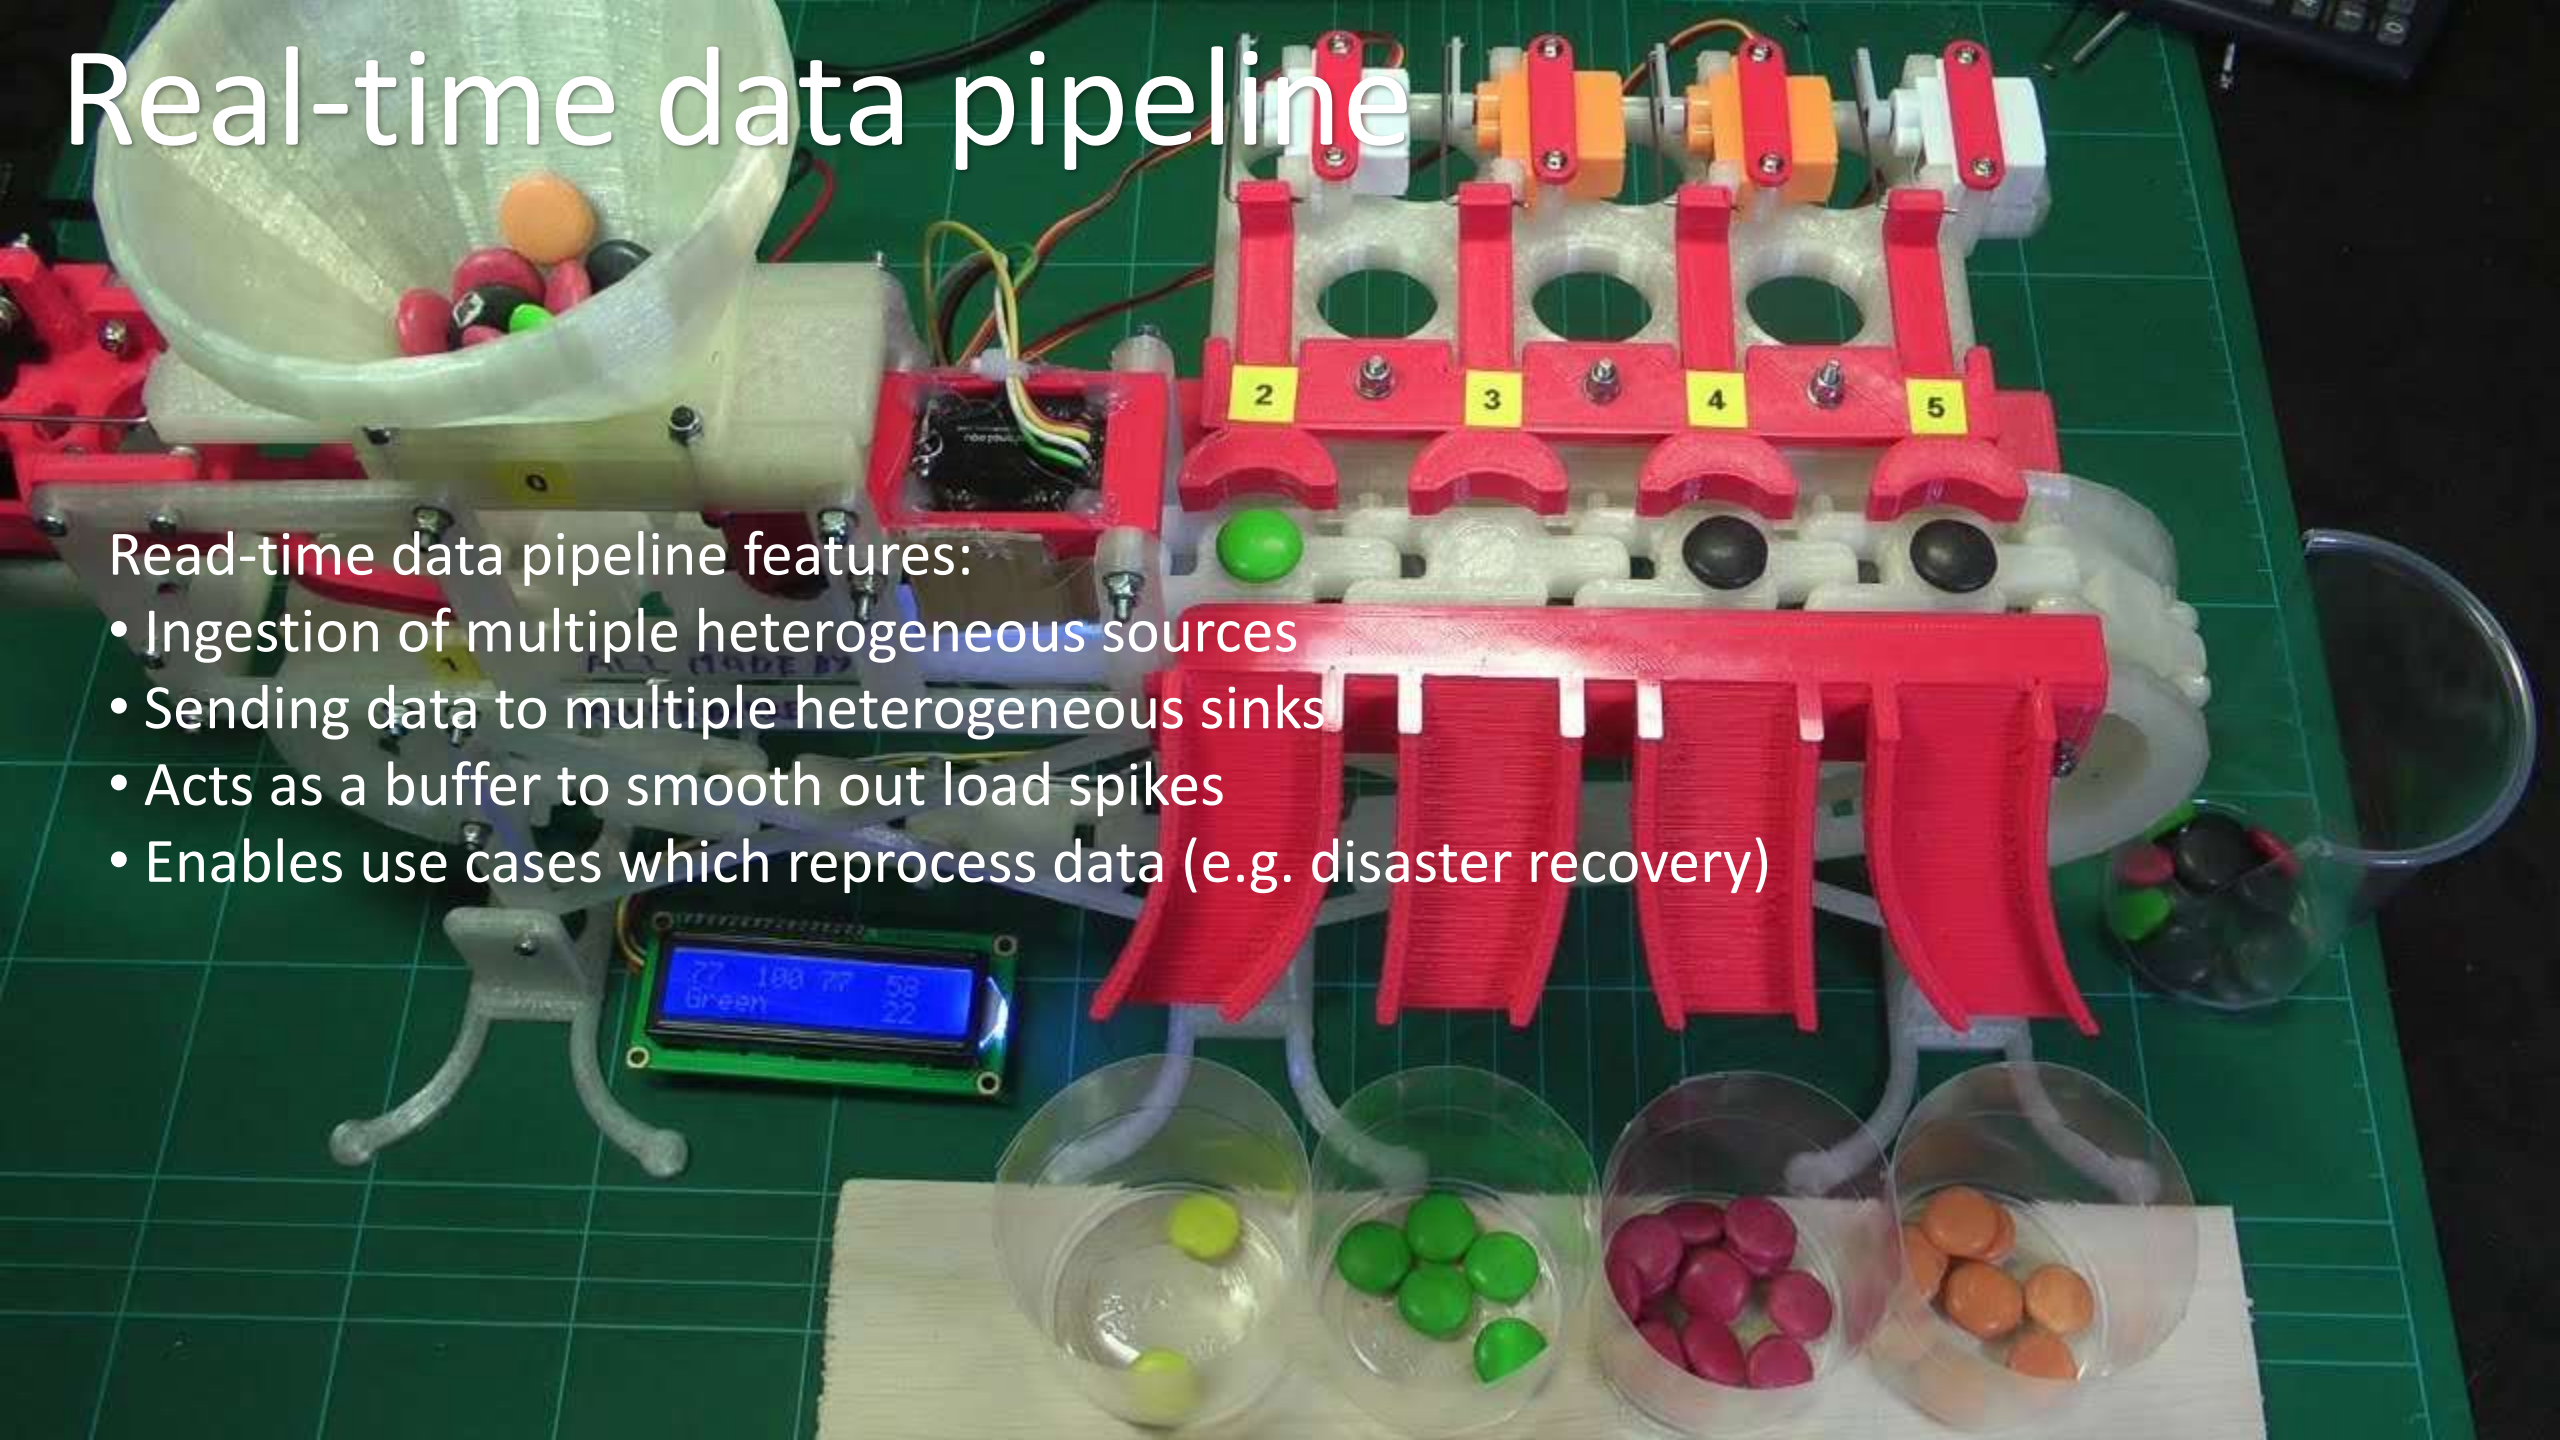

Read-time data pipeline features:

- Ingestion of multiple heterogeneous sources
- Sending data to multiple heterogeneous sinks
- Acts as a buffer to smooth out load spikes
- Enables use cases which reprocess data (e.g. disaster recovery)

Anomaly Detection Pipeline

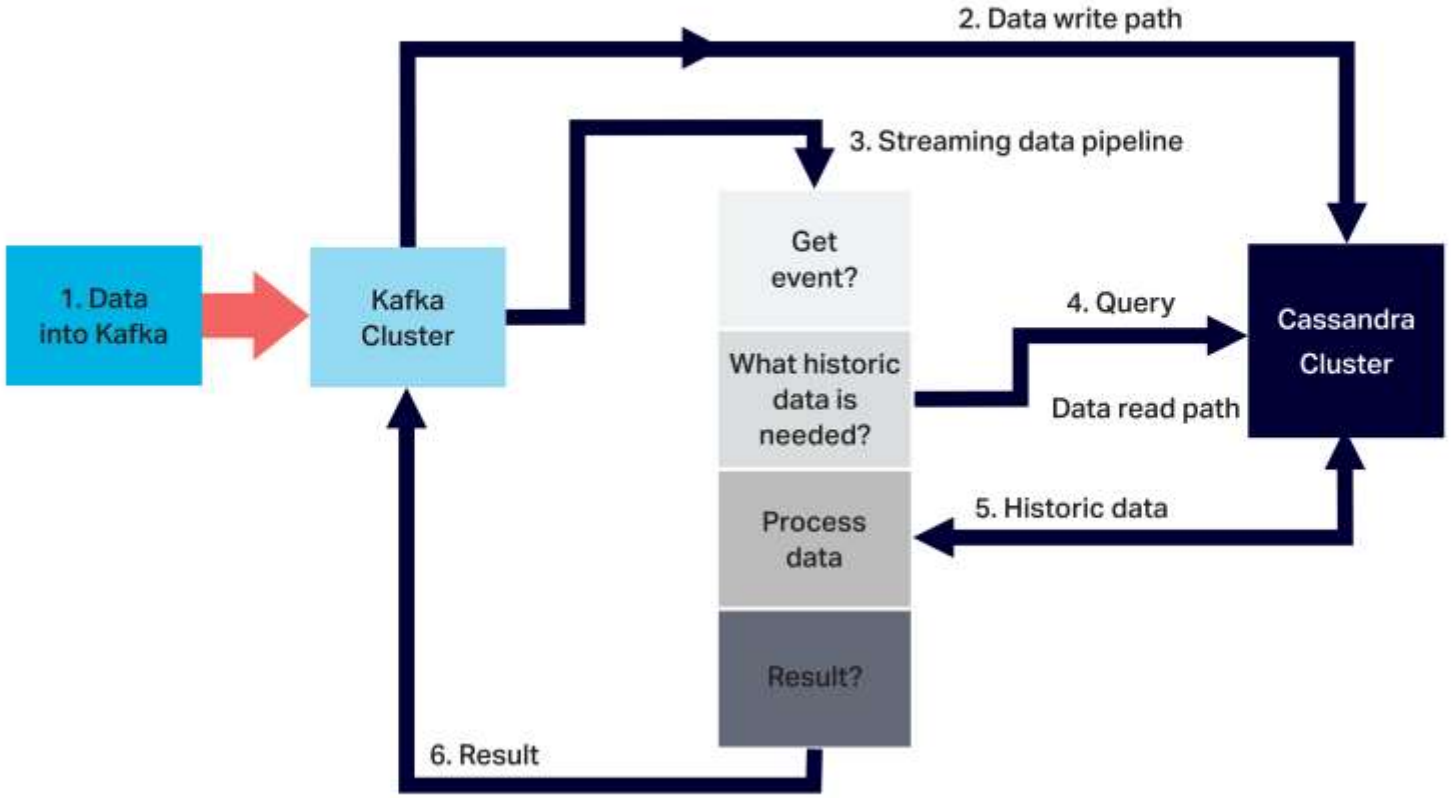

Real-time Event processing pipeline:

- Simple event driven applications (If X then Y…)
- May write and read from other data sources (e.g. Cassandra)
- New Events sent back to Kafka or to other systems
- E.g. Anomaly Detection, check out my current blog series if you are interested in this example.

Kafka Streams Processing (Kongo IoT Blog series)

Streams processing features:

- Complex streams processing (multiple events and streams)
- Time, windows, and transformations
- Uses Kafka Streams API, includes state store
- Visualization of the streams topology
- Continuously computes the loads for trucks and checks if they are overloaded.

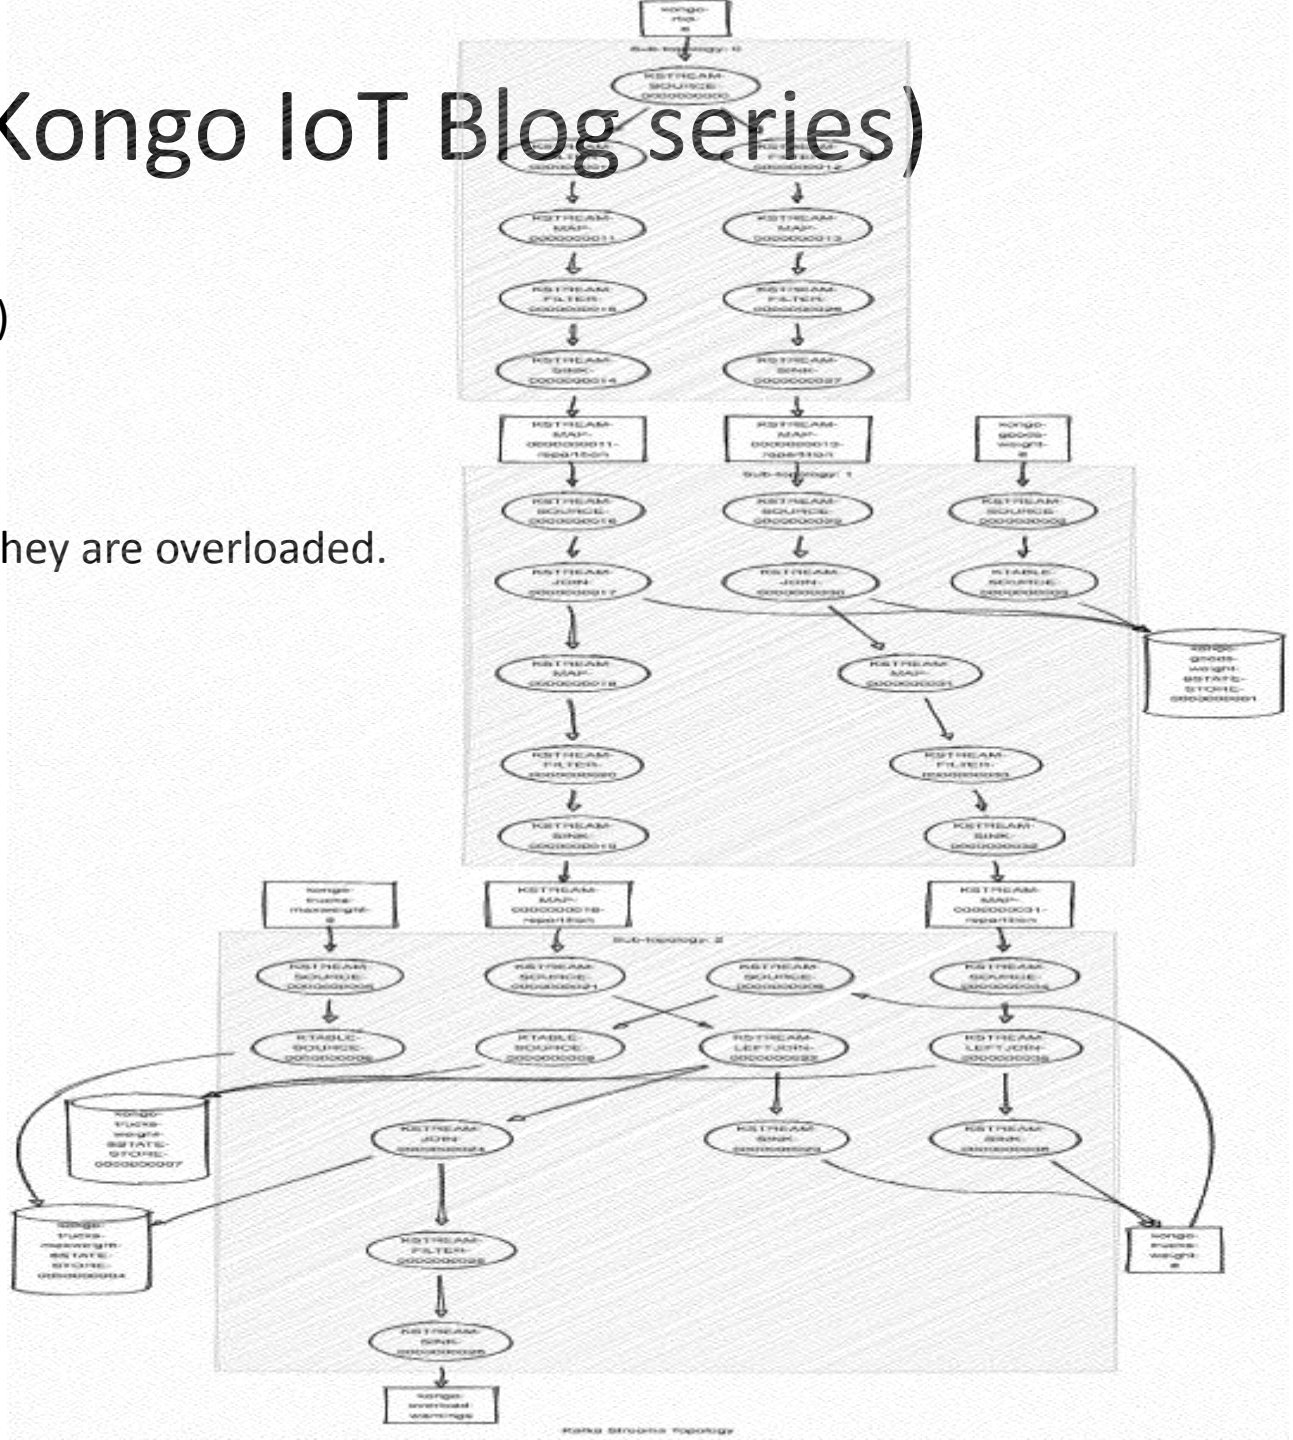

Linkedin - Before Kafka (BK)

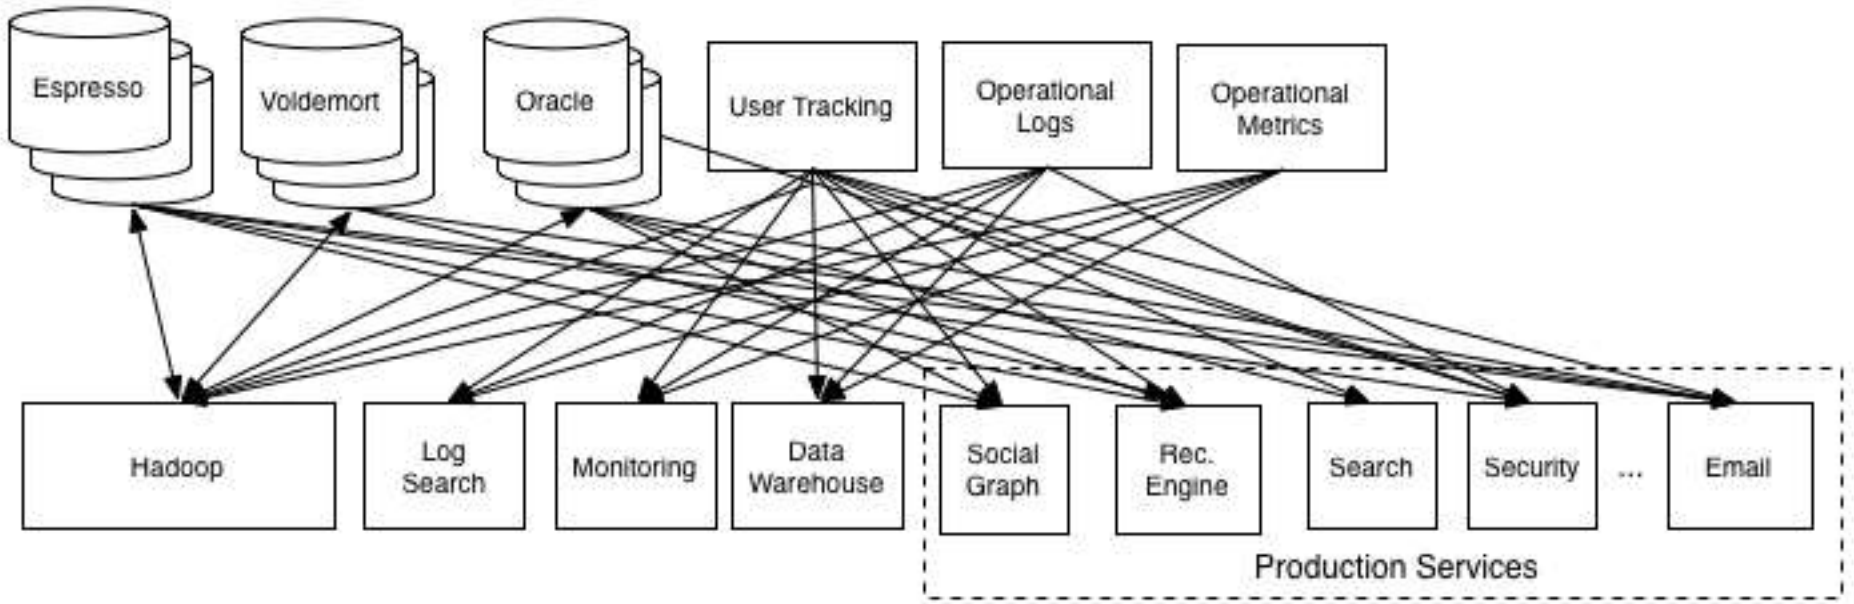

A real example from Linkedin, who developed Kafka.
Before Kafka they had spaghetti integration of monolithic applications.
To accommodate growing membership and increasing site complexity, they migrated from a monolithic application infrastructure to one based on microservices, *which made the integration even more complex!*

After Kafka (AK)

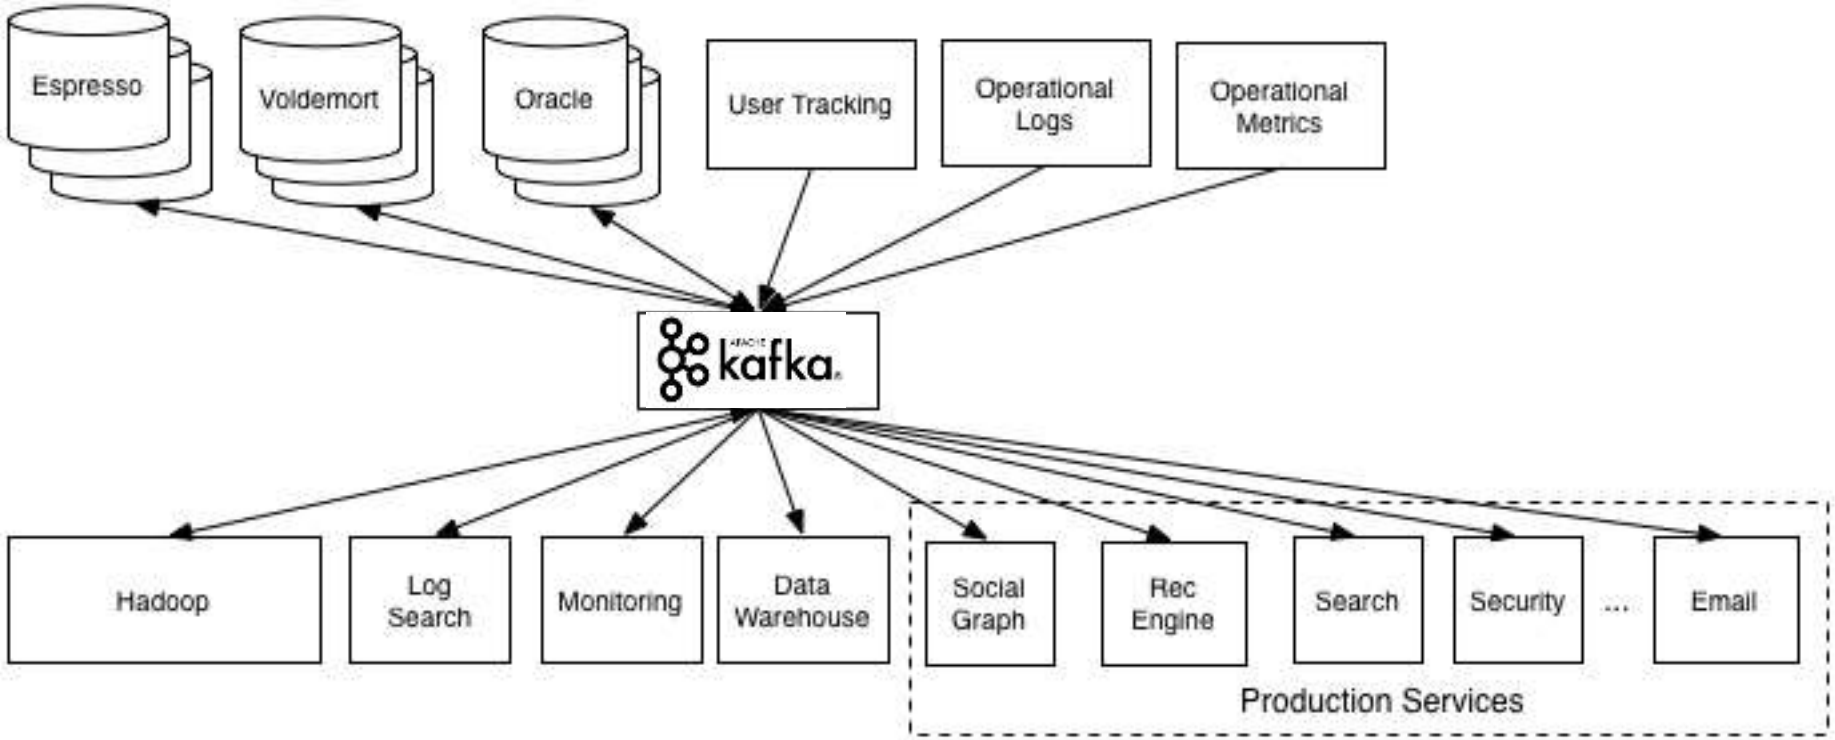

Rather than maintaining and scaling each pipeline individually, they invested in the development of a single, distributed pub-sub platform - Kafka was born.
The main benefit was *better Service decoupling and independent scaling*.

The End (of the introduction) -

Find out more

instaclustr

Apache Kafka: https://kafka.apache.org/

Instaclustr blogs

- Mix of Cassandra, Spark, Zeppelin and Kafka

https://www.instaclustr.com/paul-brebner/

- Kafka introduction

https://insidebigdata.com/2018/04/12/developing-deeper-understanding-apache-kafka-architecture/
https://insidebigdata.com/2018/04/19/developing-deeper-understanding-apache-kafka-architecture-part-2-

- Kongo – Kafka IoT logistics application blog series

https://www.instaclustr.com/instaclustr-kongo-iot-logistics-streaming-demo-application/

- Anomaly detection with Kafka and Cassandra (and Kubernetes), current blog series

https://www.instaclustr.com/anomalia-machina-1-massively-scalable-anomaly-detection-with-apache-kafka-

Instaclustr's Managed Kafka (Free trial)

https://www.instaclustr.com/solutions/managed-apache-kafka/